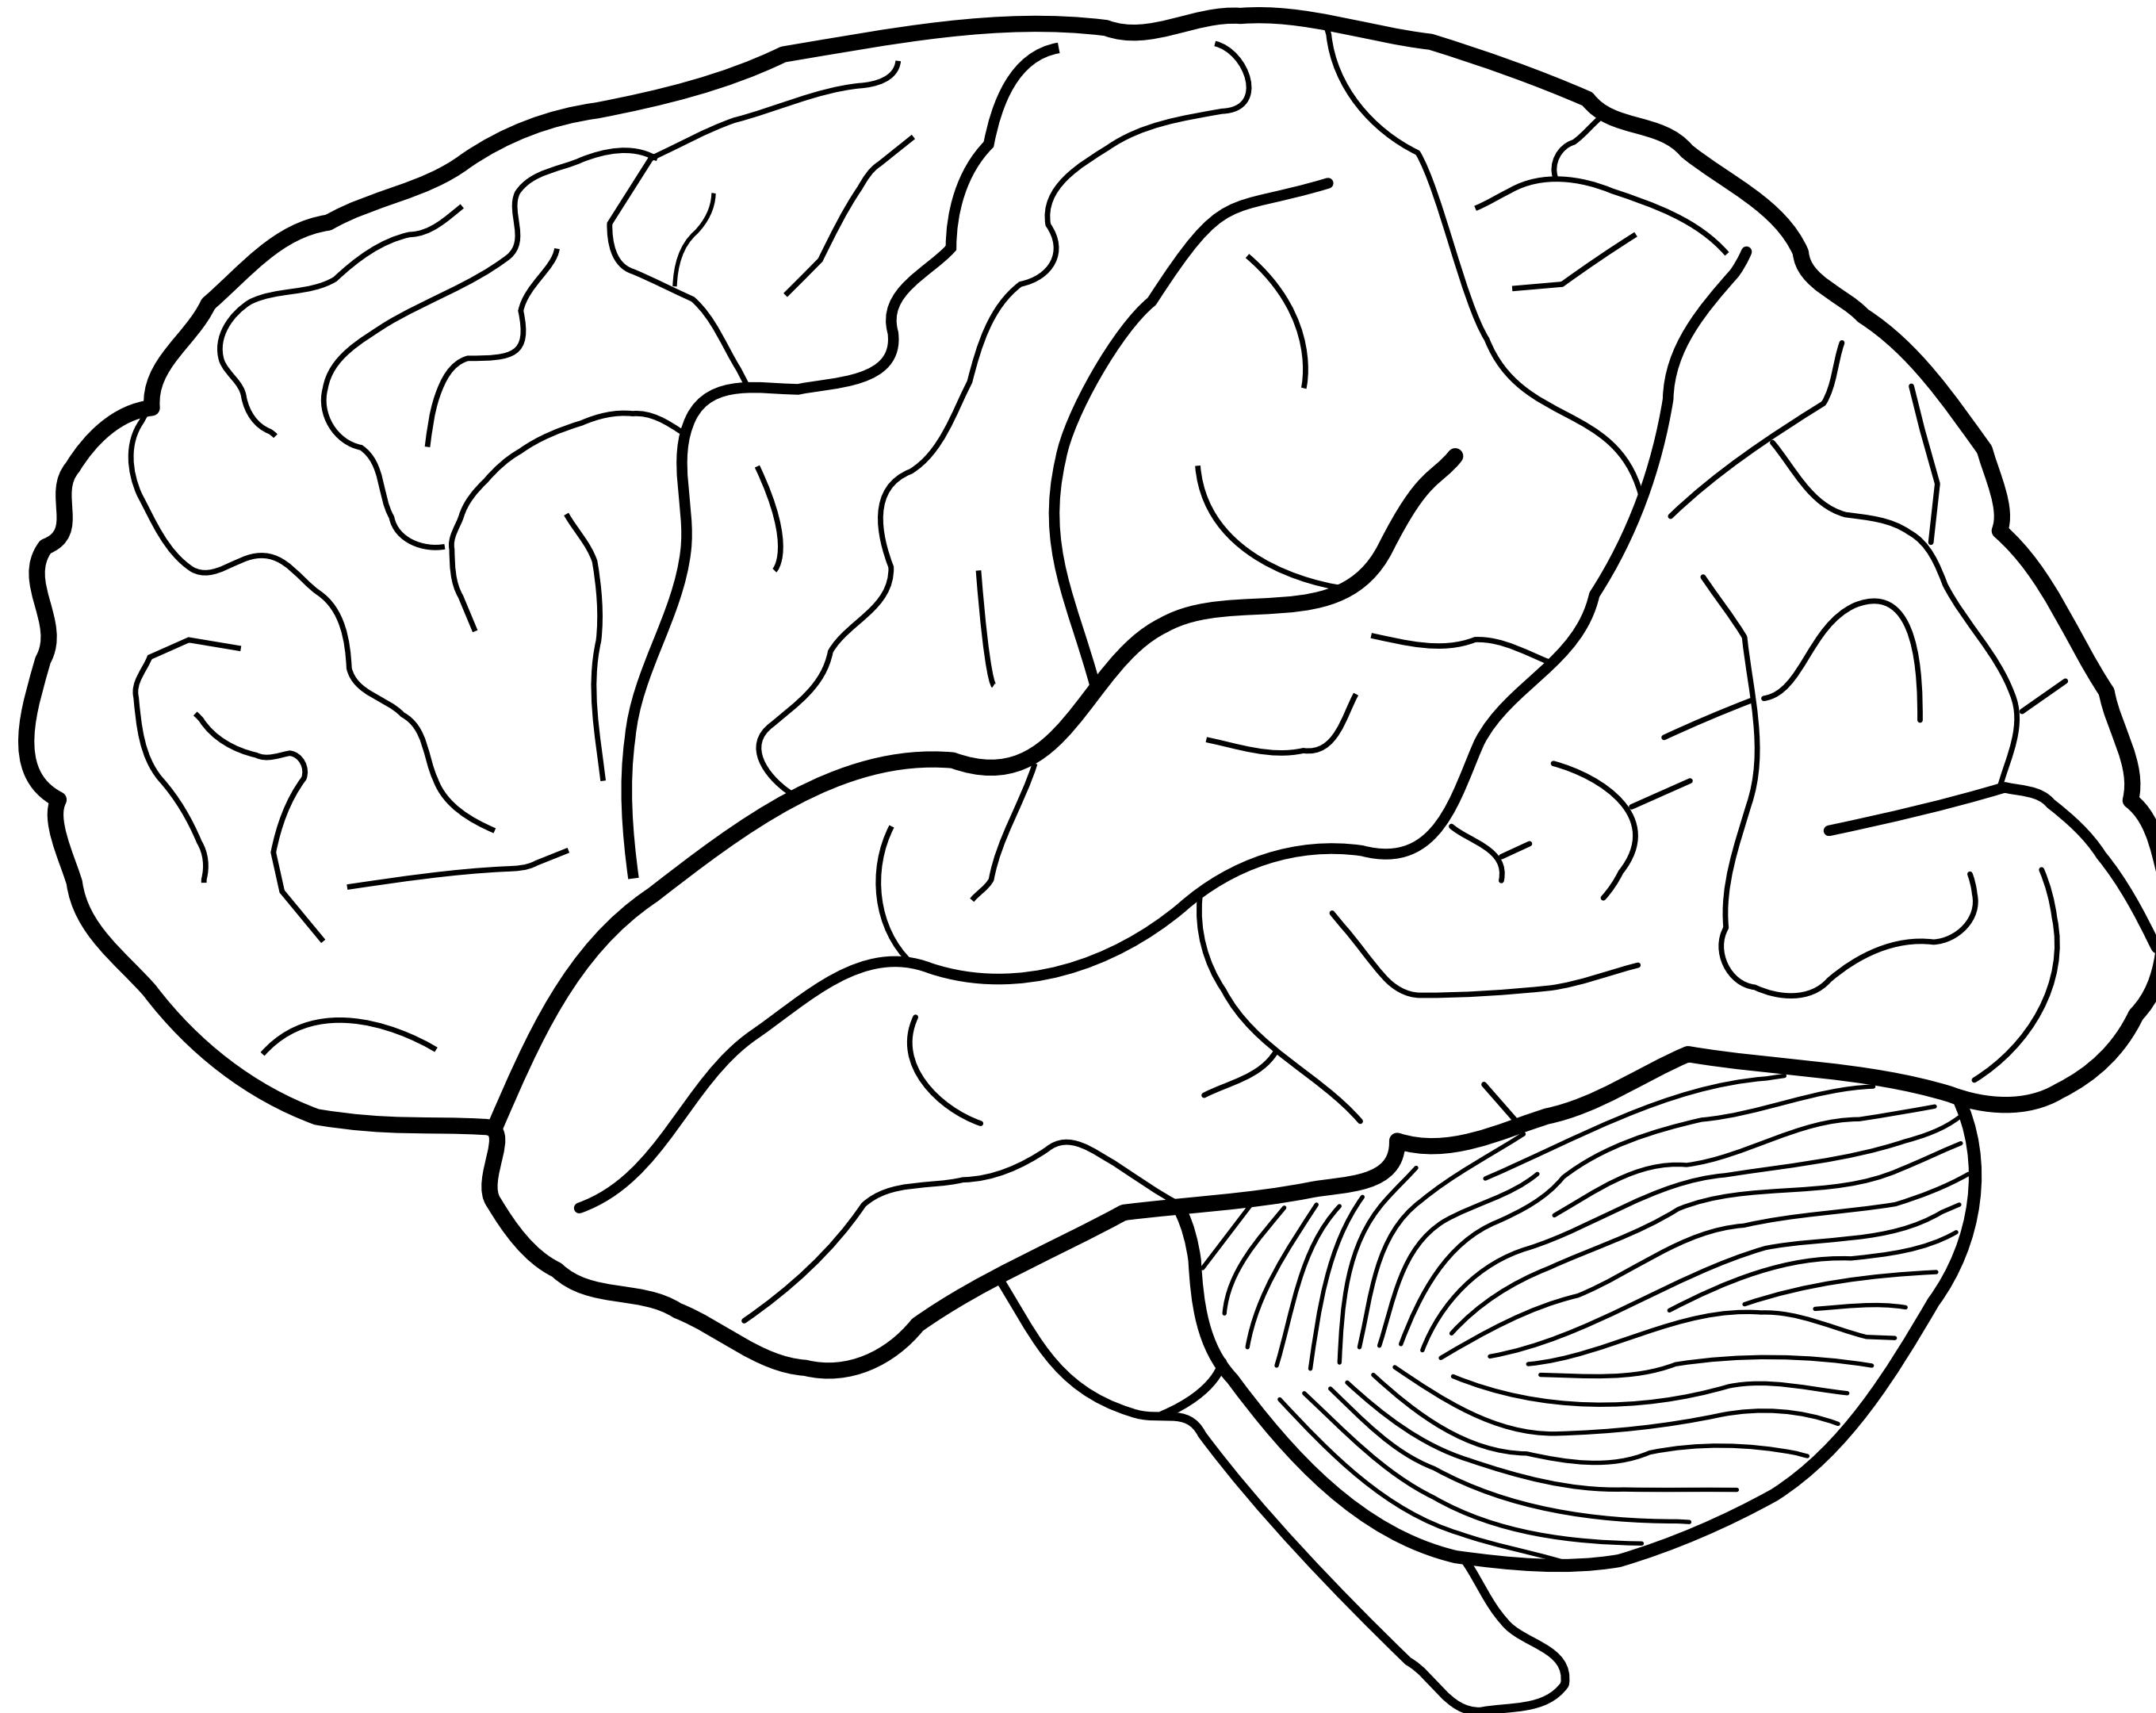


Neuroscience Introduction

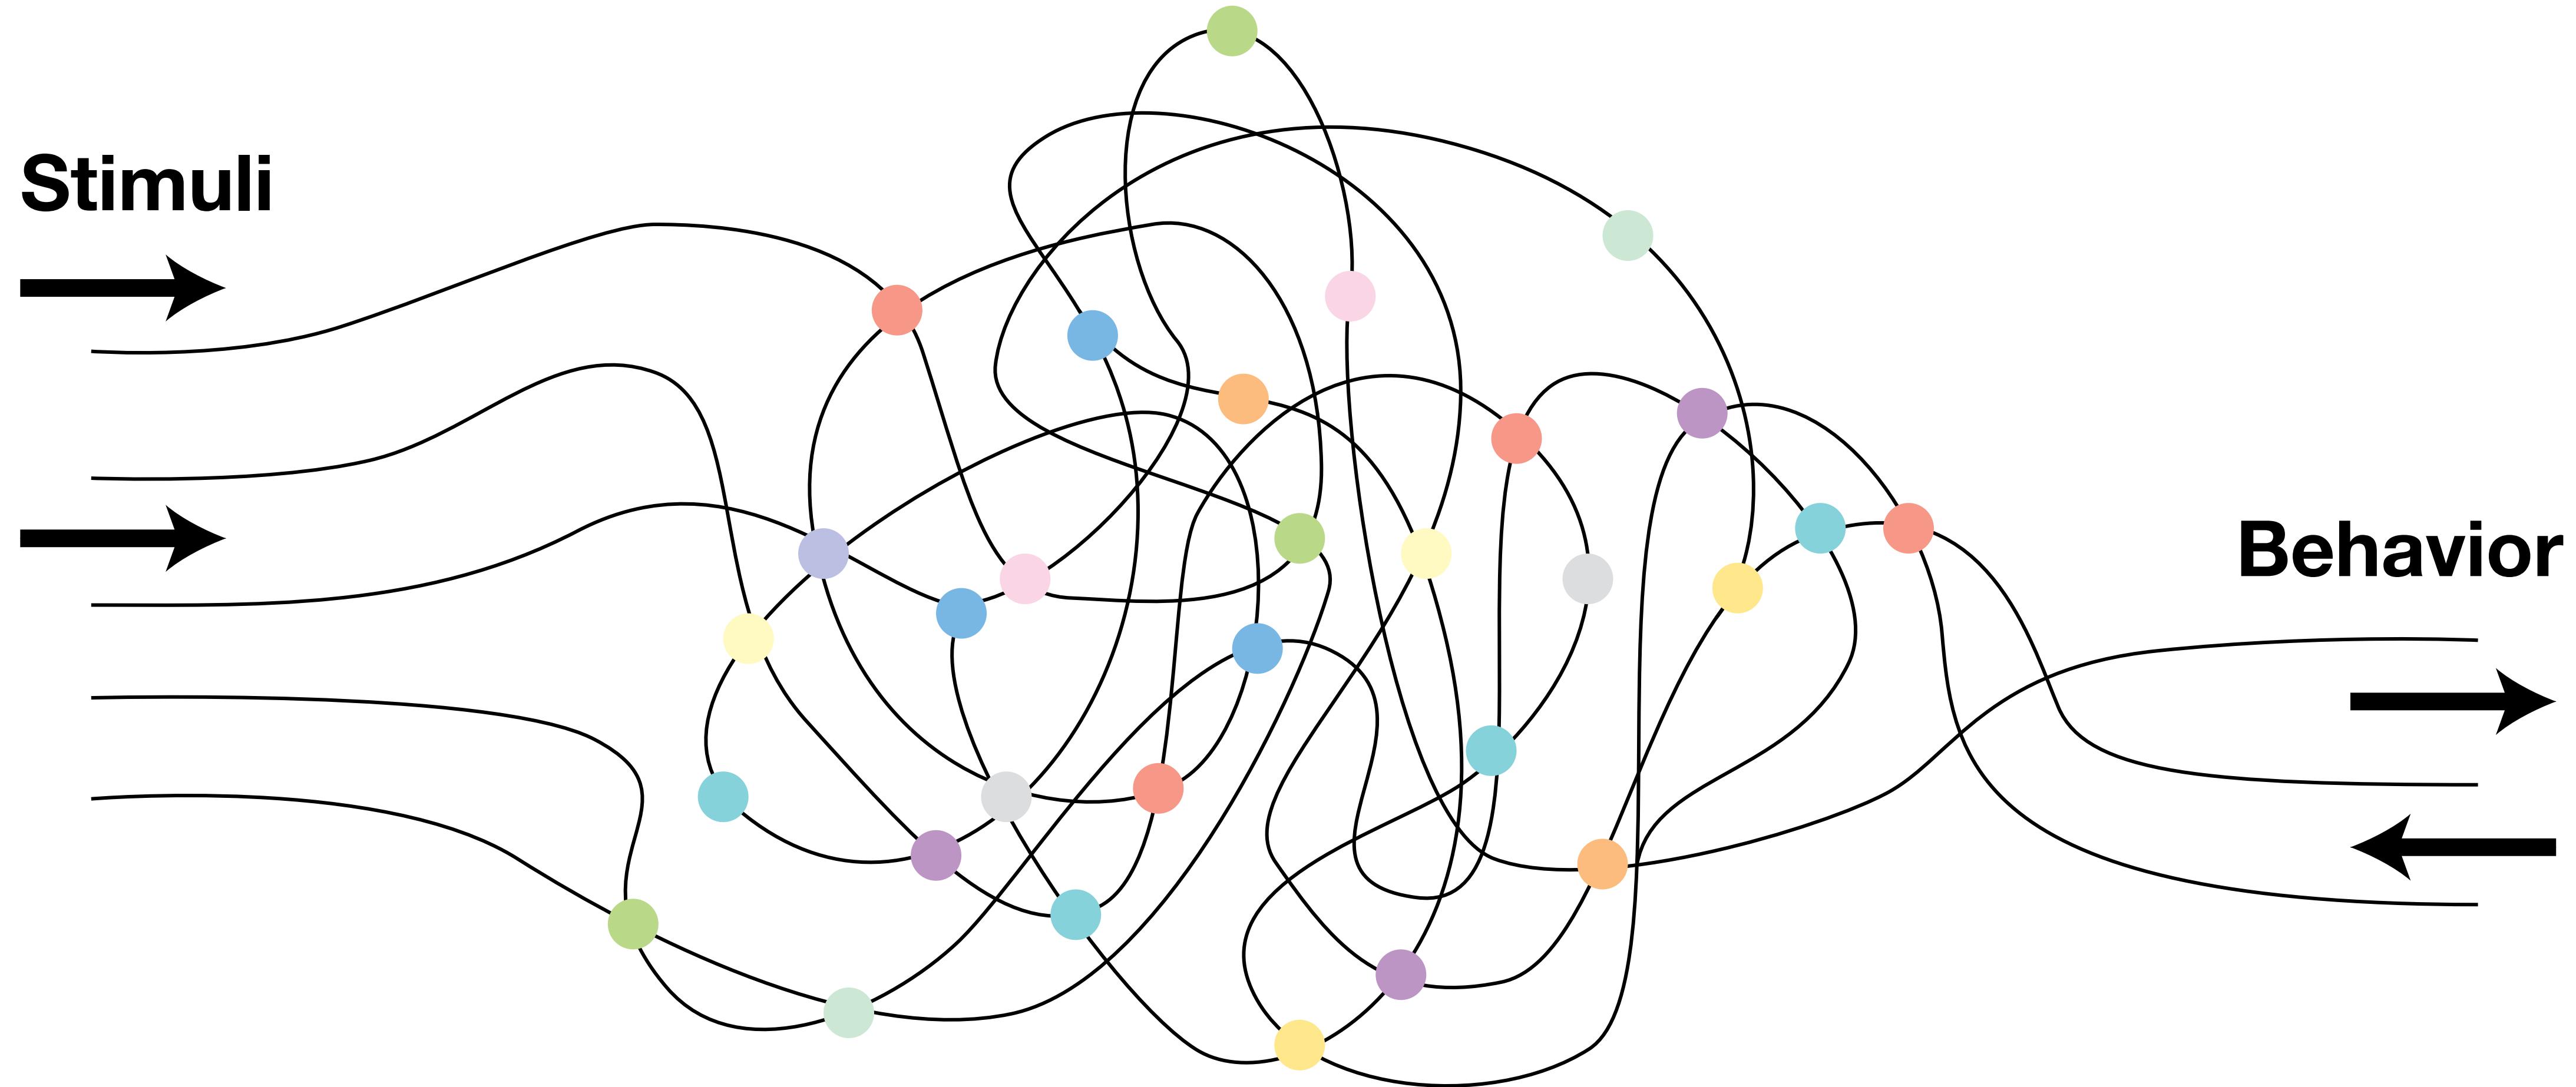
The brain



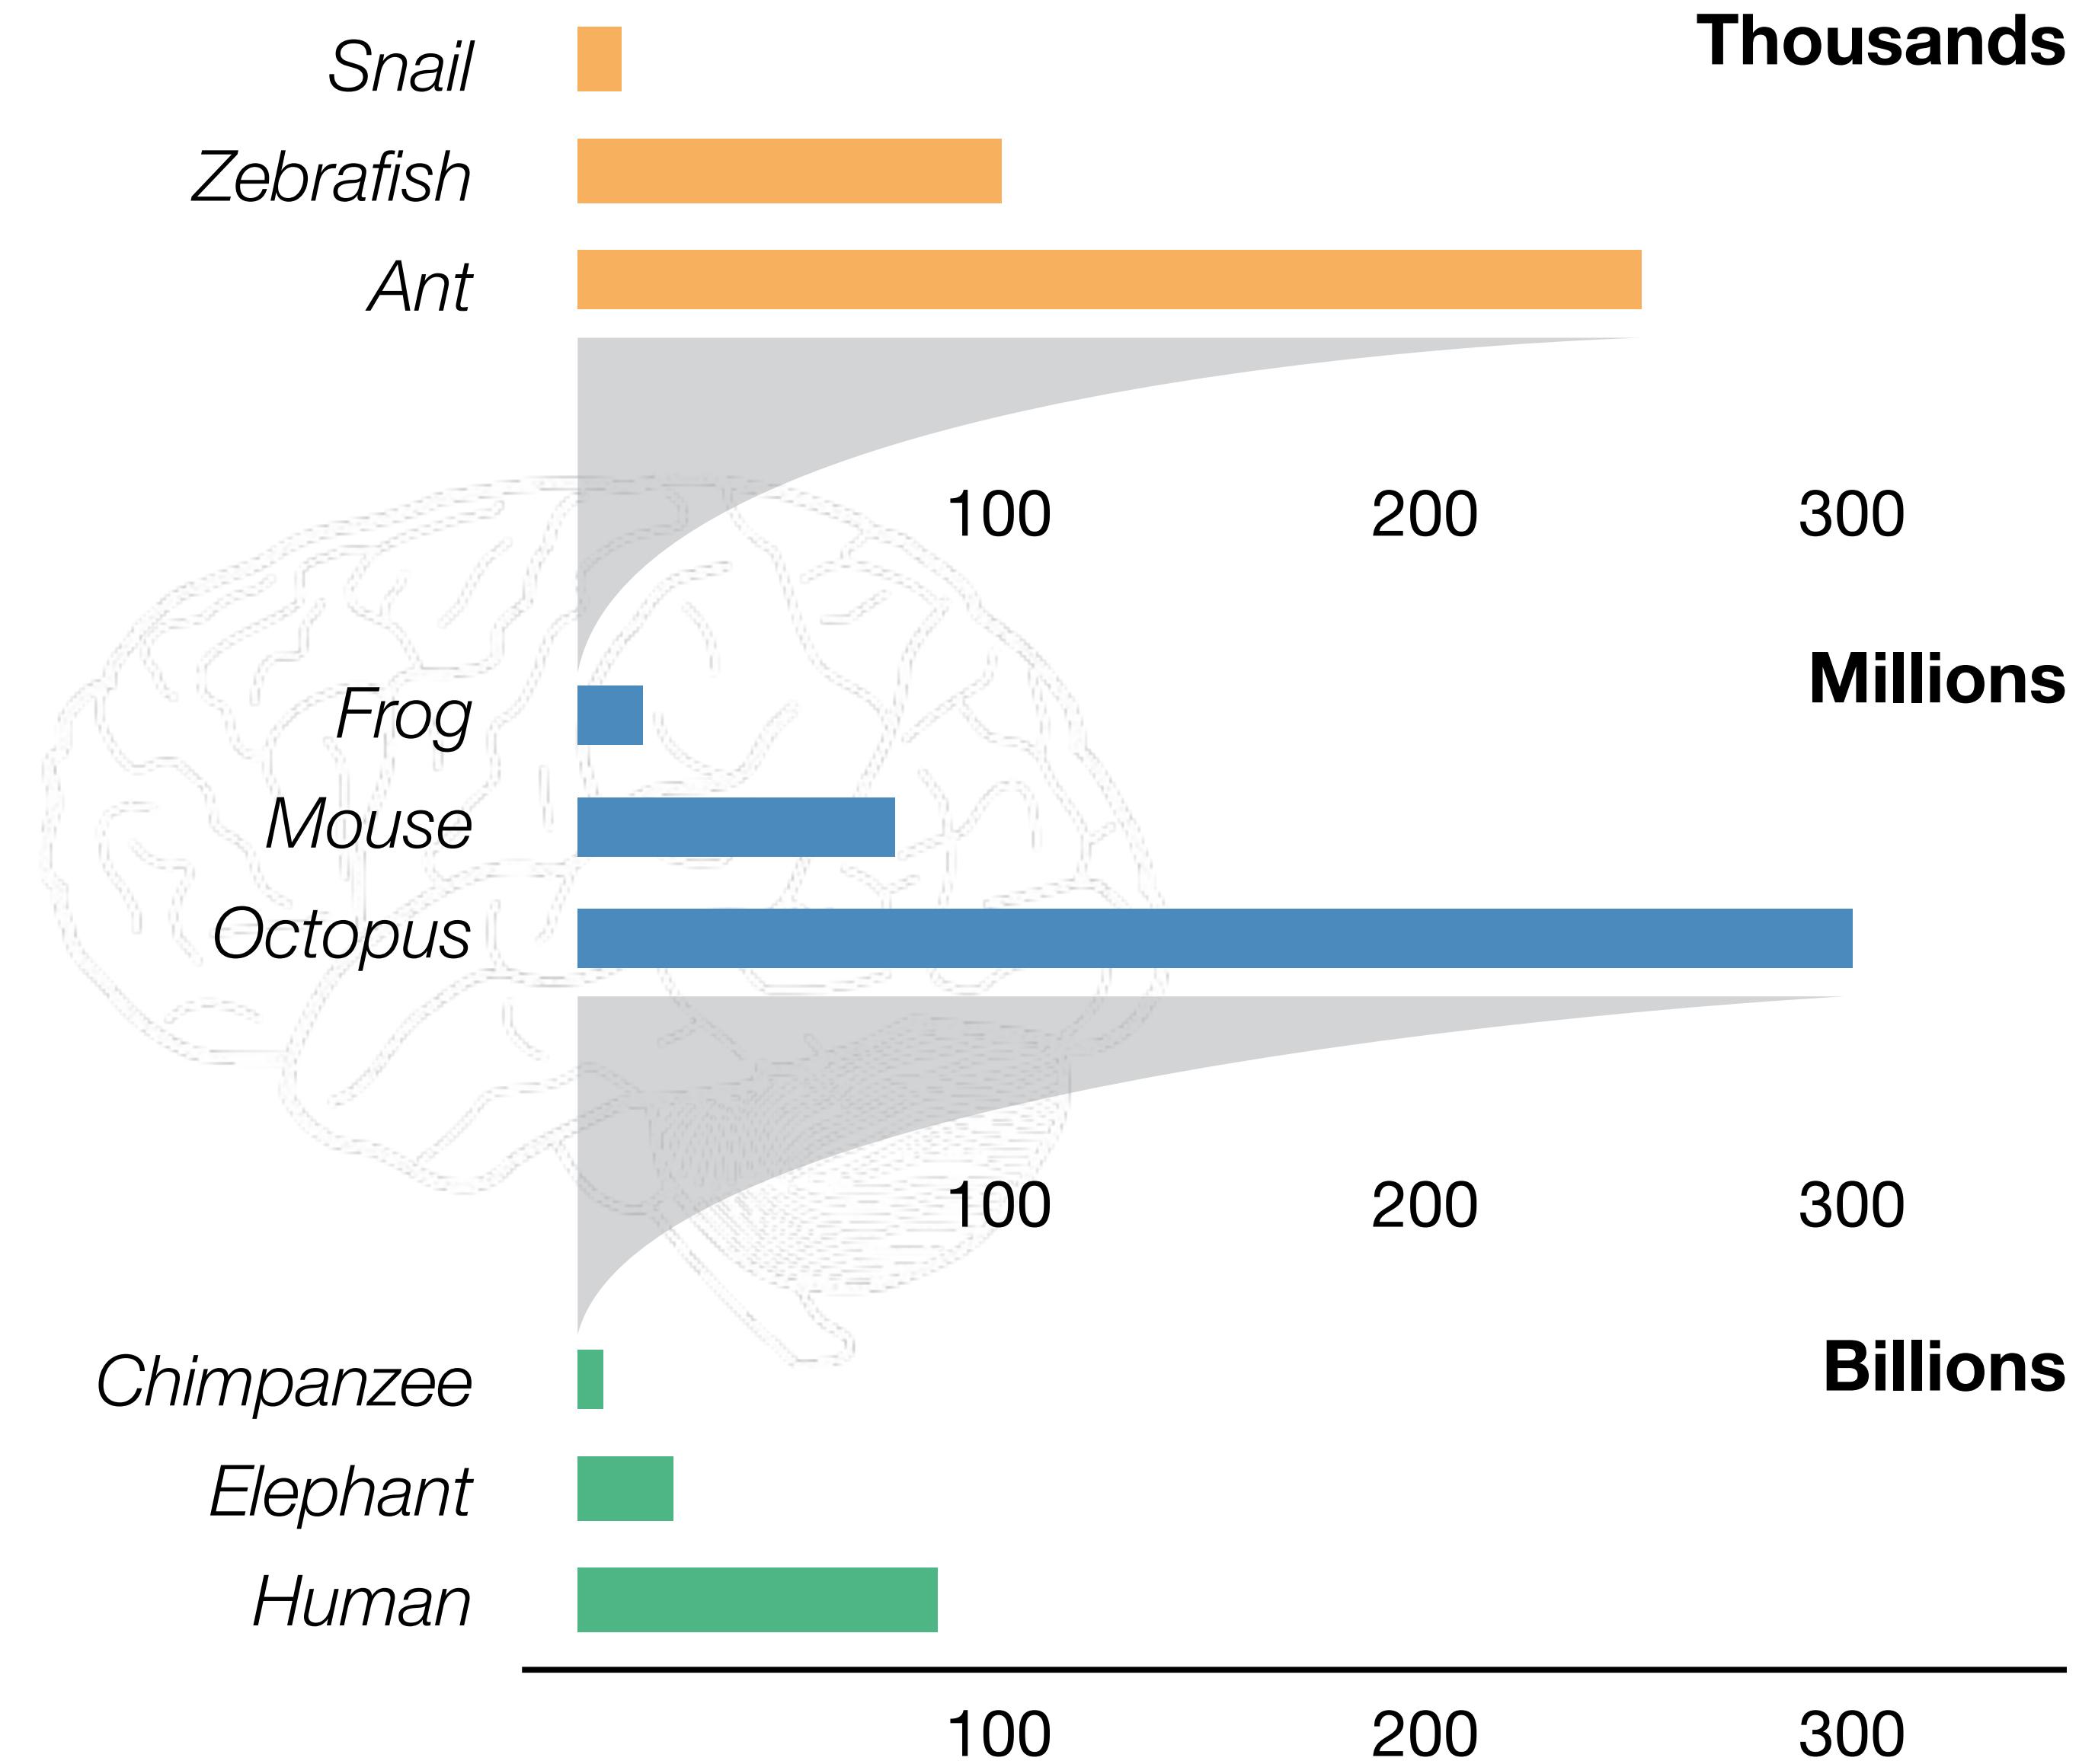
“As humans, we can identify galaxies light years away, we can study particles smaller than an atom. **But we still haven’t unlocked the mystery of the three pounds of matter that sits between our ears.”**

President Obama

The brain



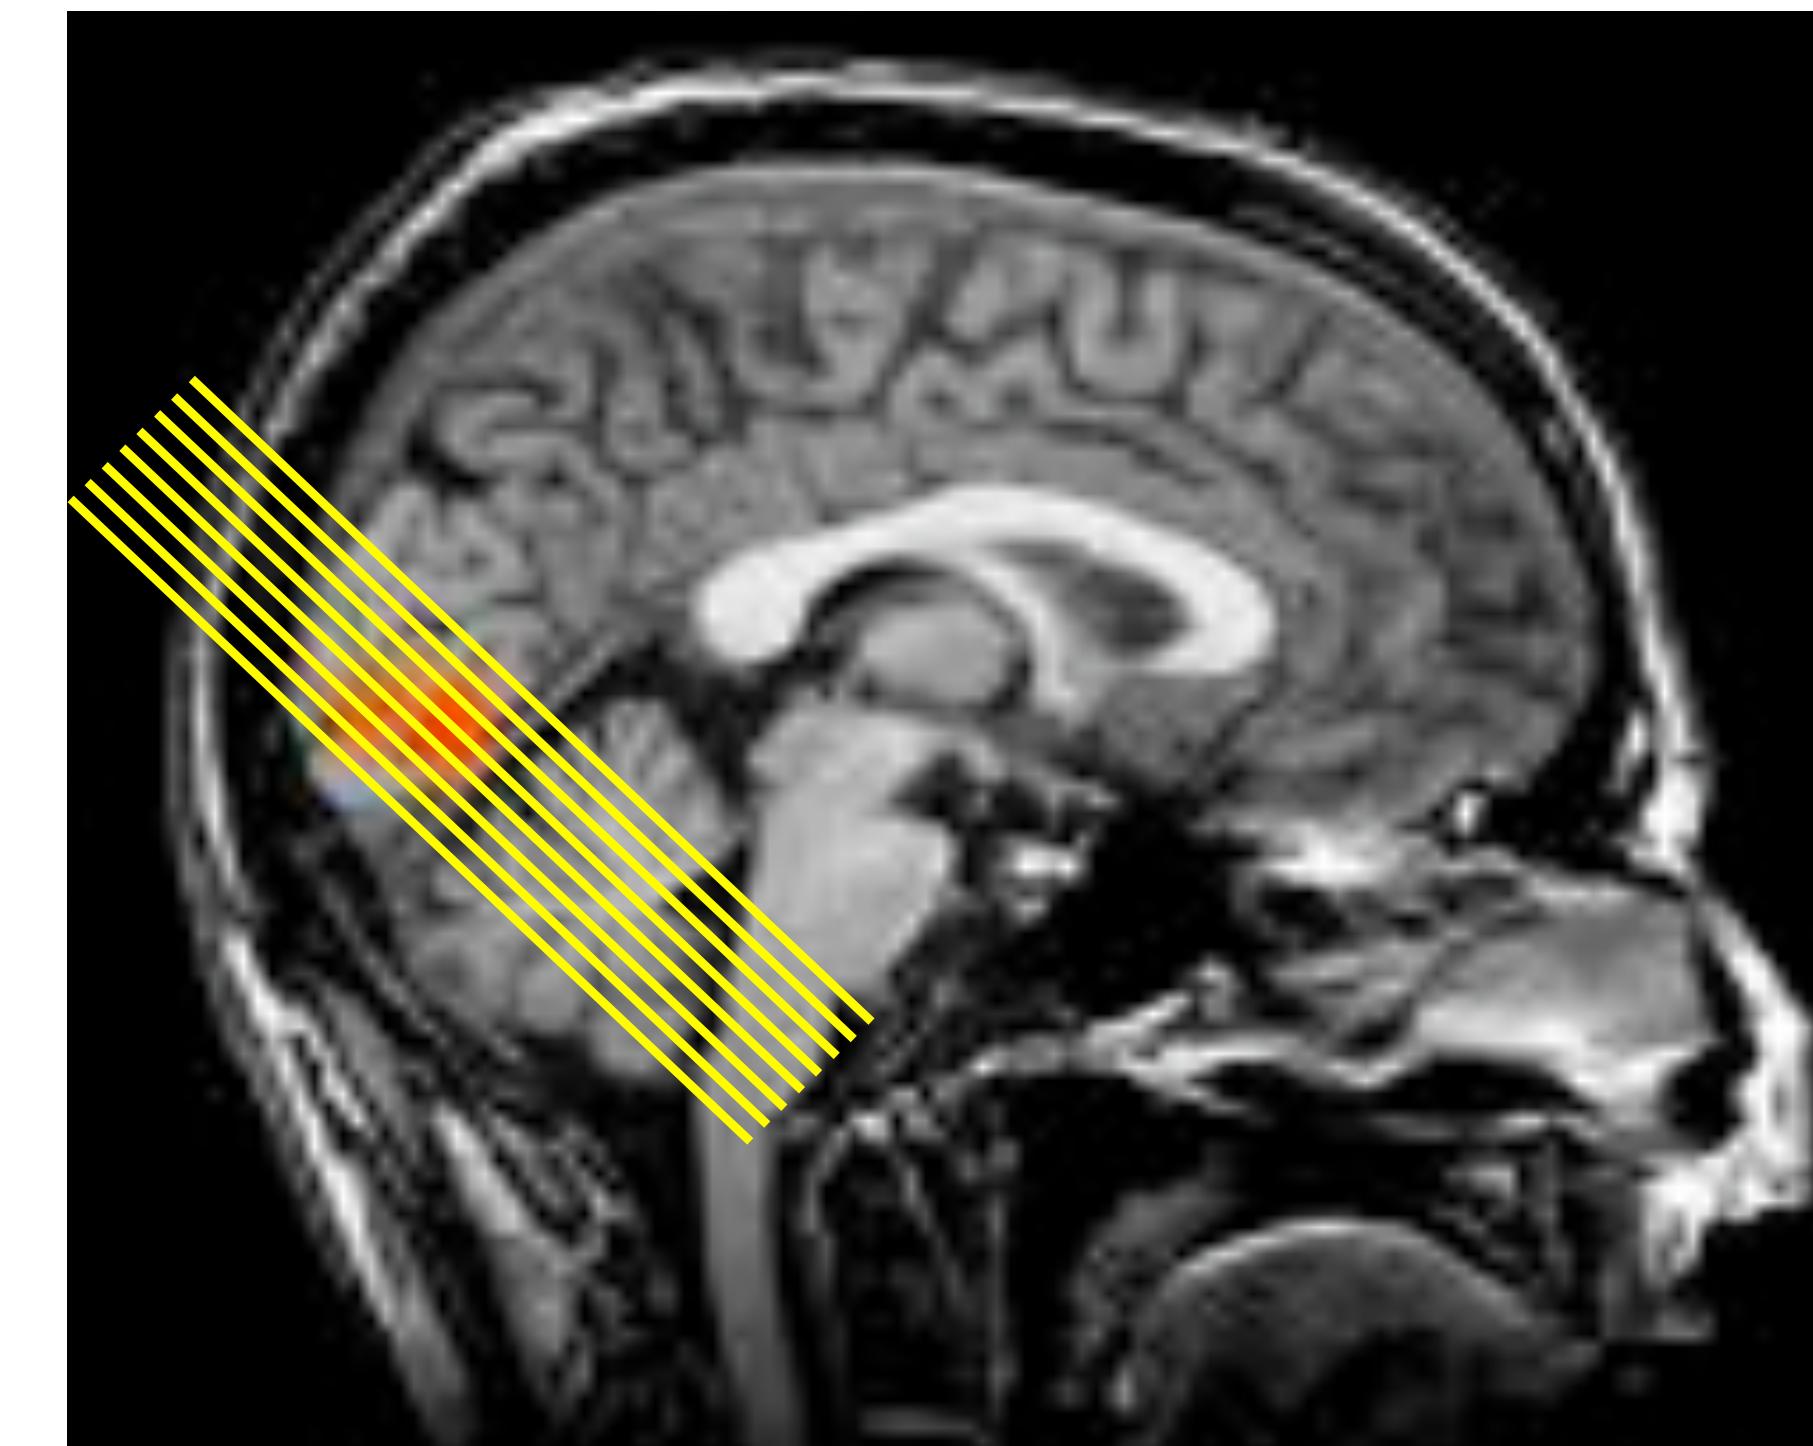
Numbers of neurons



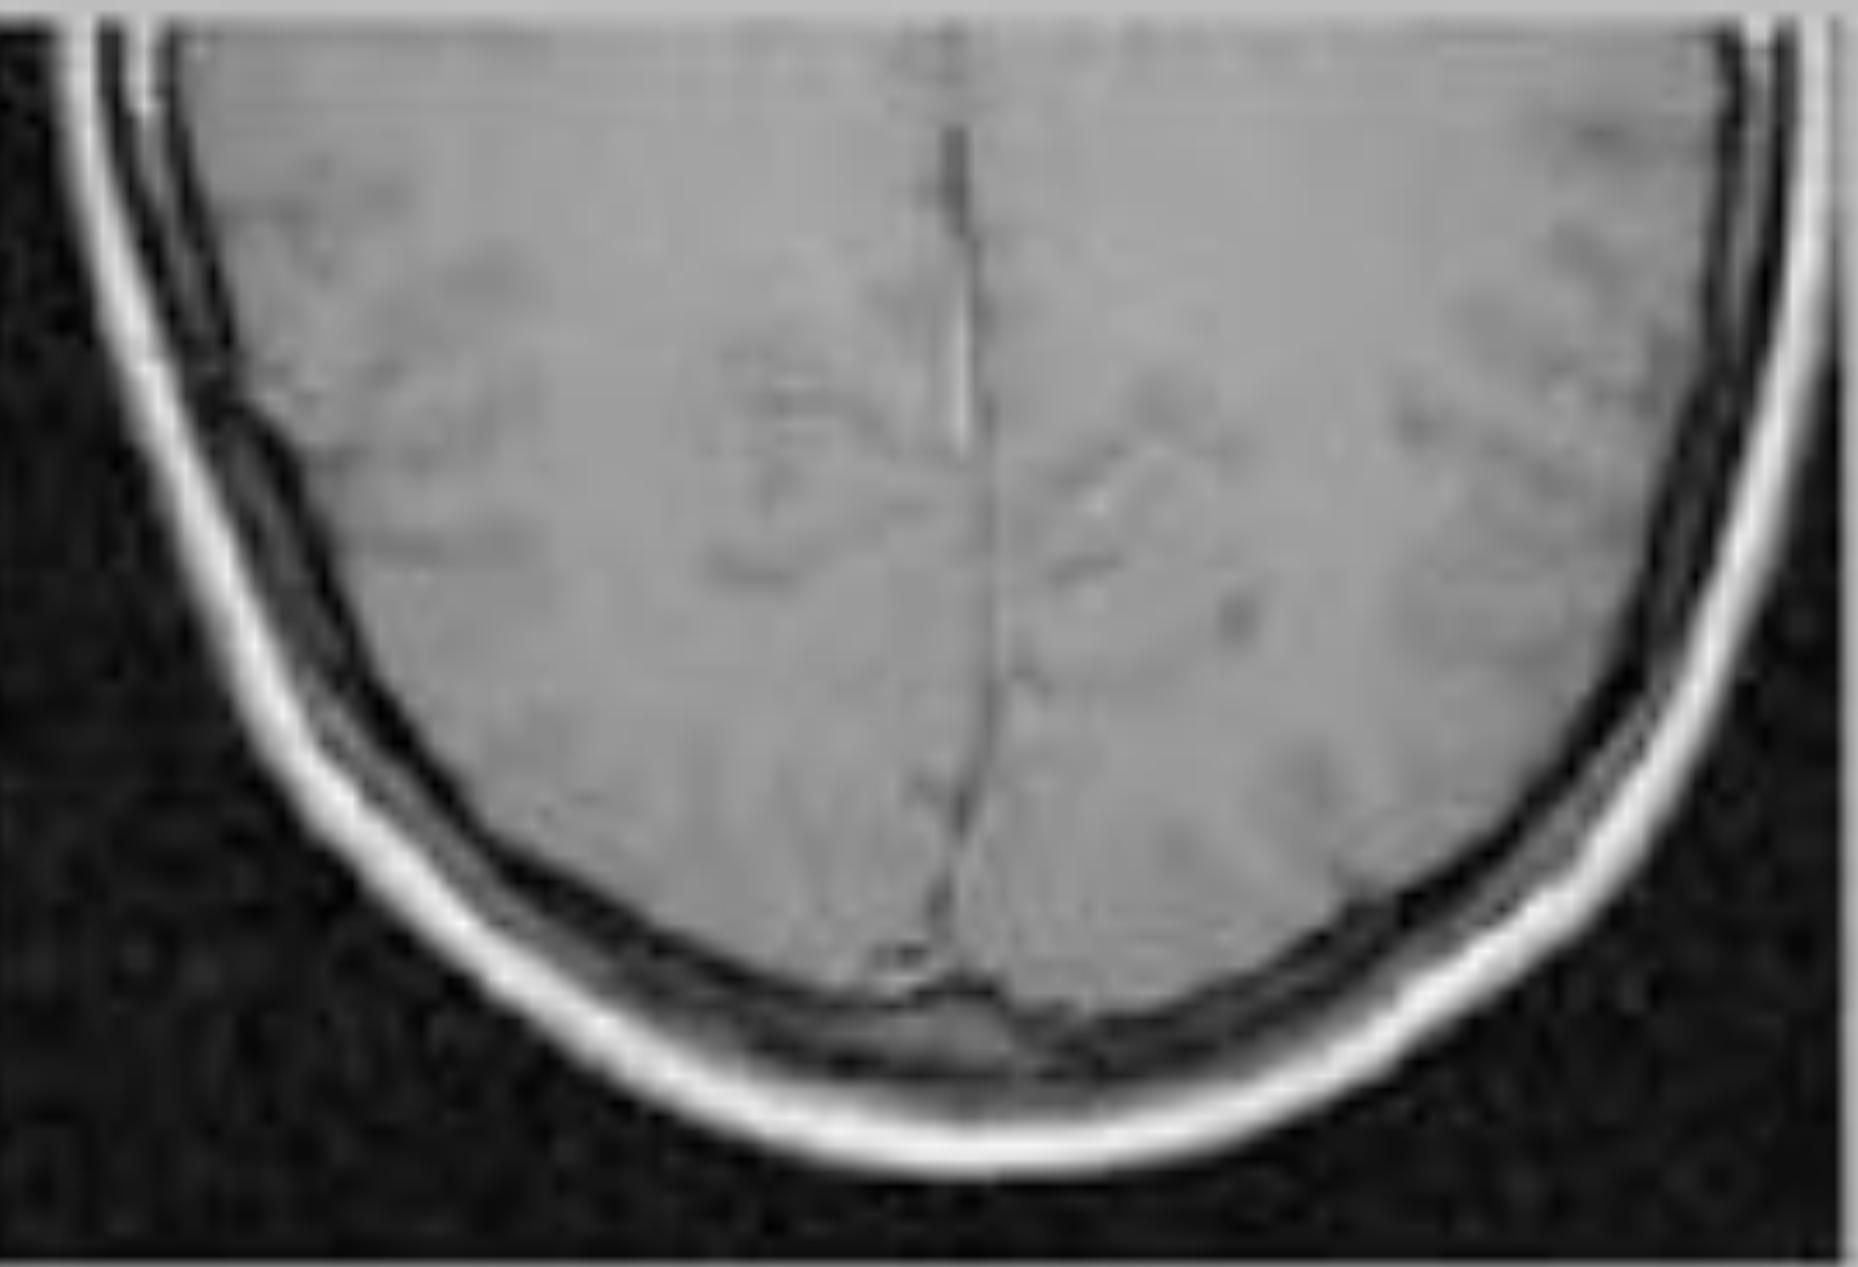
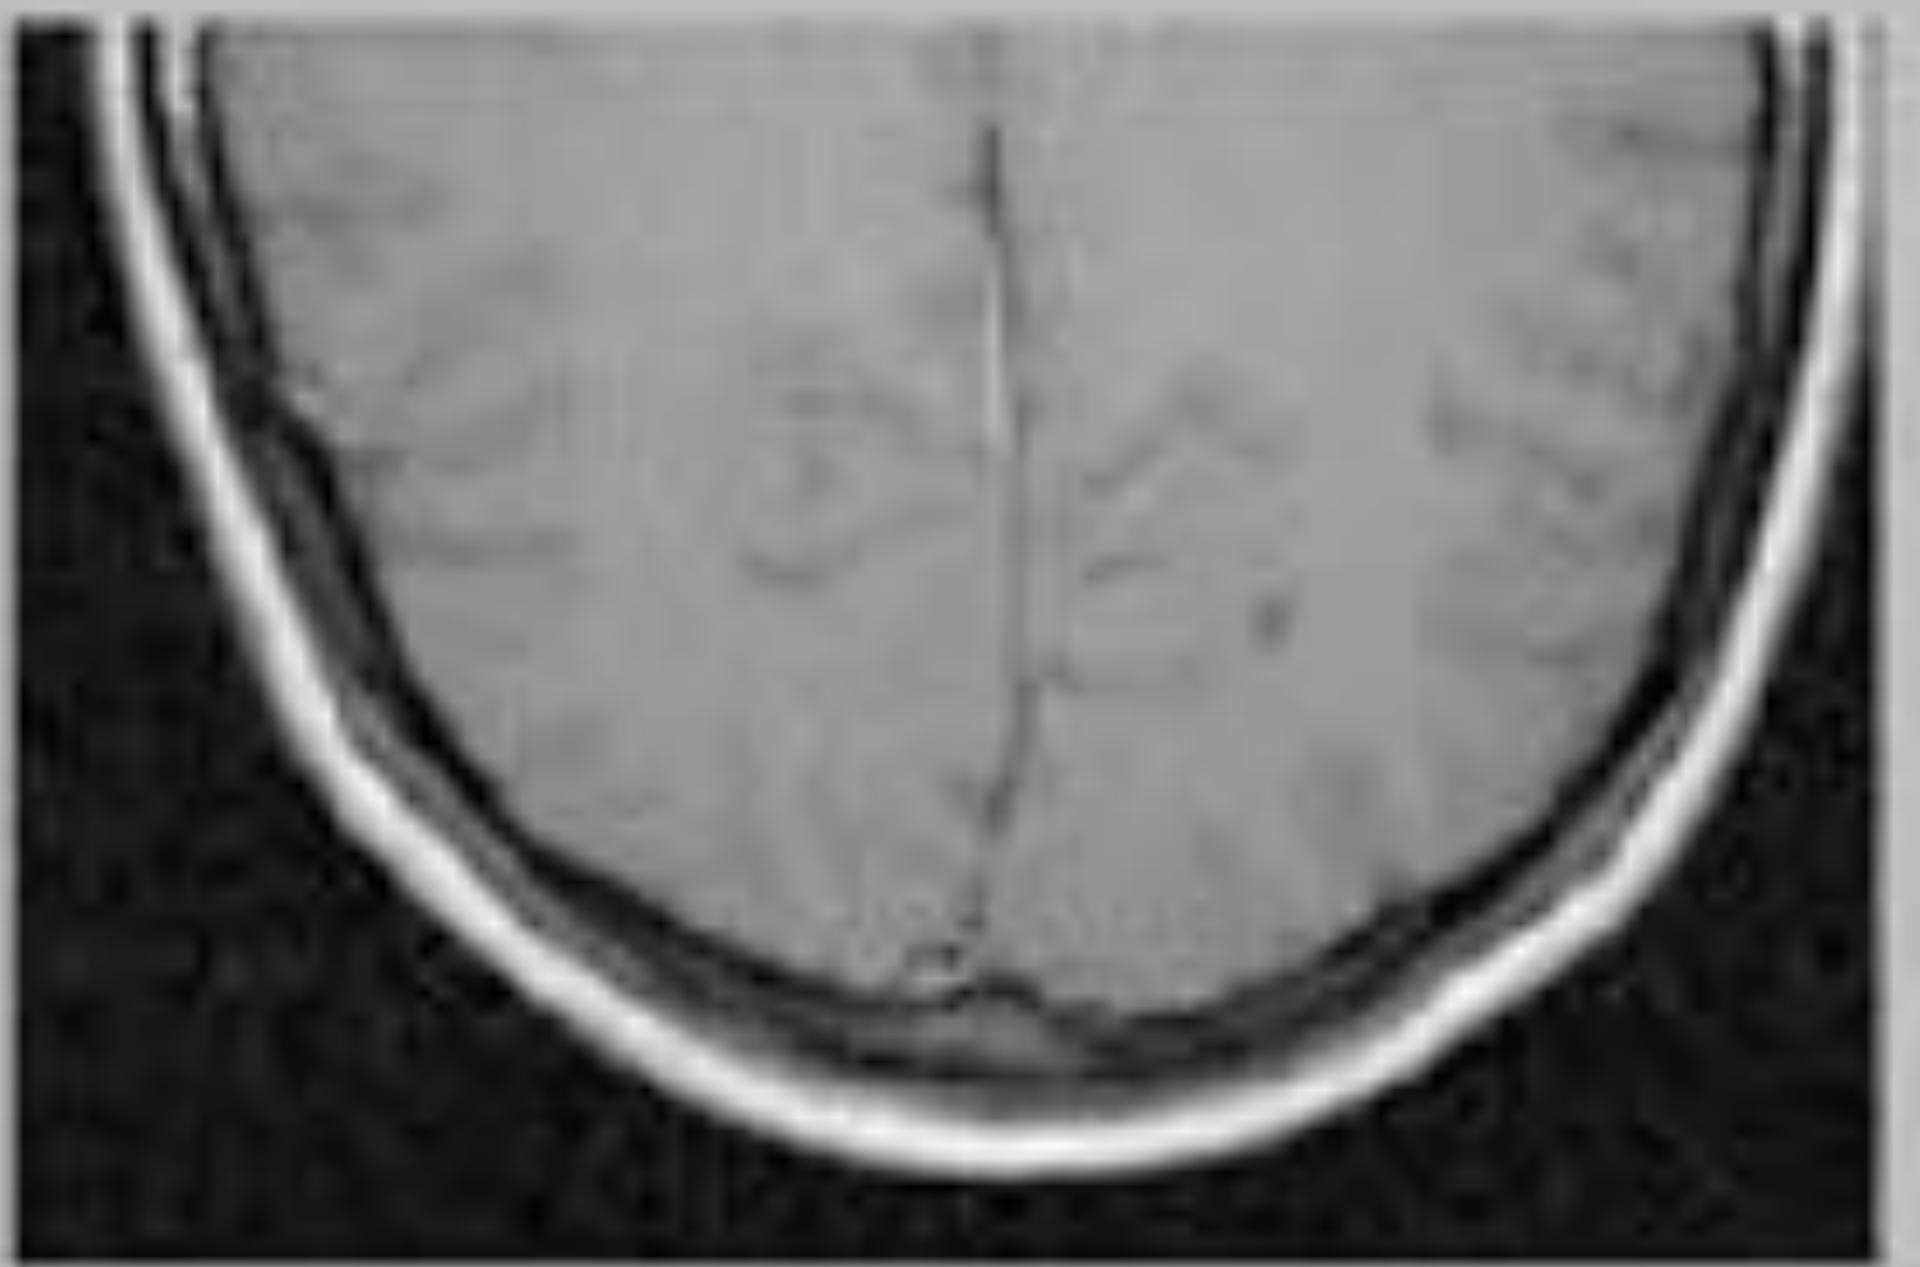
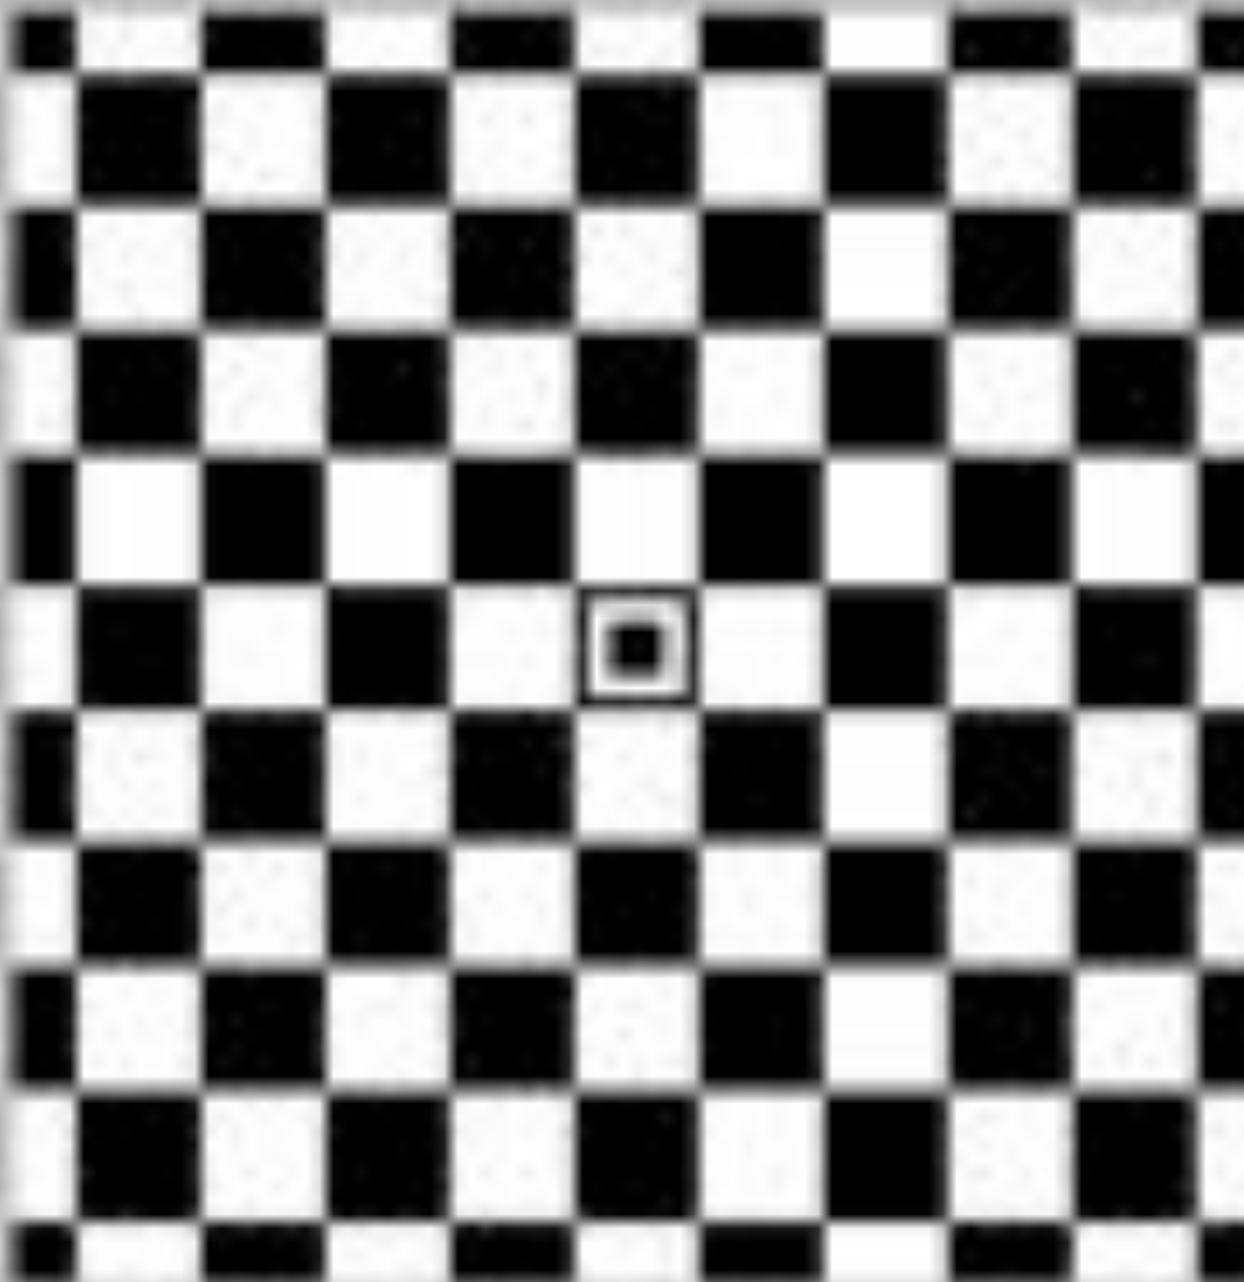
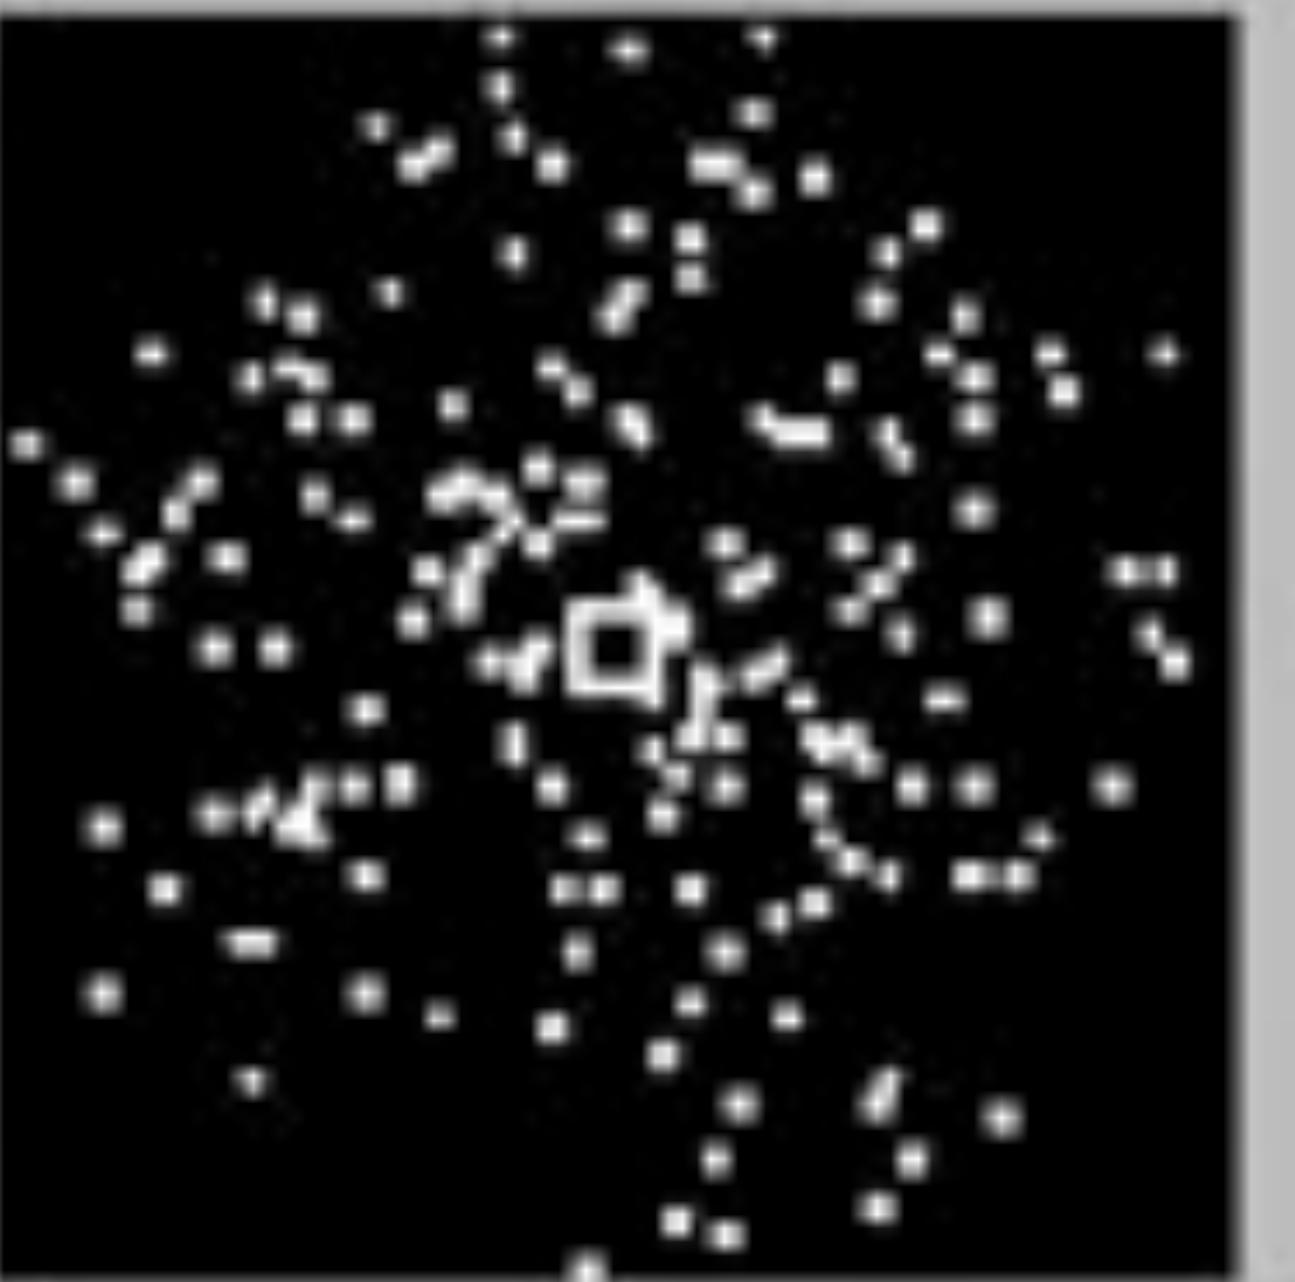
Studying the brain in humans

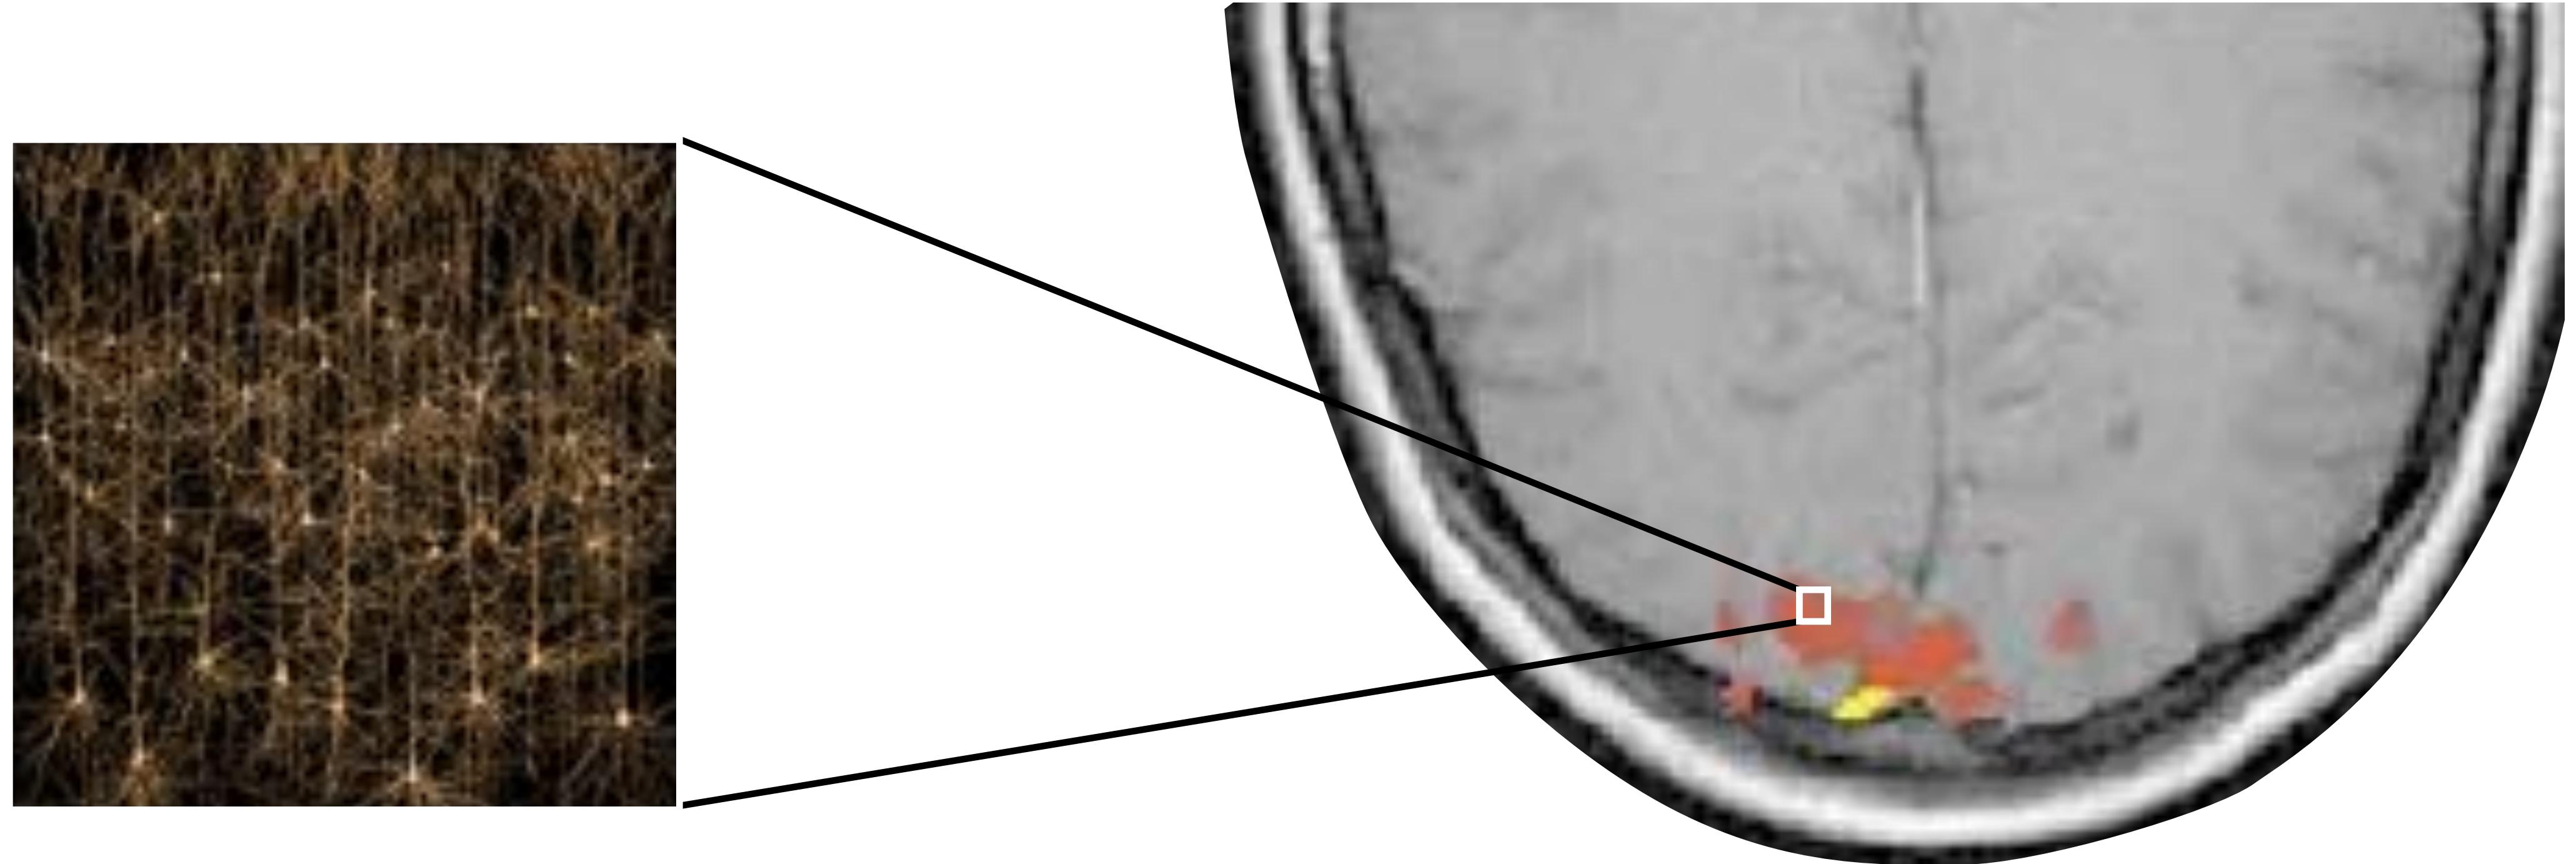


fMRI scanner



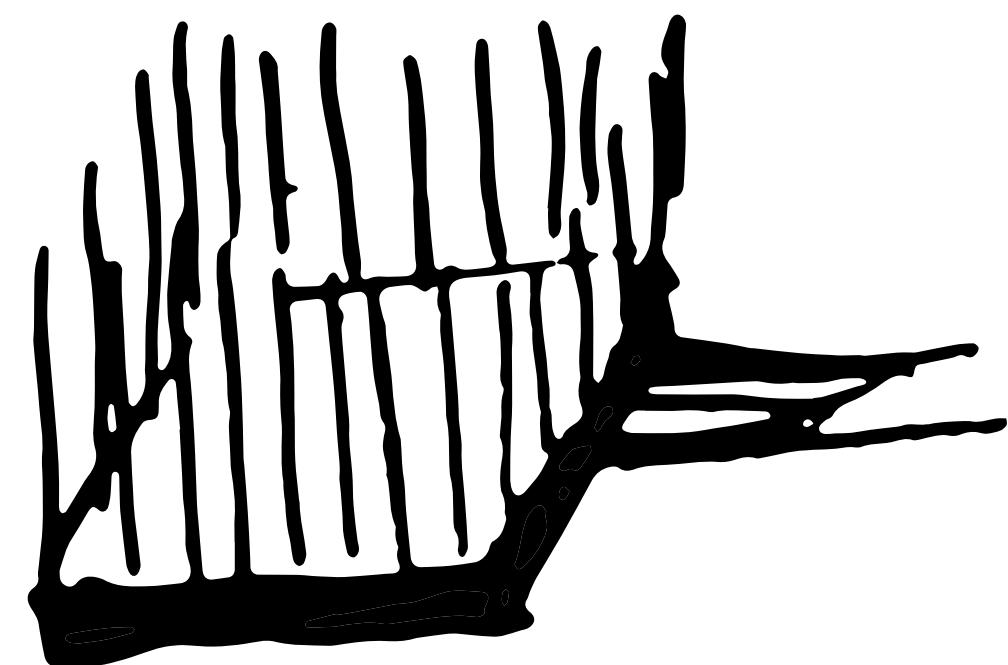
human brain



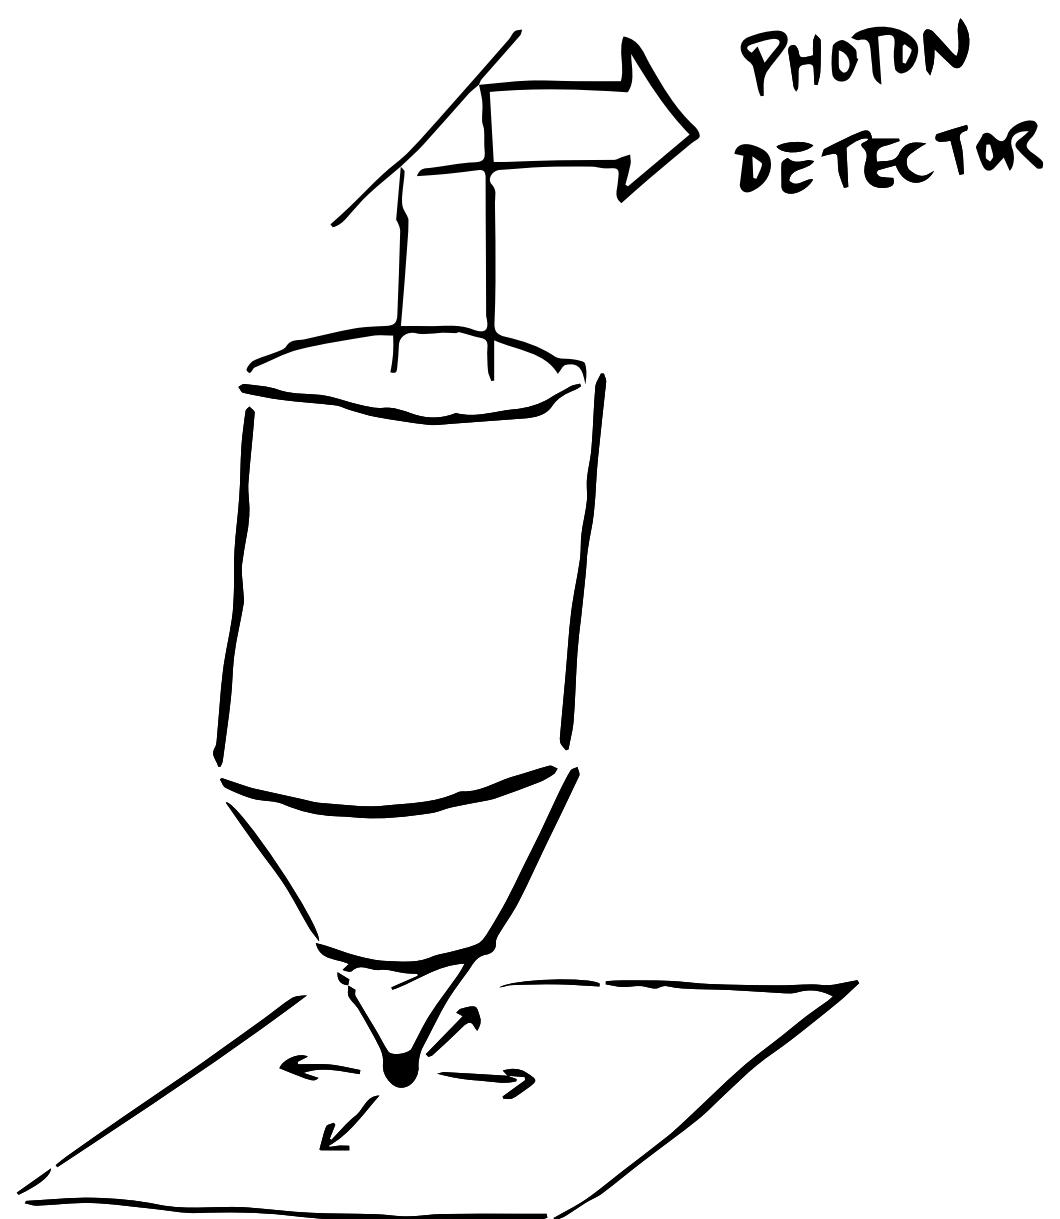


~50,000 neurons per cubic millimeter
-> need higher resolution!

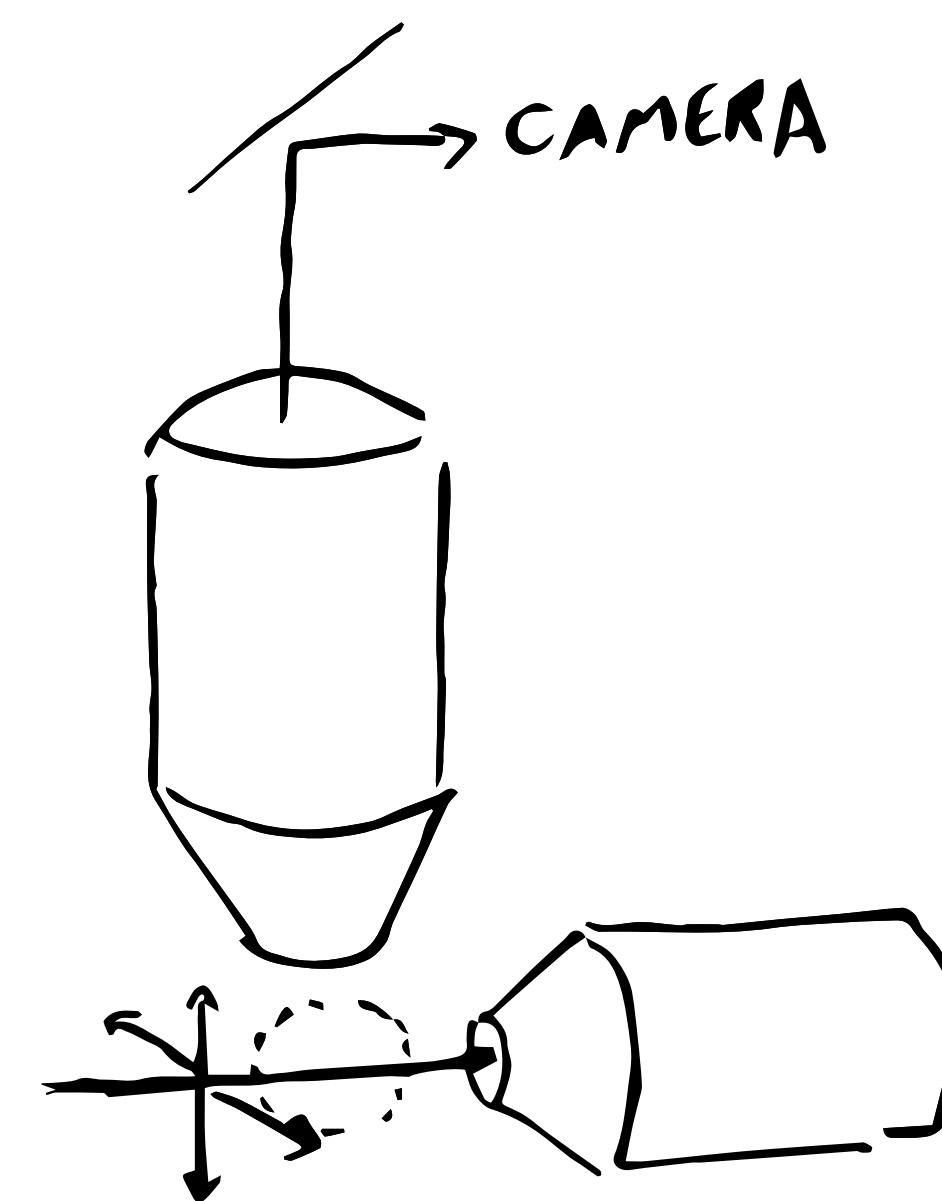
multielectrode
10-100

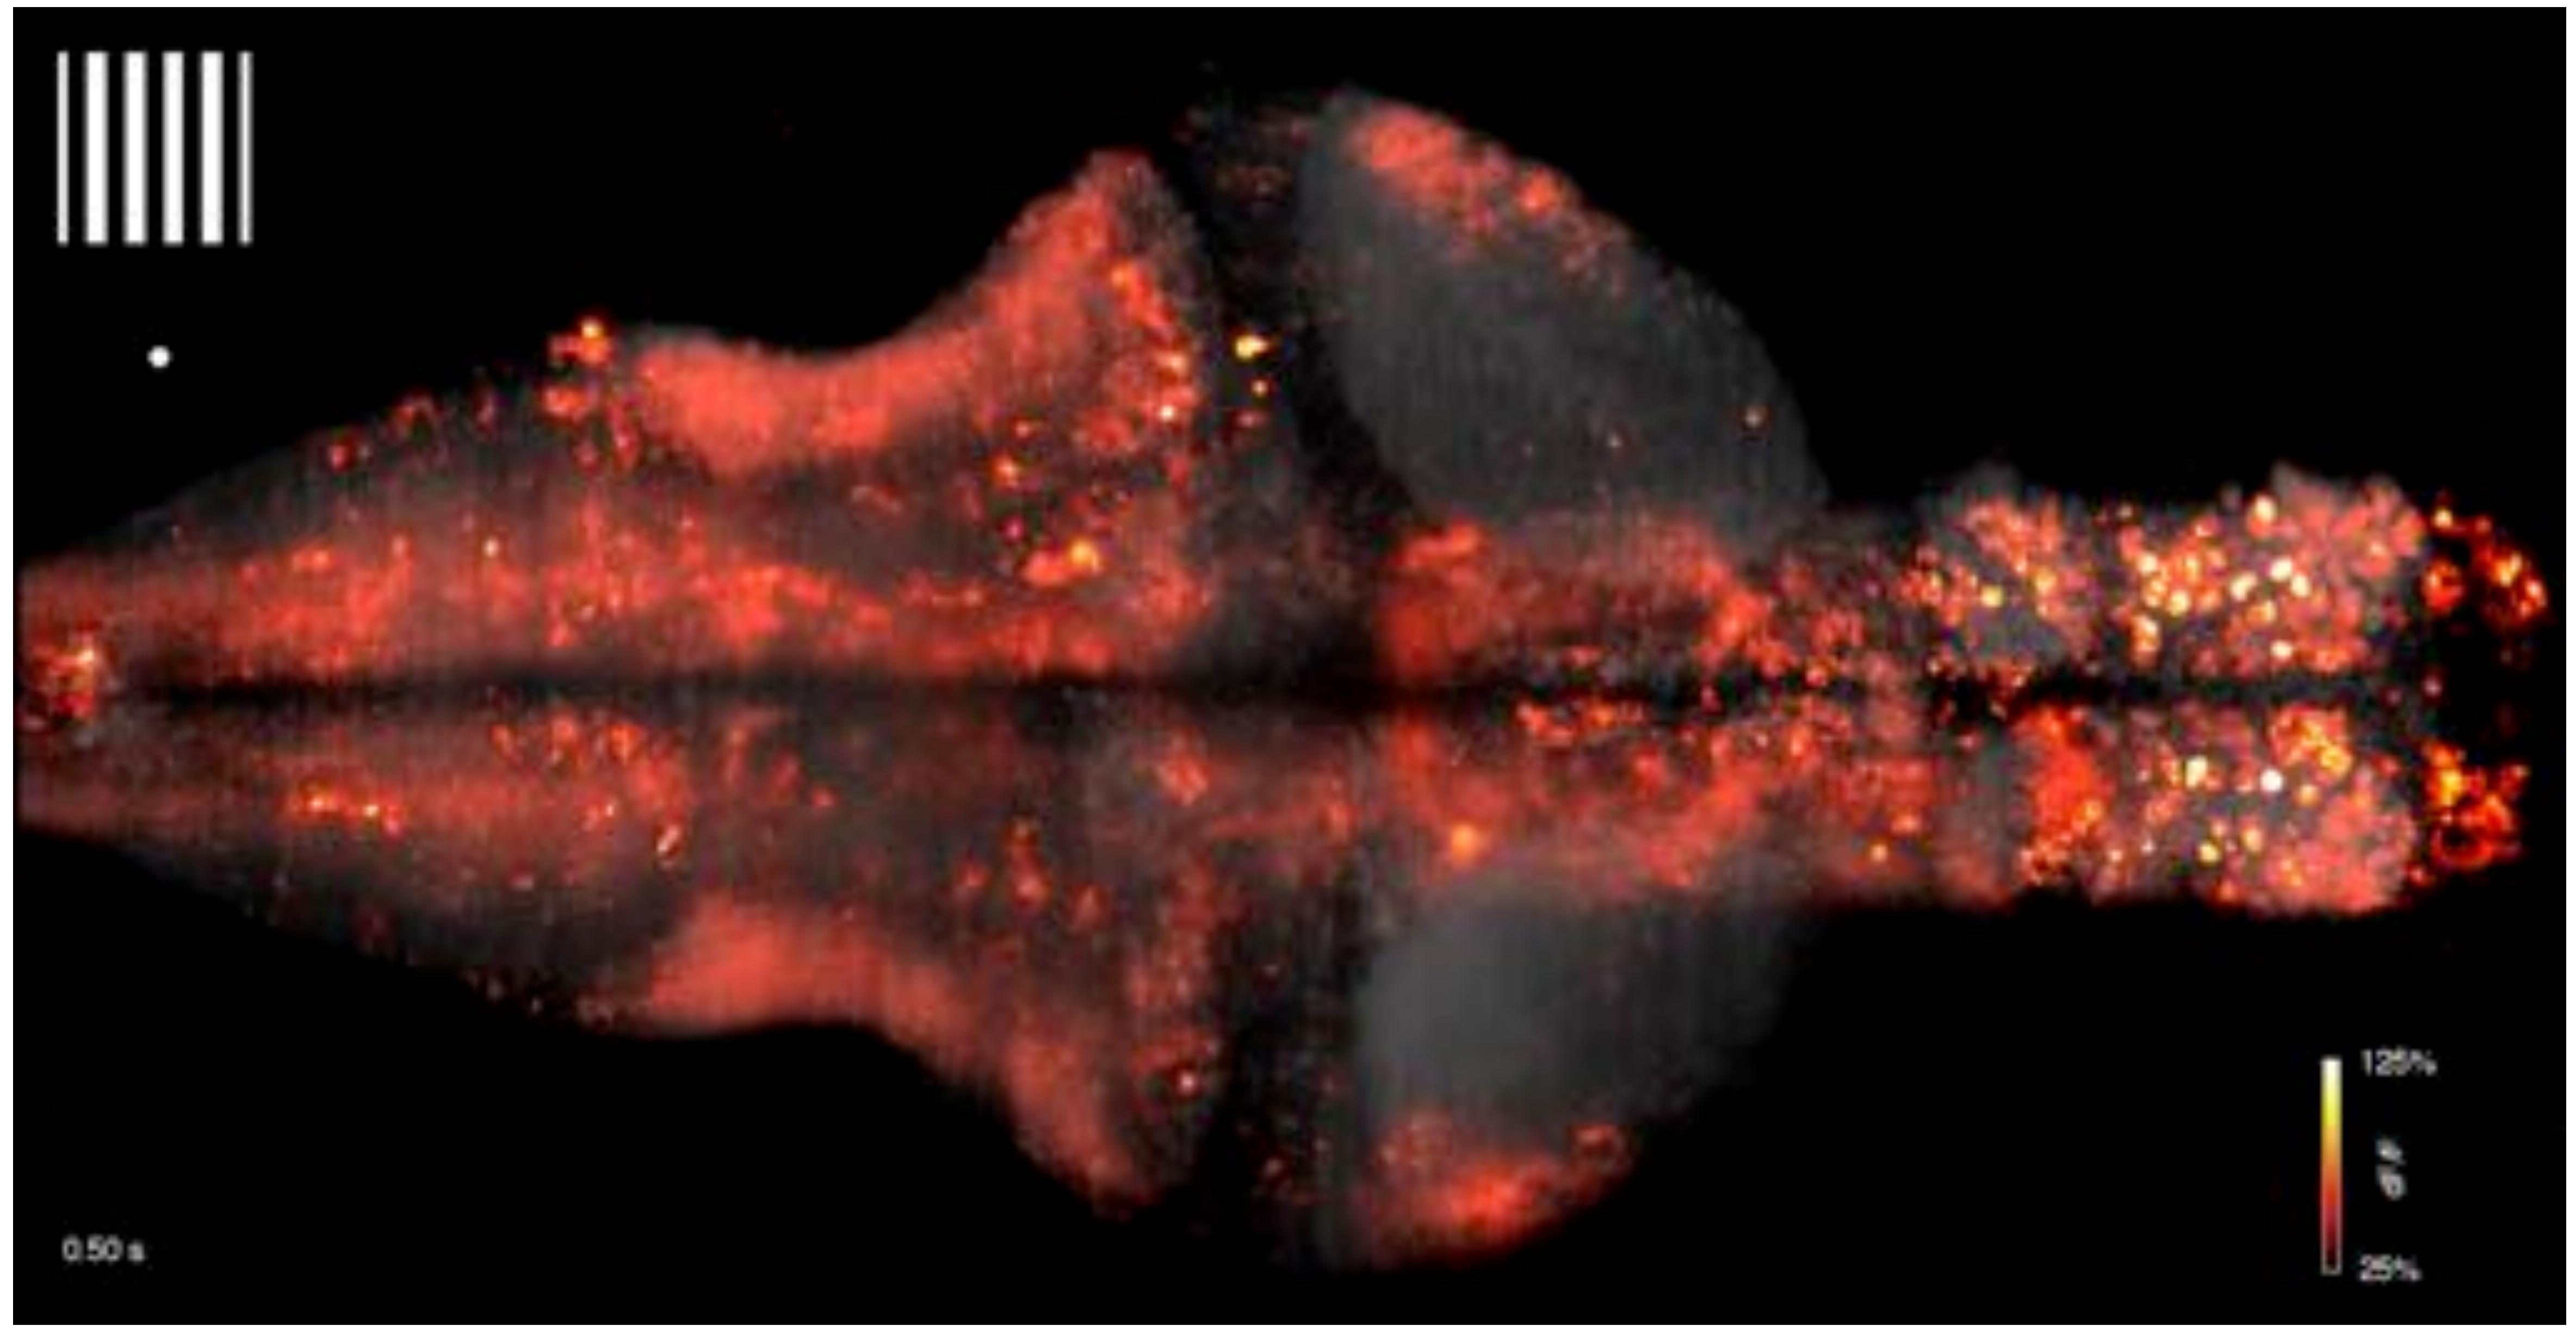


two-photon
100-1000

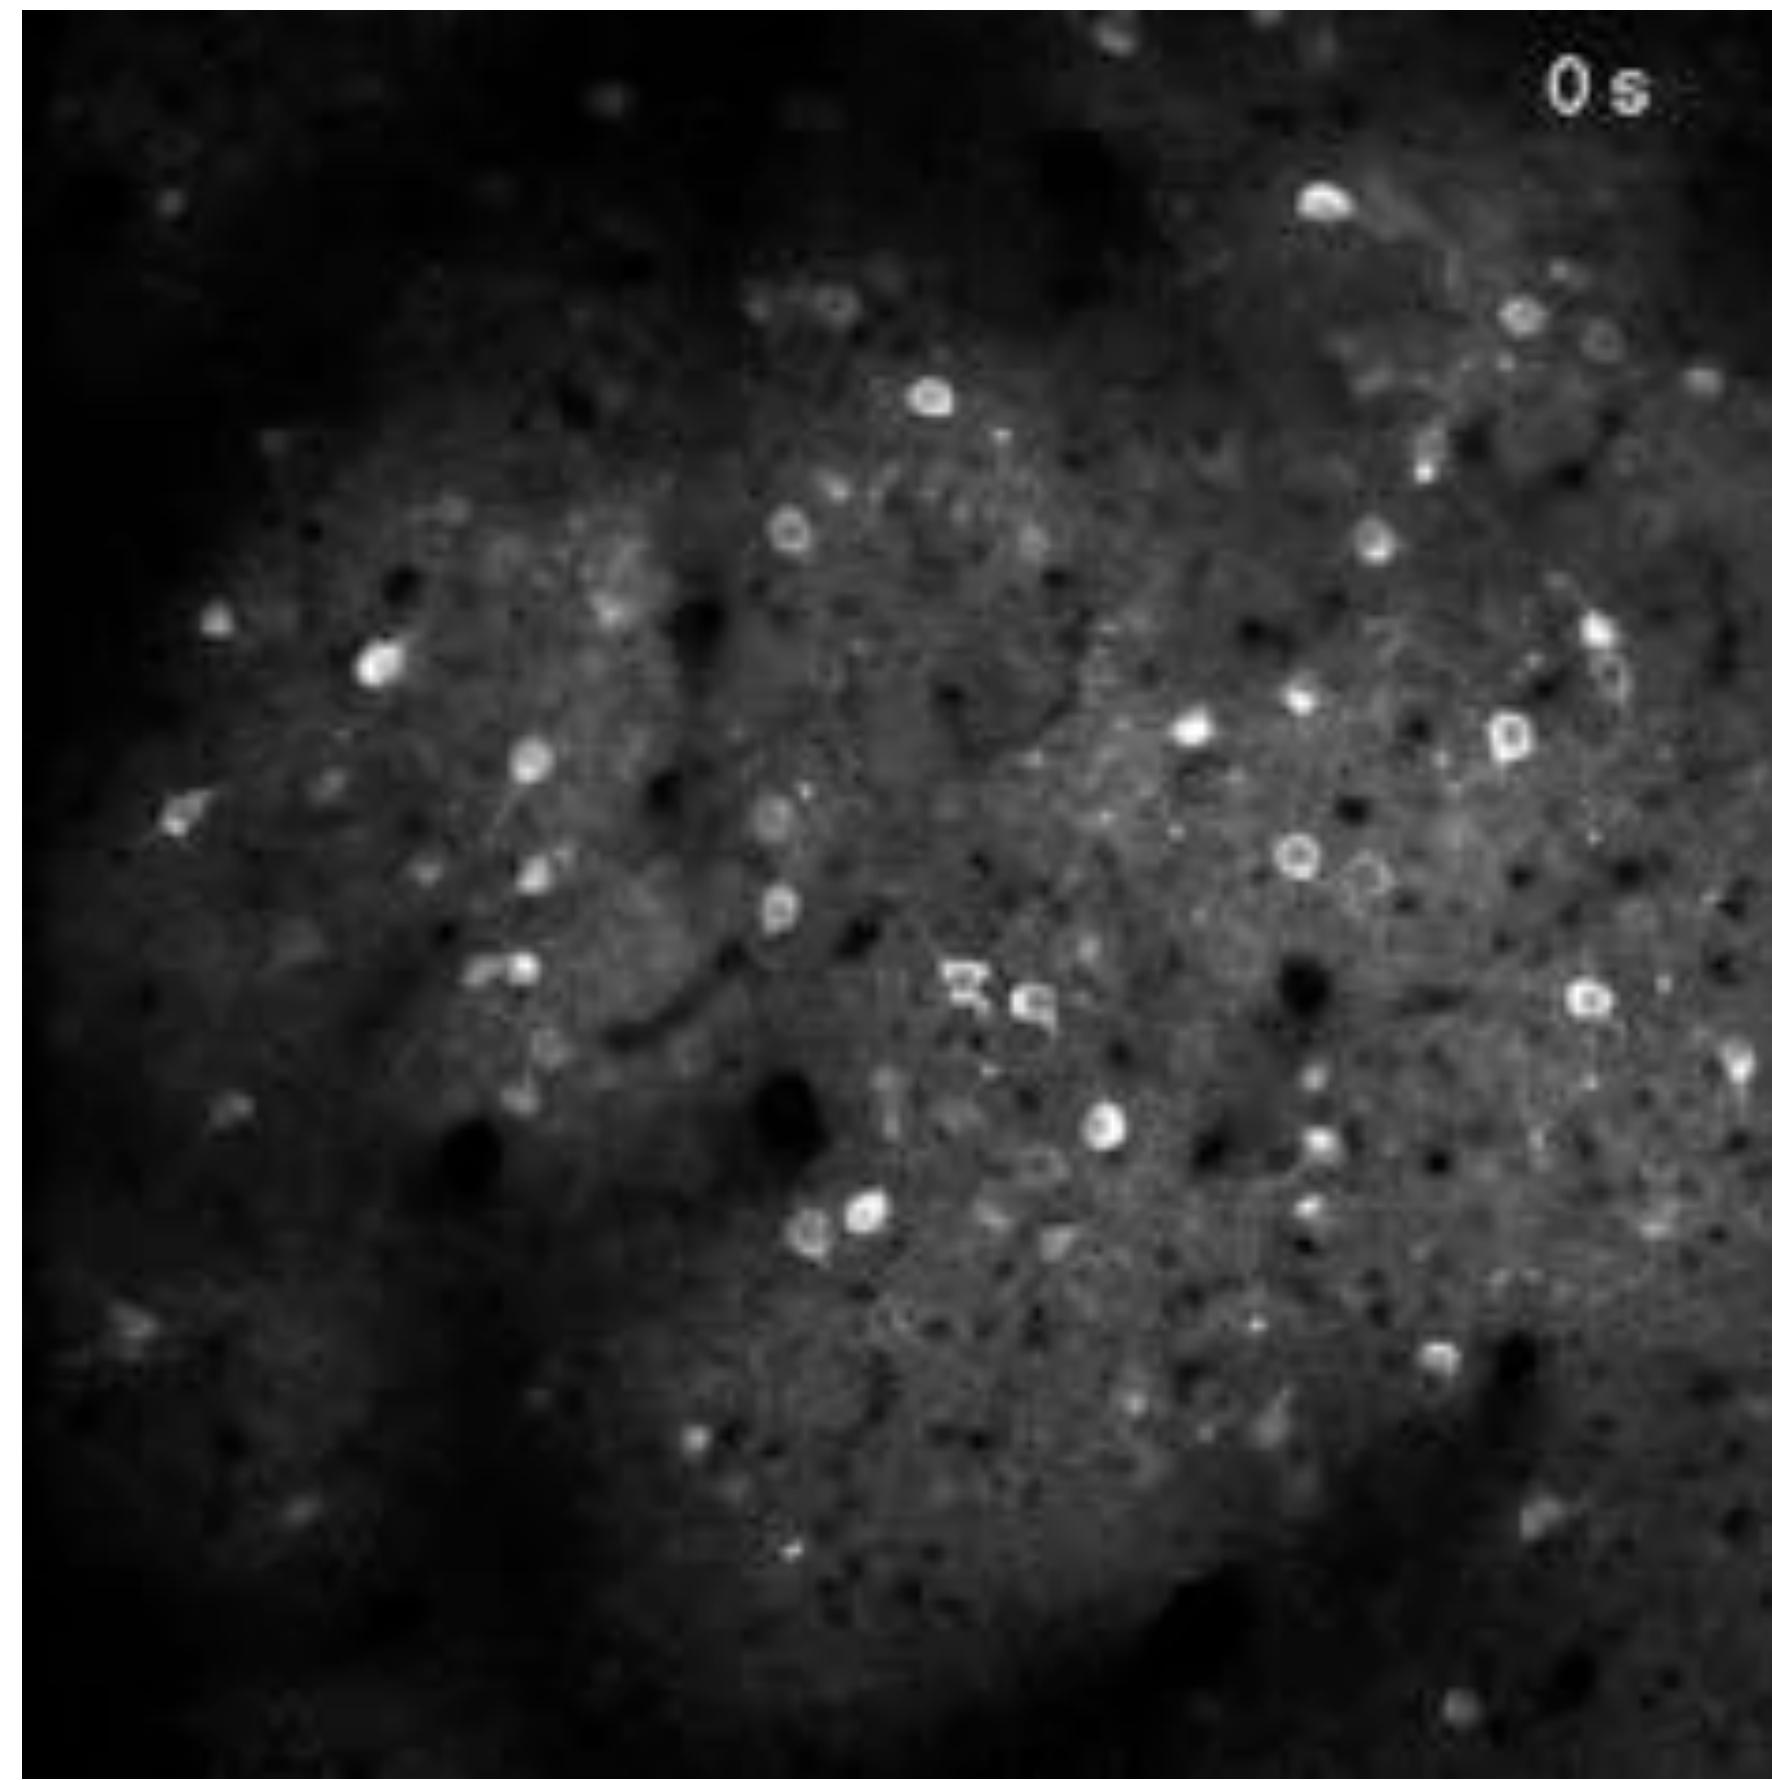


light-sheet
100000



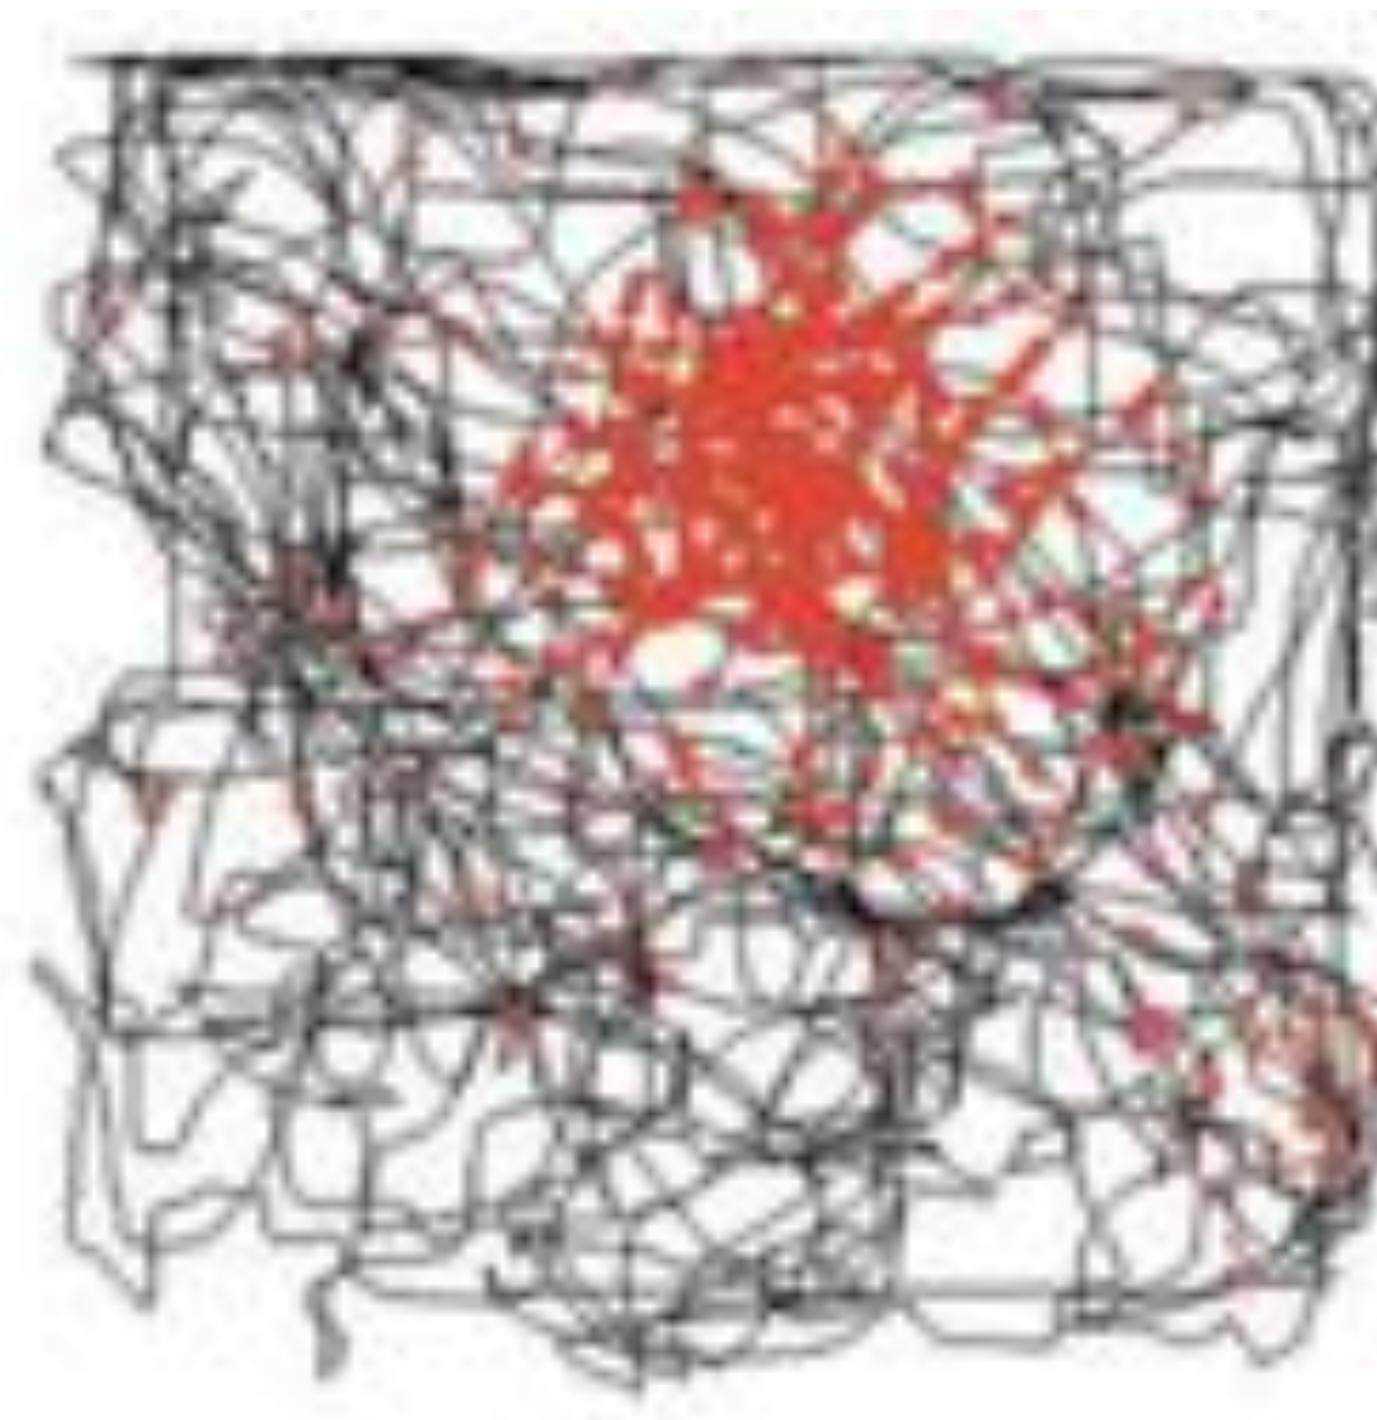


Vladimirov, et al., 2014

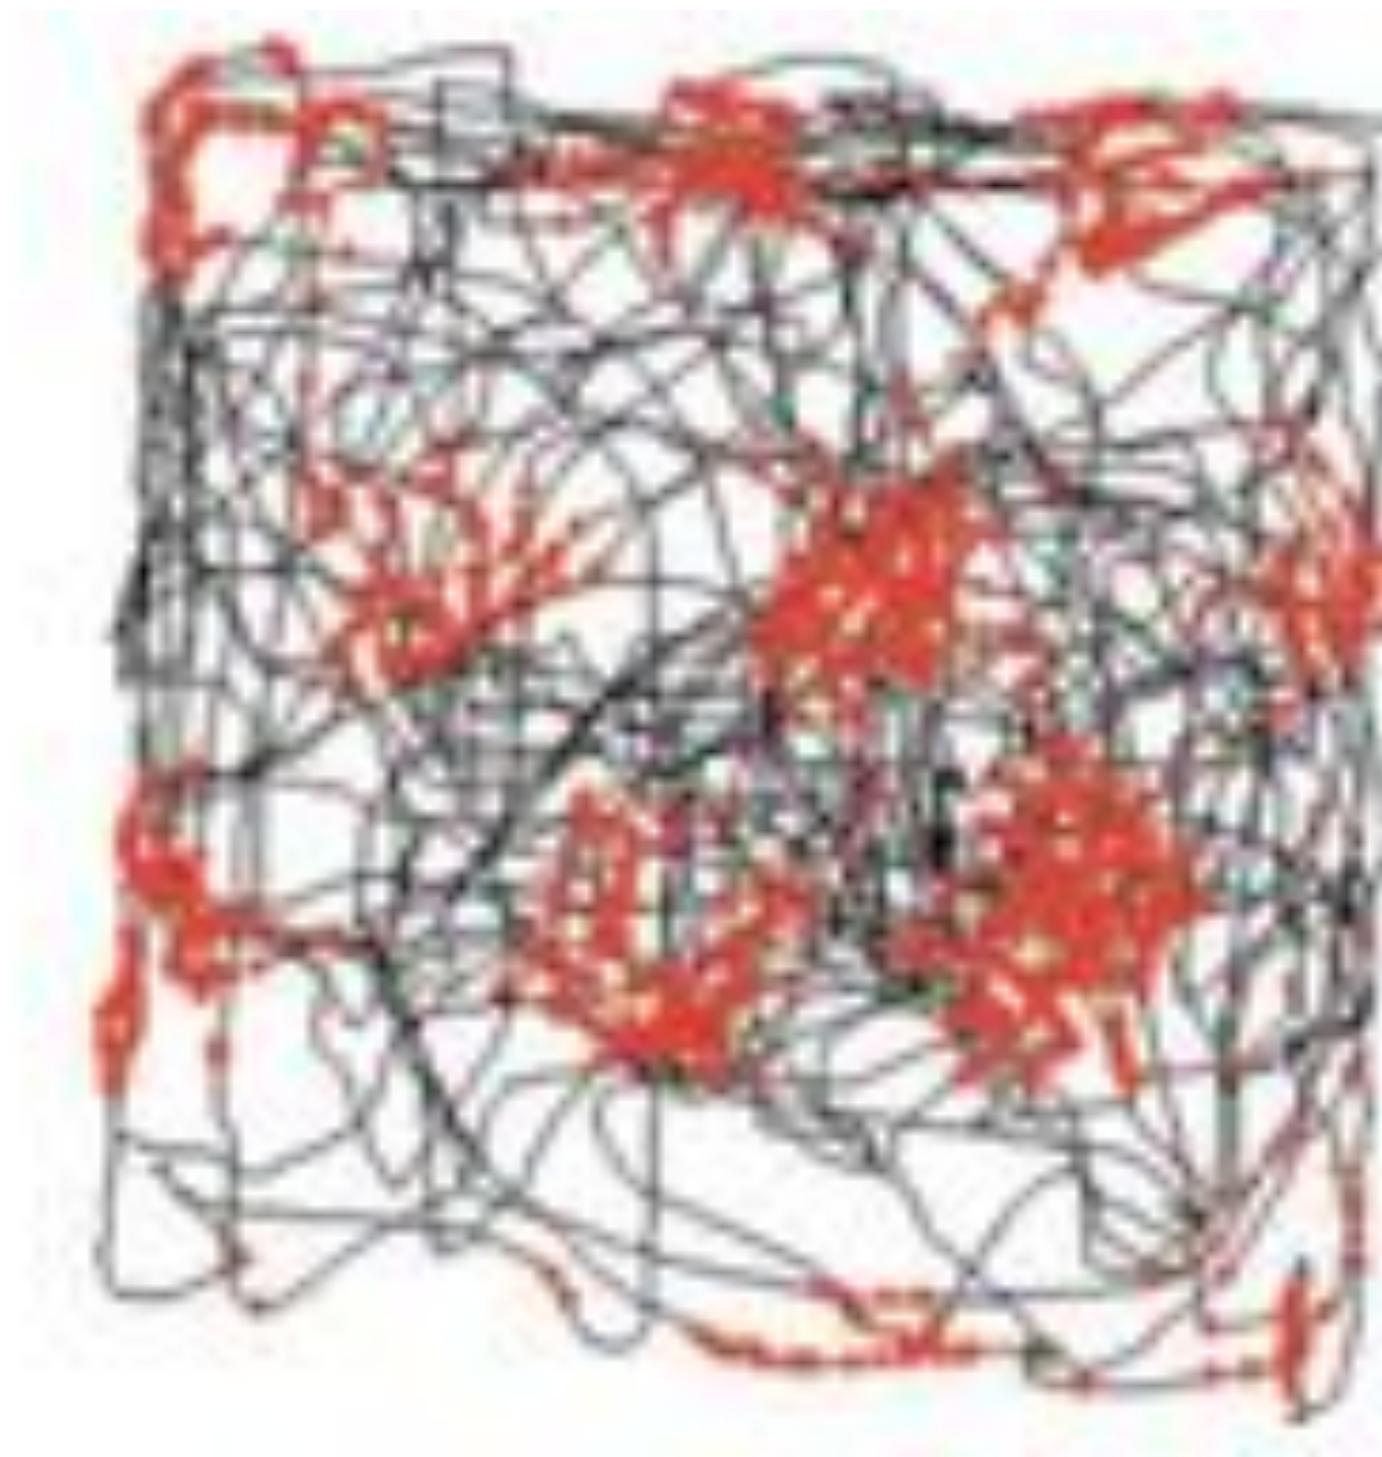


Sofroniew, et al., 2014

relating neuronal responses to properties of an animal and its environment



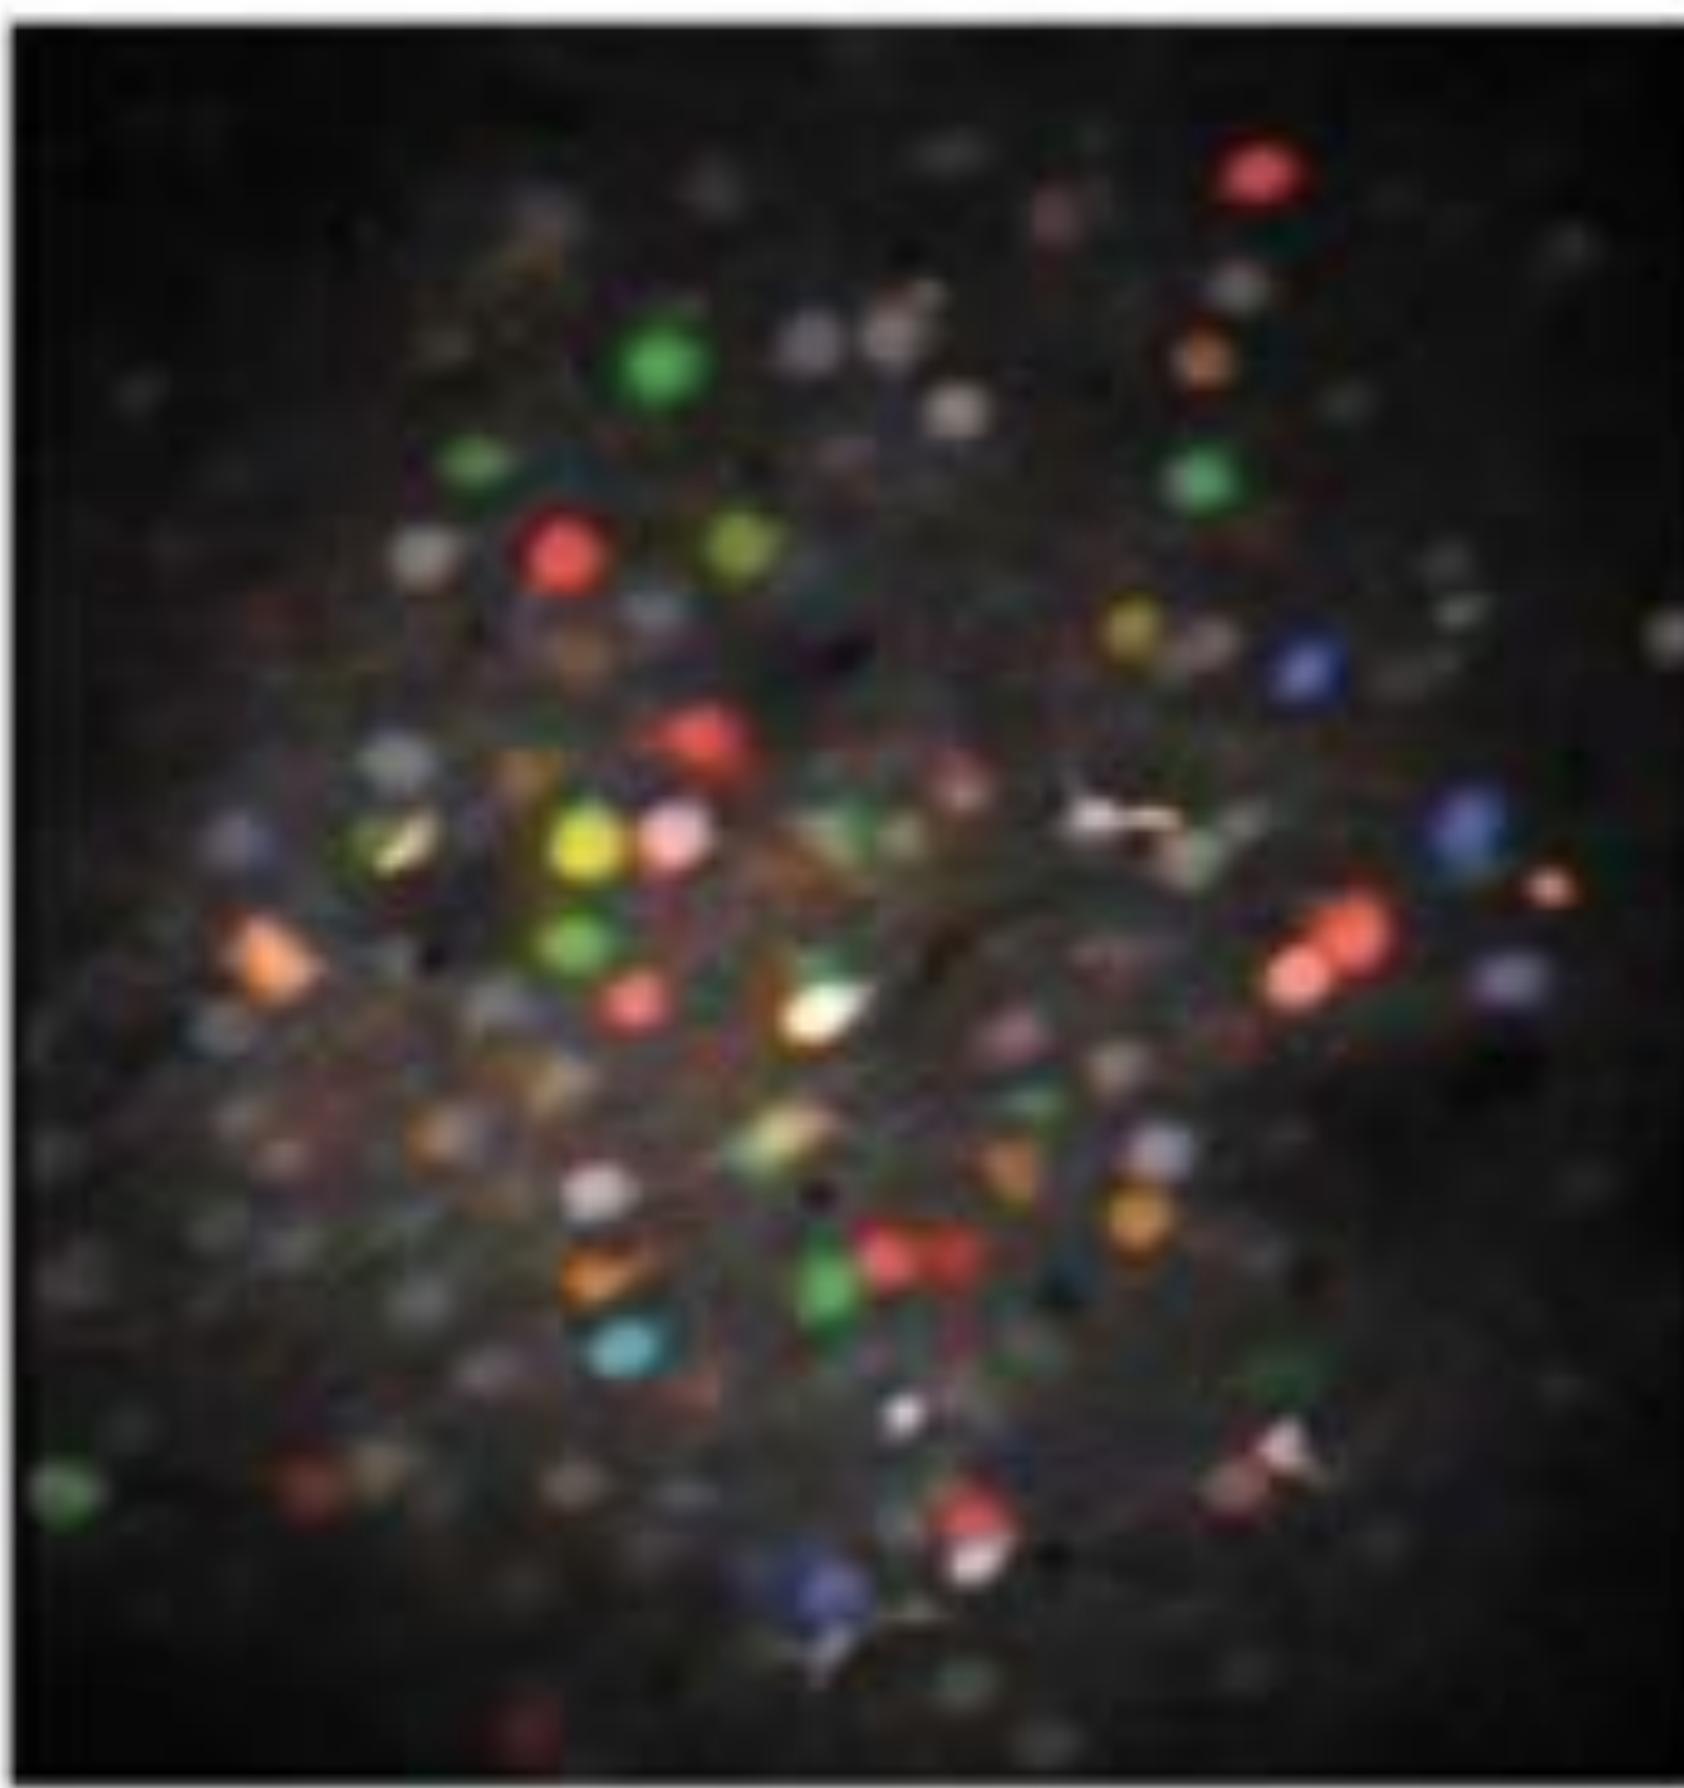
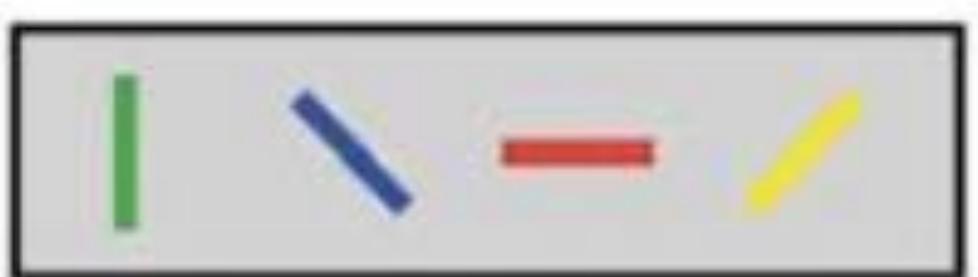
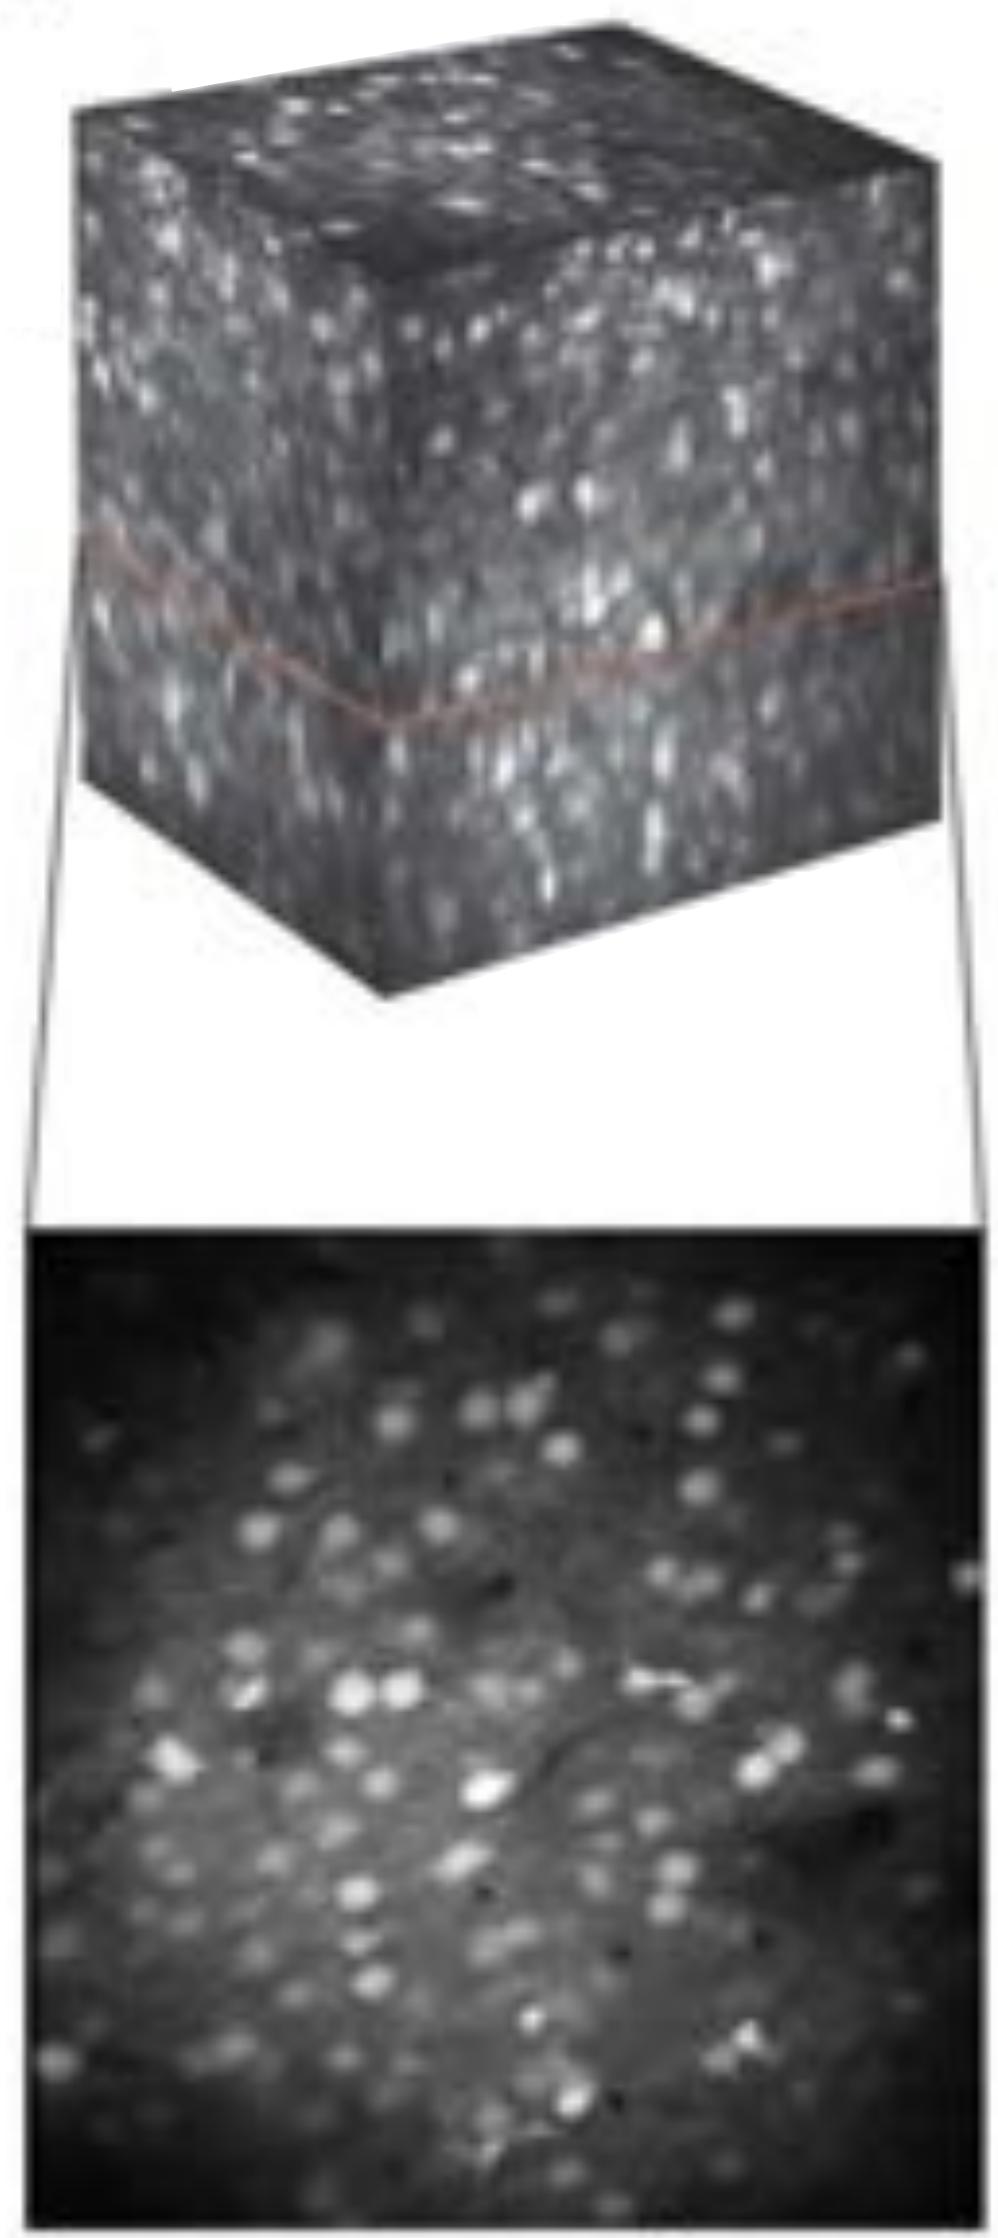
“place cell”



“grid cell”

- position of a mouse in maze

Moser et al., 2008



fine-scale
sensory
tuning

Hubel & Weisel, 1959
Ohki et al., 2006

Mouse, somatosensory cortex
~1,000 neurons

Larval zebrafish, whole-brain
~100,000 neurons

* Entire mouse brain
~80,000,000 neurons

* hypothetical

| 0.1 TB / experiment

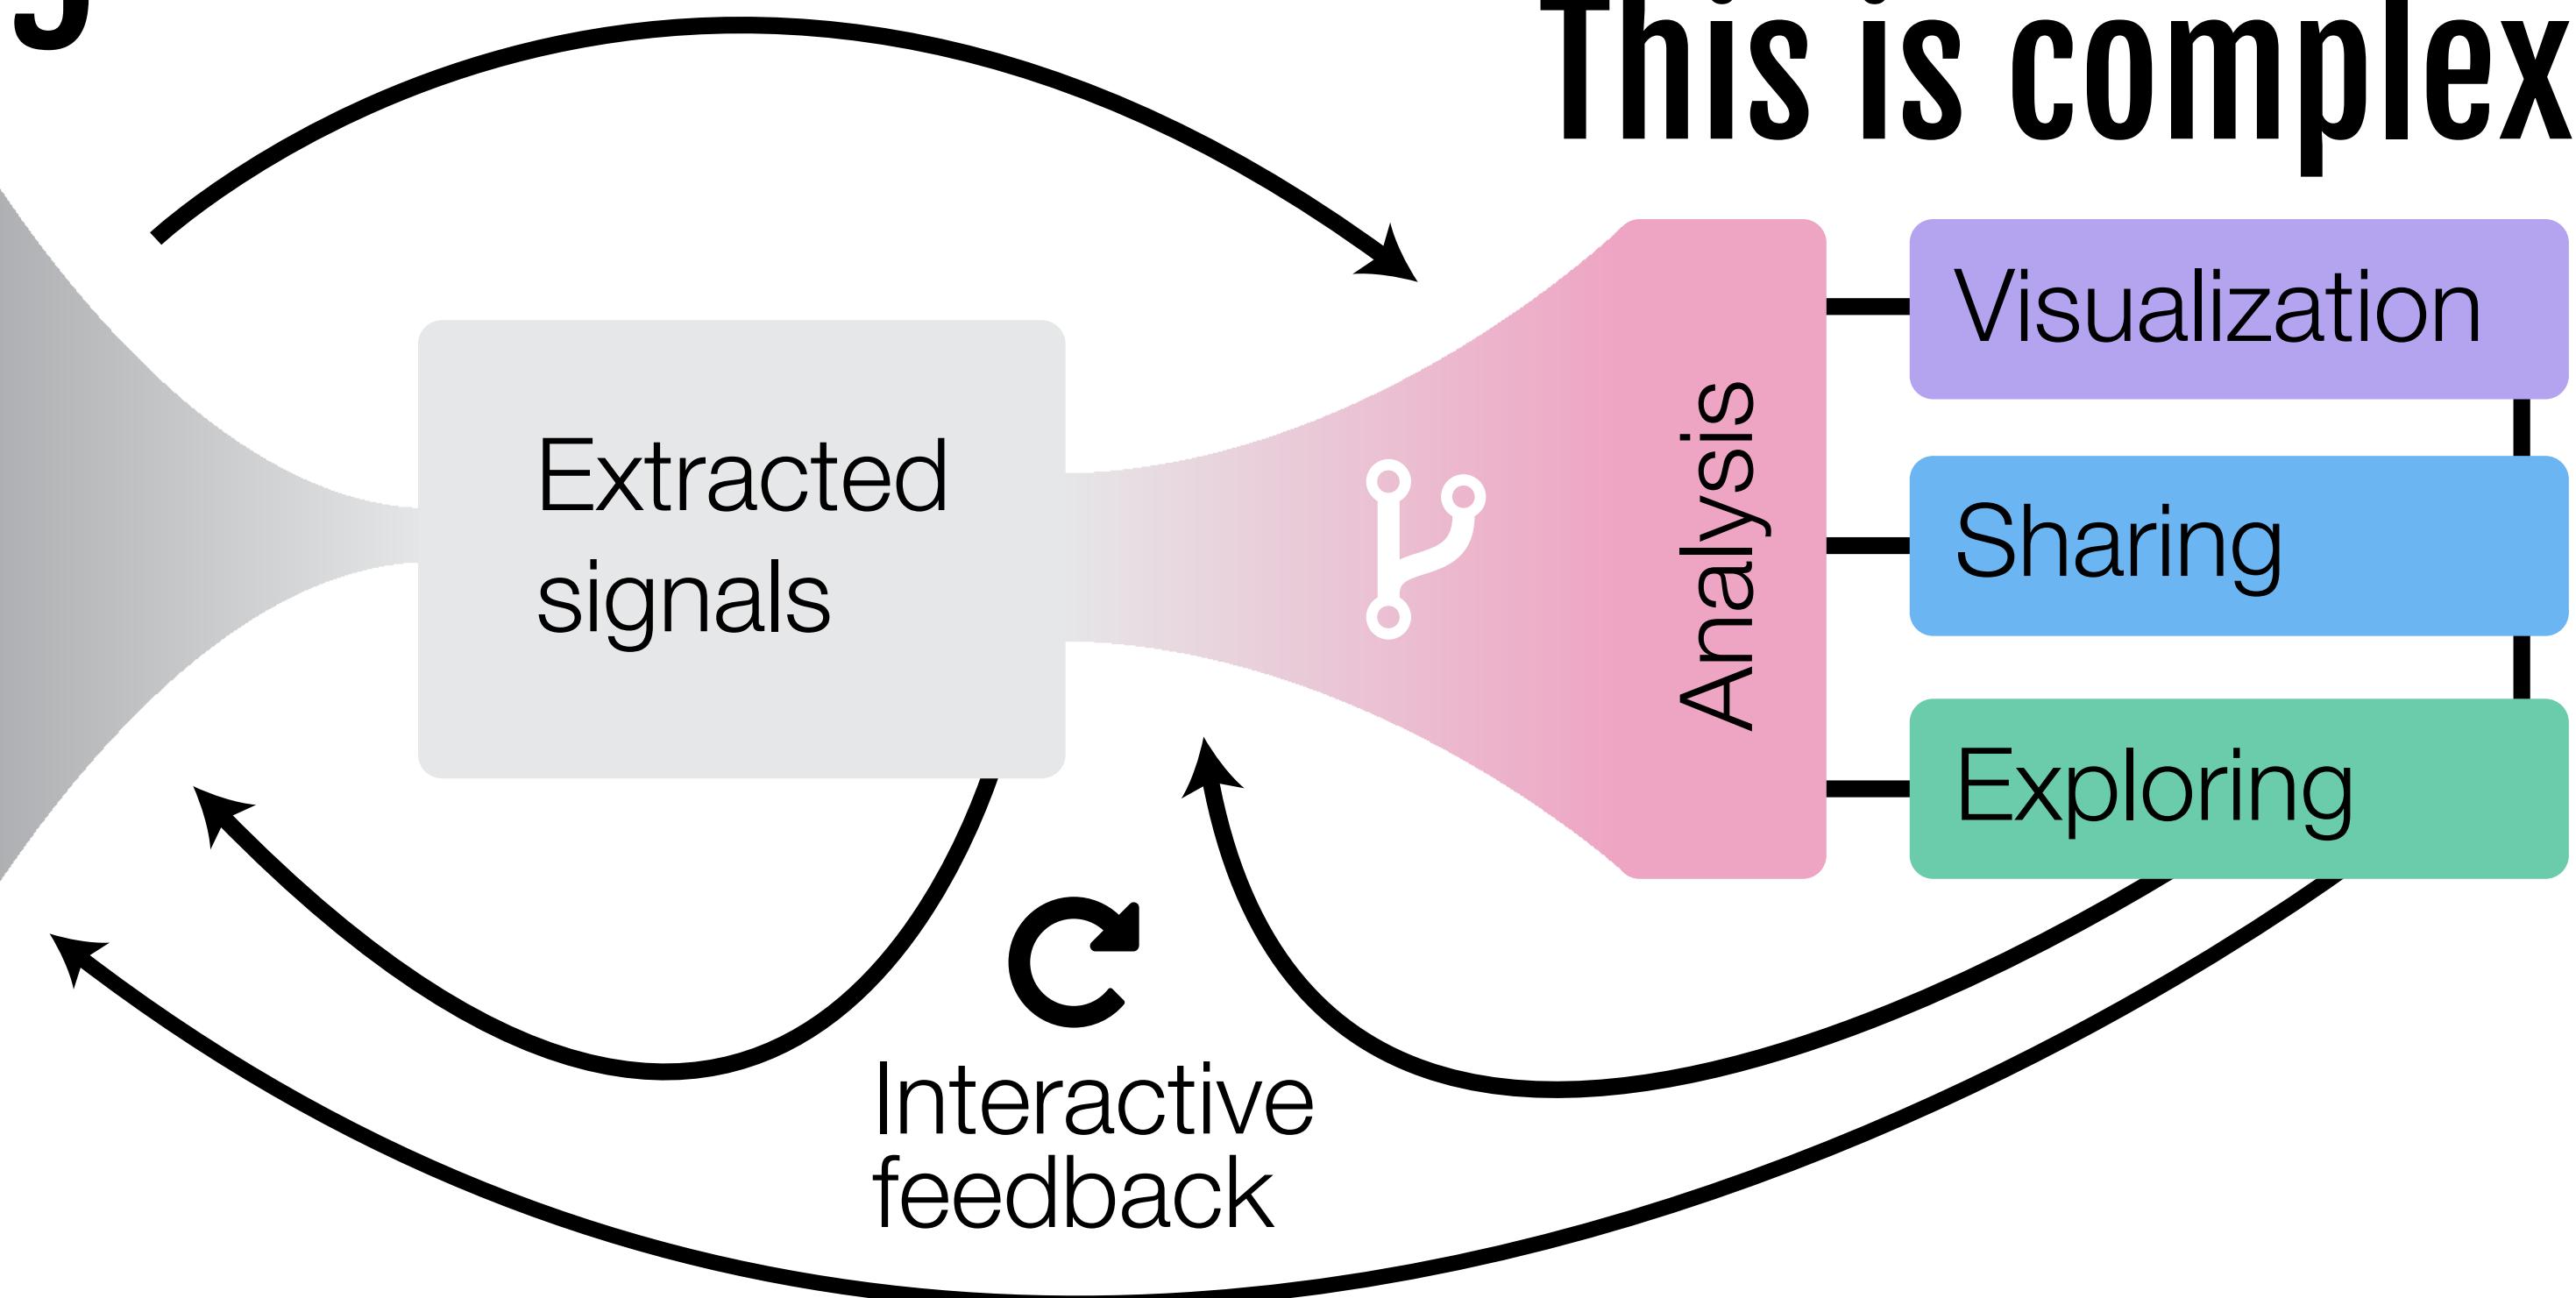
■ 1 TB

■ // ■ >100 TB

Exploratory Data Analysis

This is really big

Raw
data



This is complex

Supervised methods

$$\mathbf{y} = f(\mathbf{X})$$

predict
our data

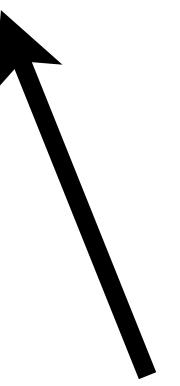
as a function
of other data

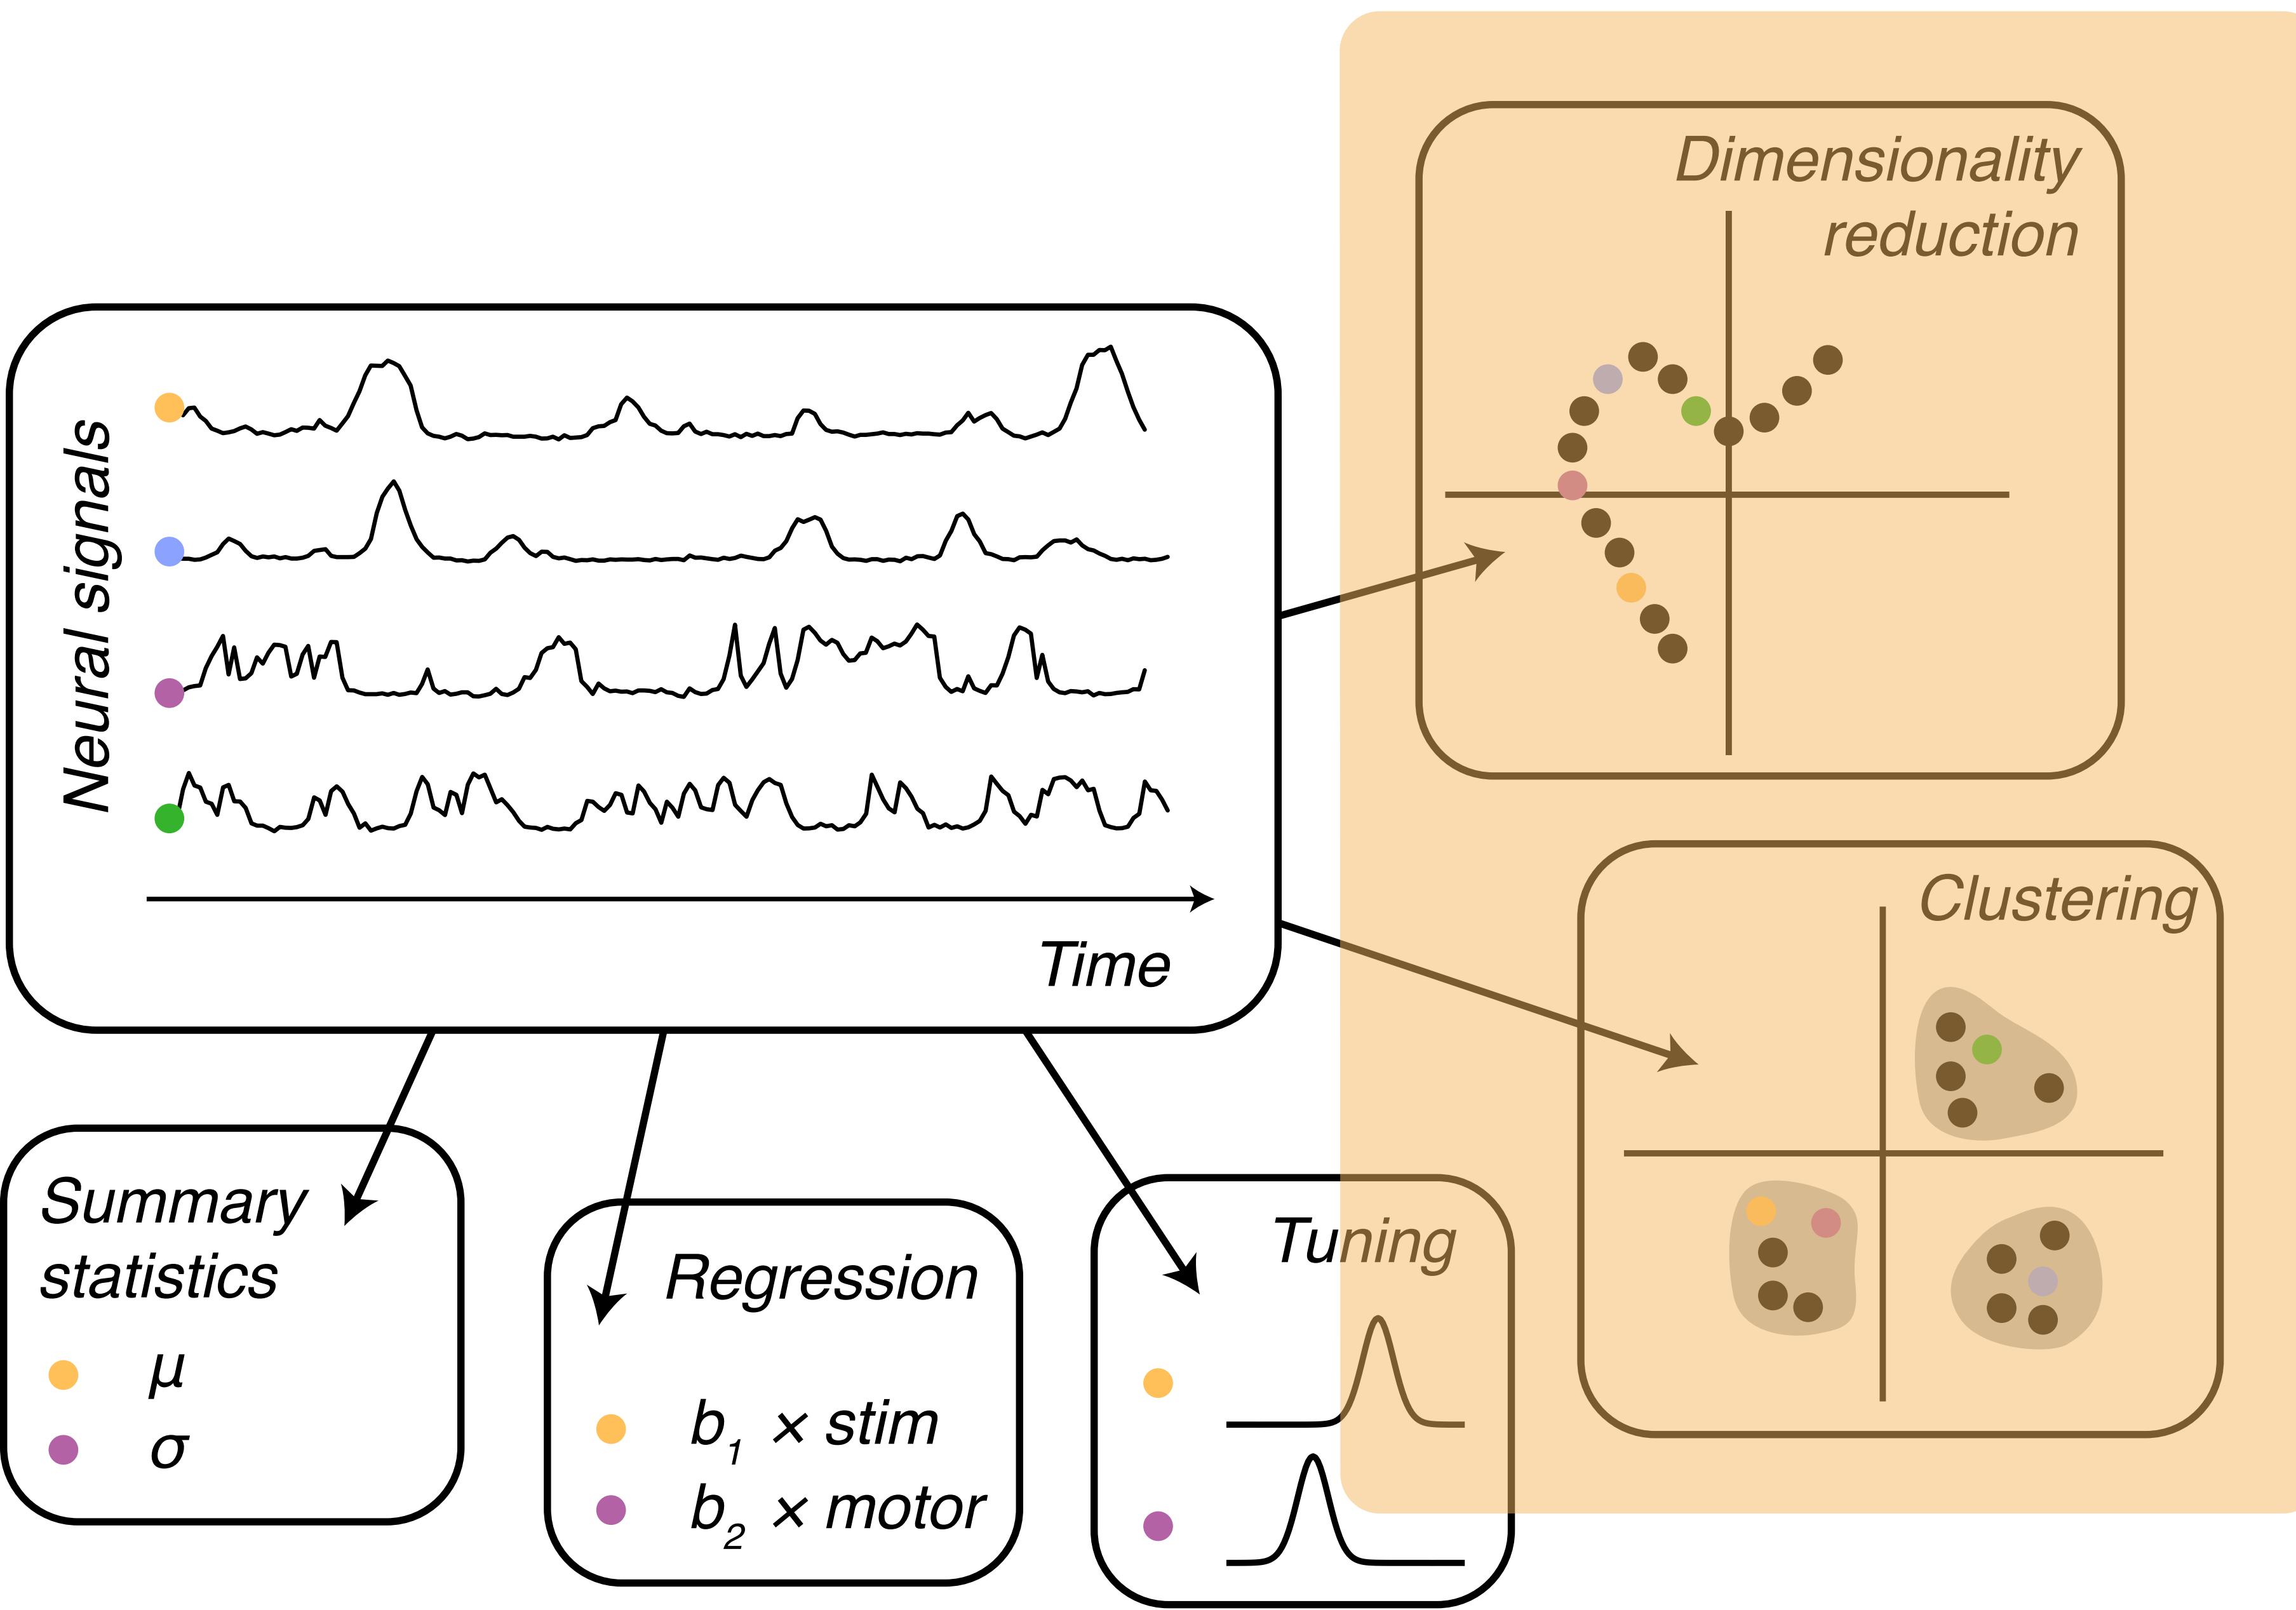


Unsupervised methods

$$f(\mathbf{X})$$

find structure
in the data
on its own





Clustering for preprocessing

- Raw data is complex and high-dimensional
- Clustering finds collections of inputs that are similar to one another
- These groups of clusters may be the more meaningful “unit” of measurement

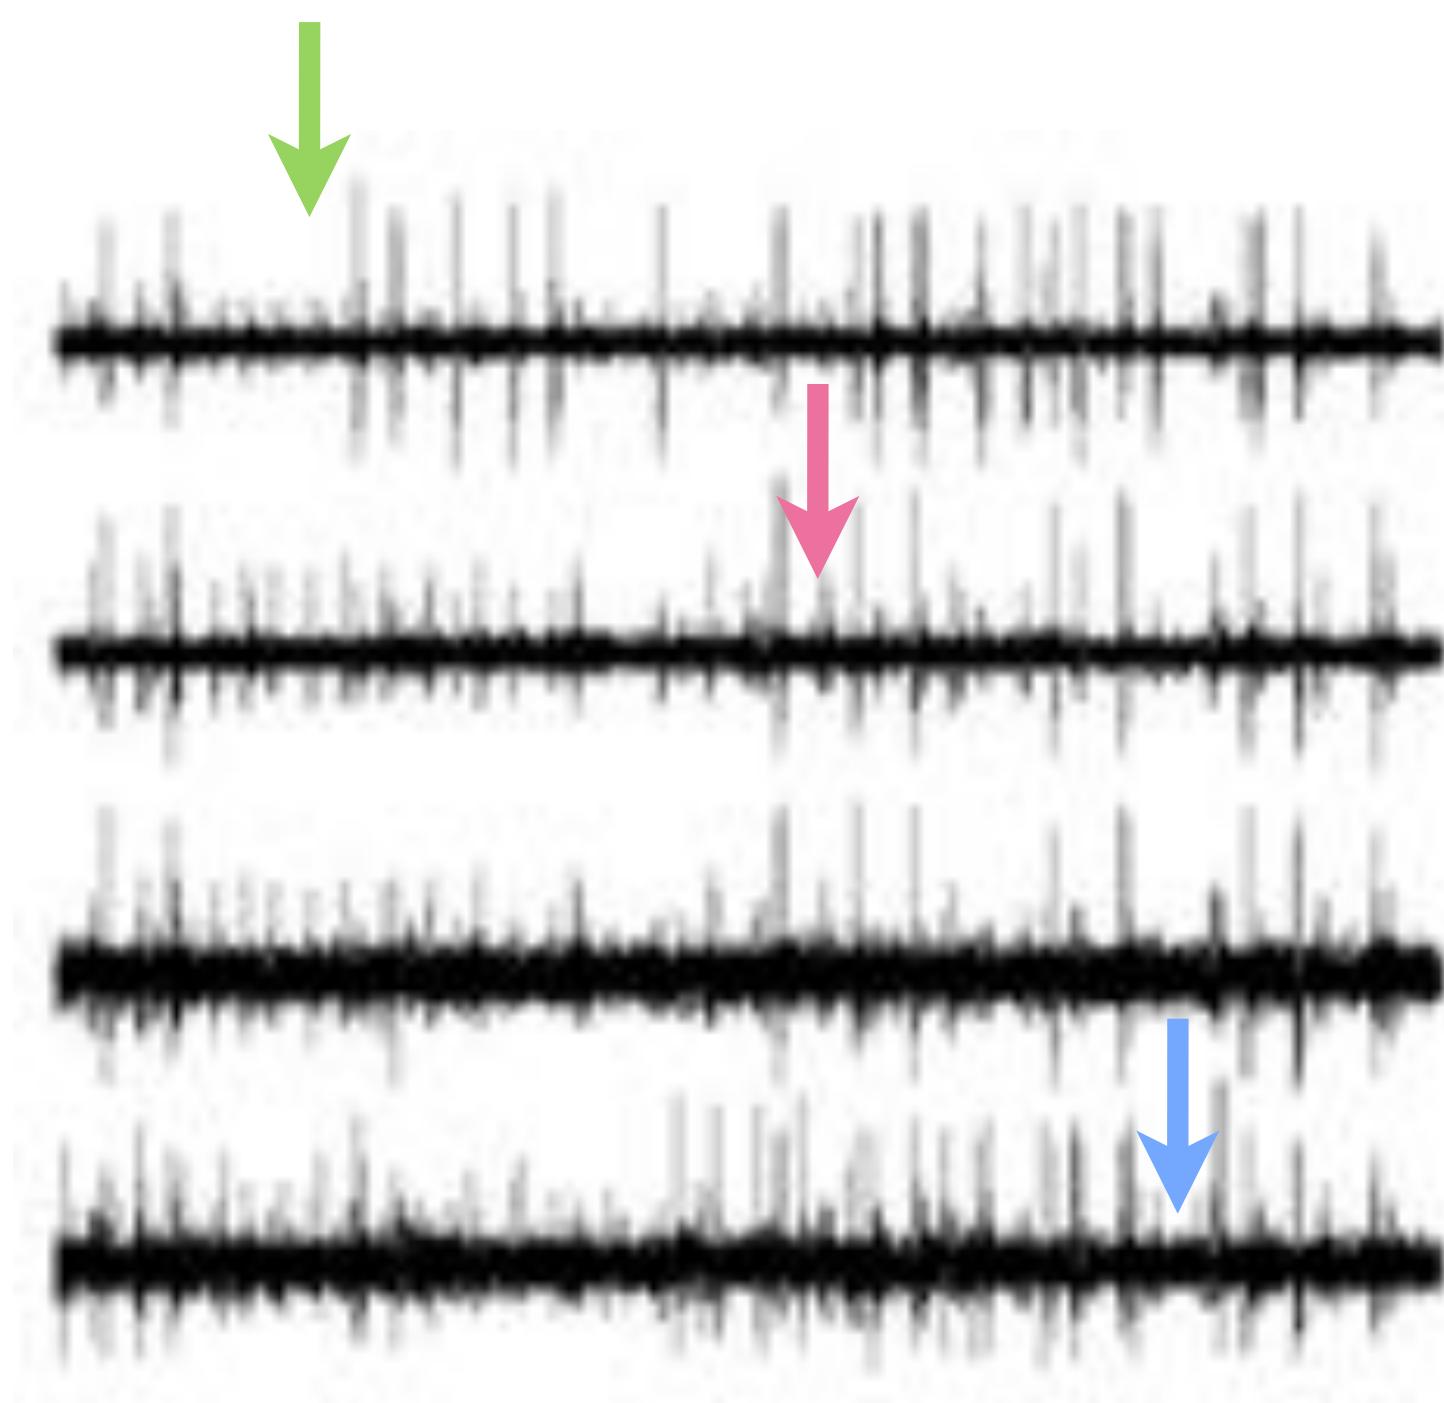
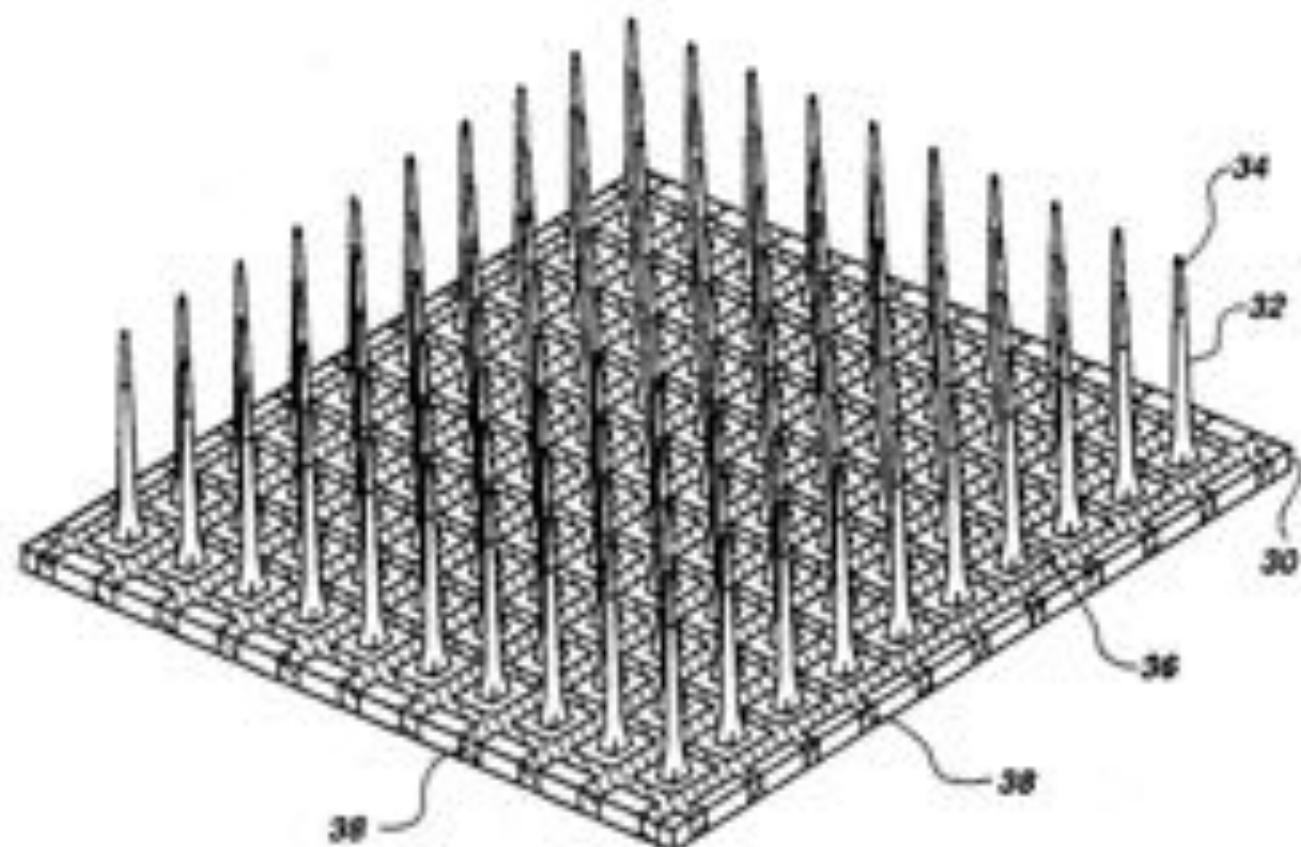
Raw data



Clustered data



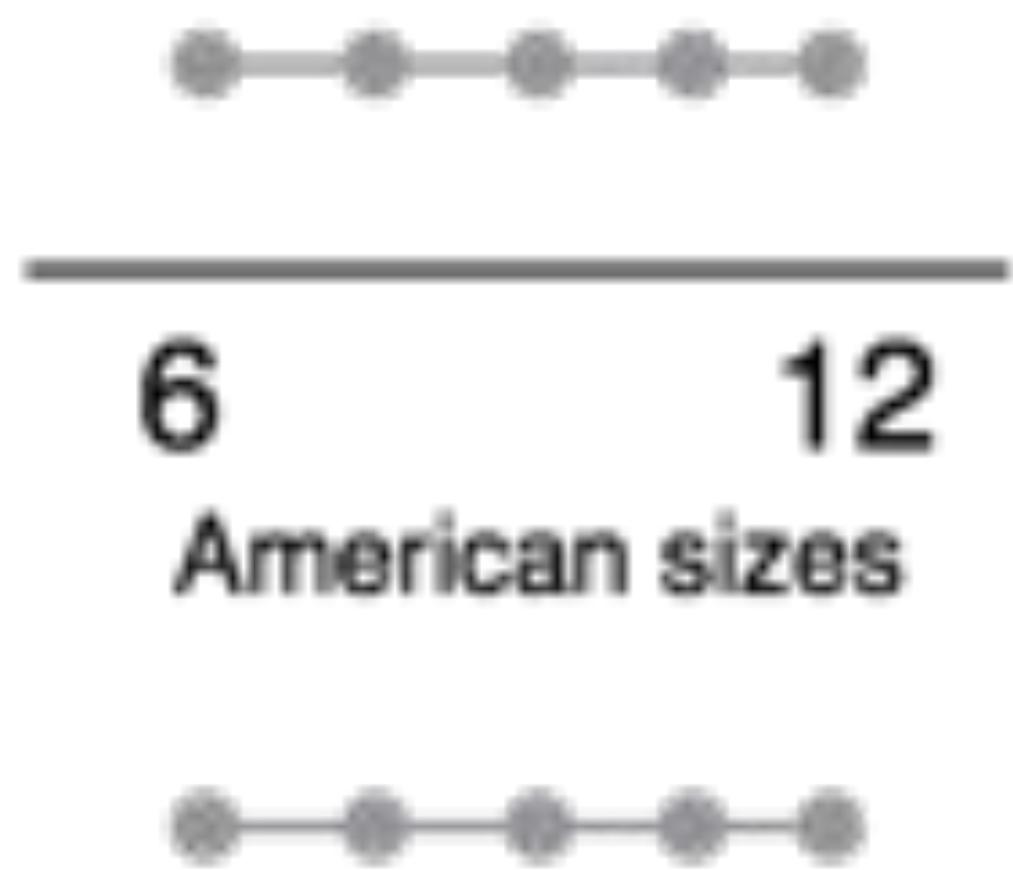
Clustering to find waveforms associated with individual neurons based on their traces across multiple electrodes



Dimensionality reduction for insight

- Raw data is complex and high-dimensional
- Dimensionality reduction describes the data using a simpler, more compact representation
- This representation may make interesting patterns in the data more clear or easier to see

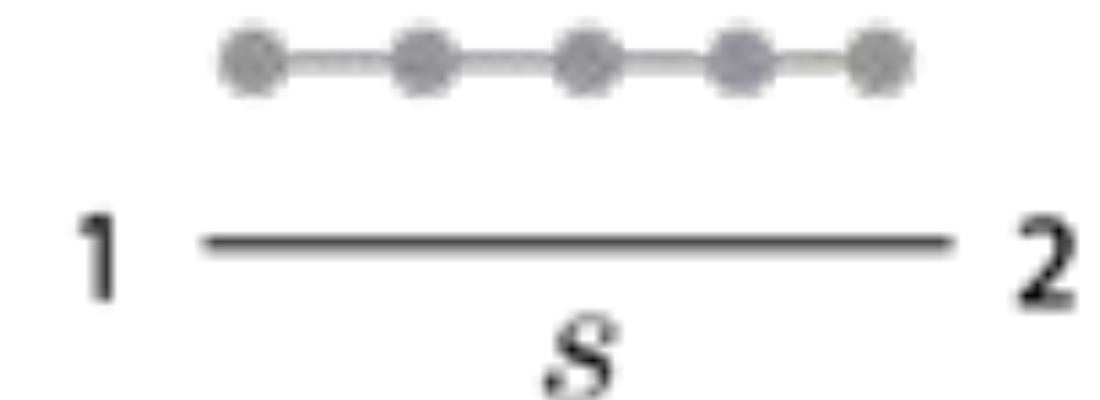
“Shoe store” space



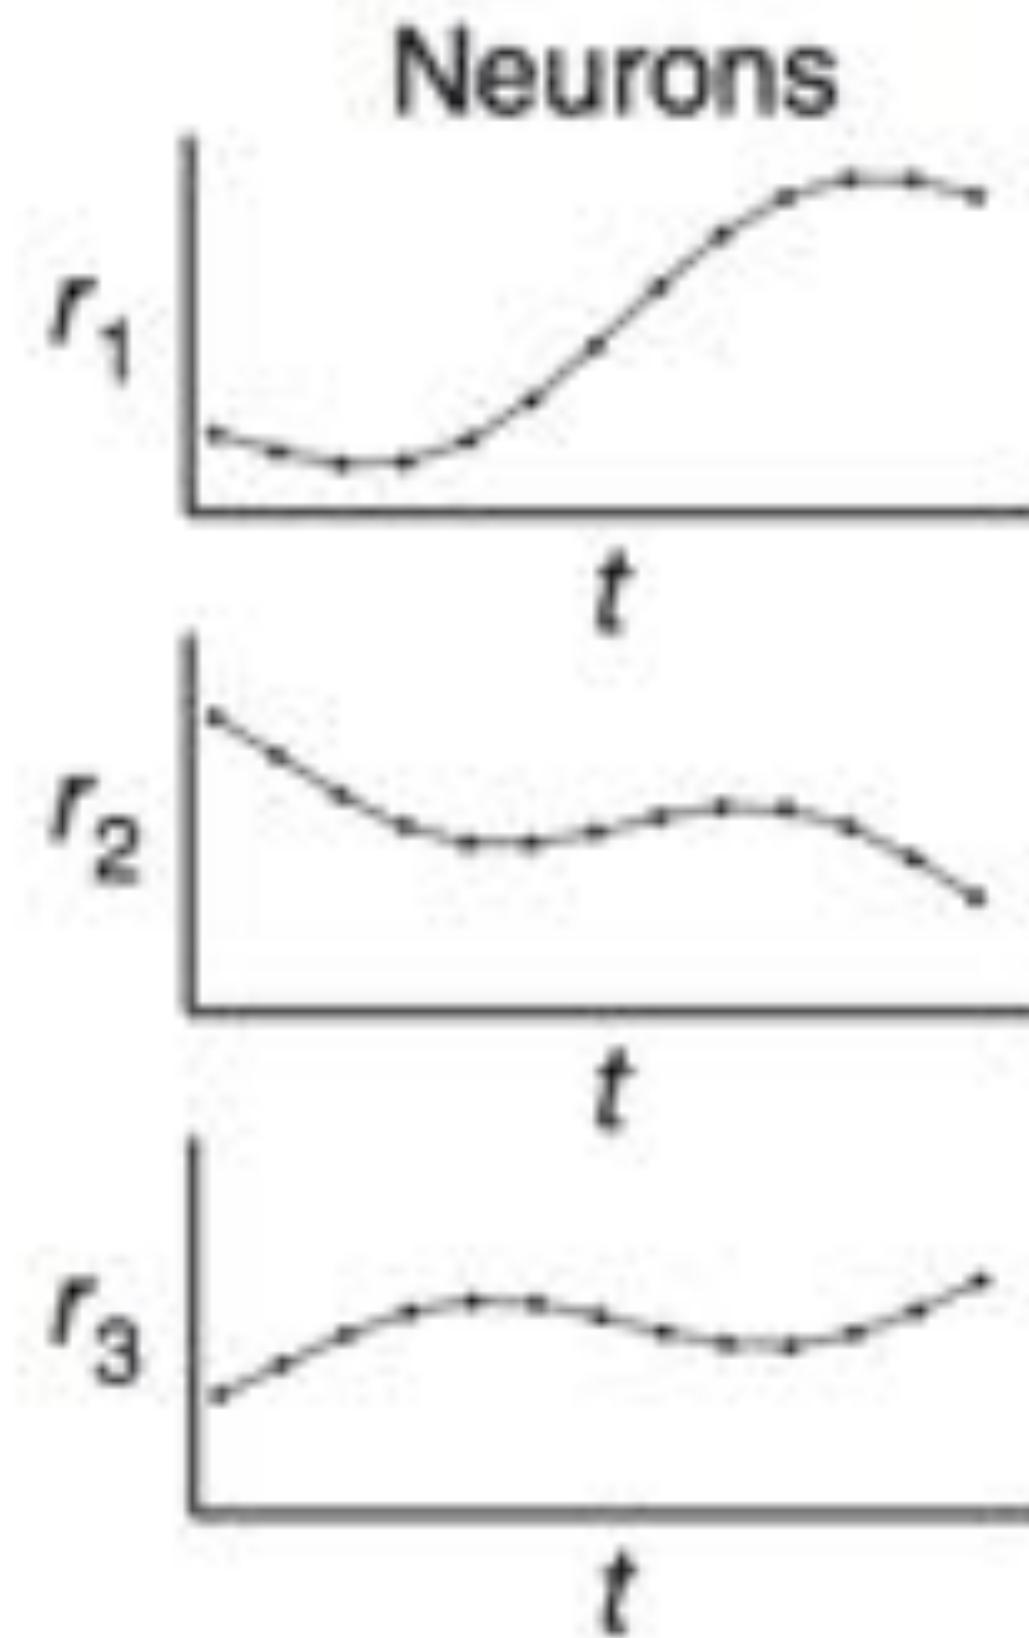
weights

$$= \begin{bmatrix} 6s \\ 6s + 33 \end{bmatrix} * \begin{bmatrix} s \end{bmatrix} \leftrightarrow$$

“Size” space

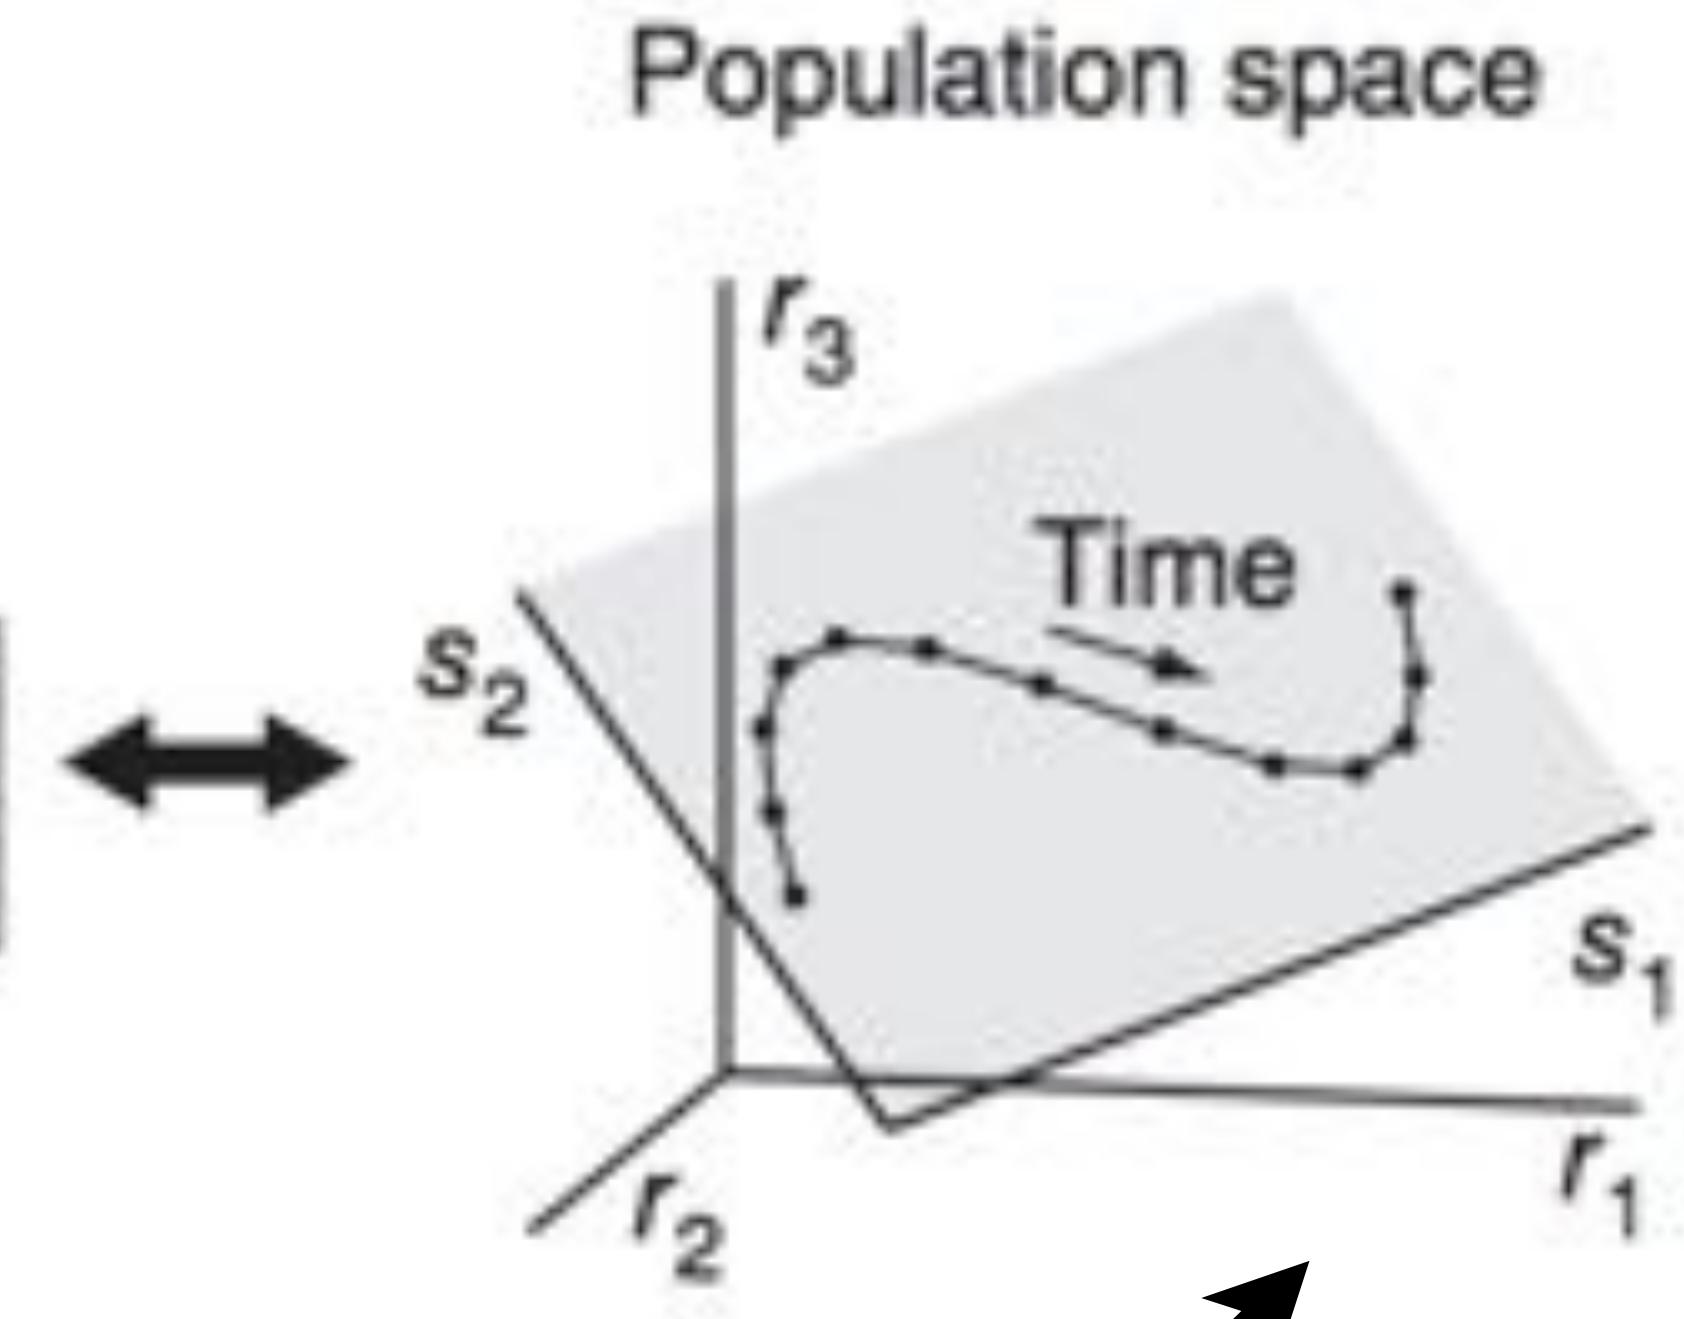


Dimensionality
reduction

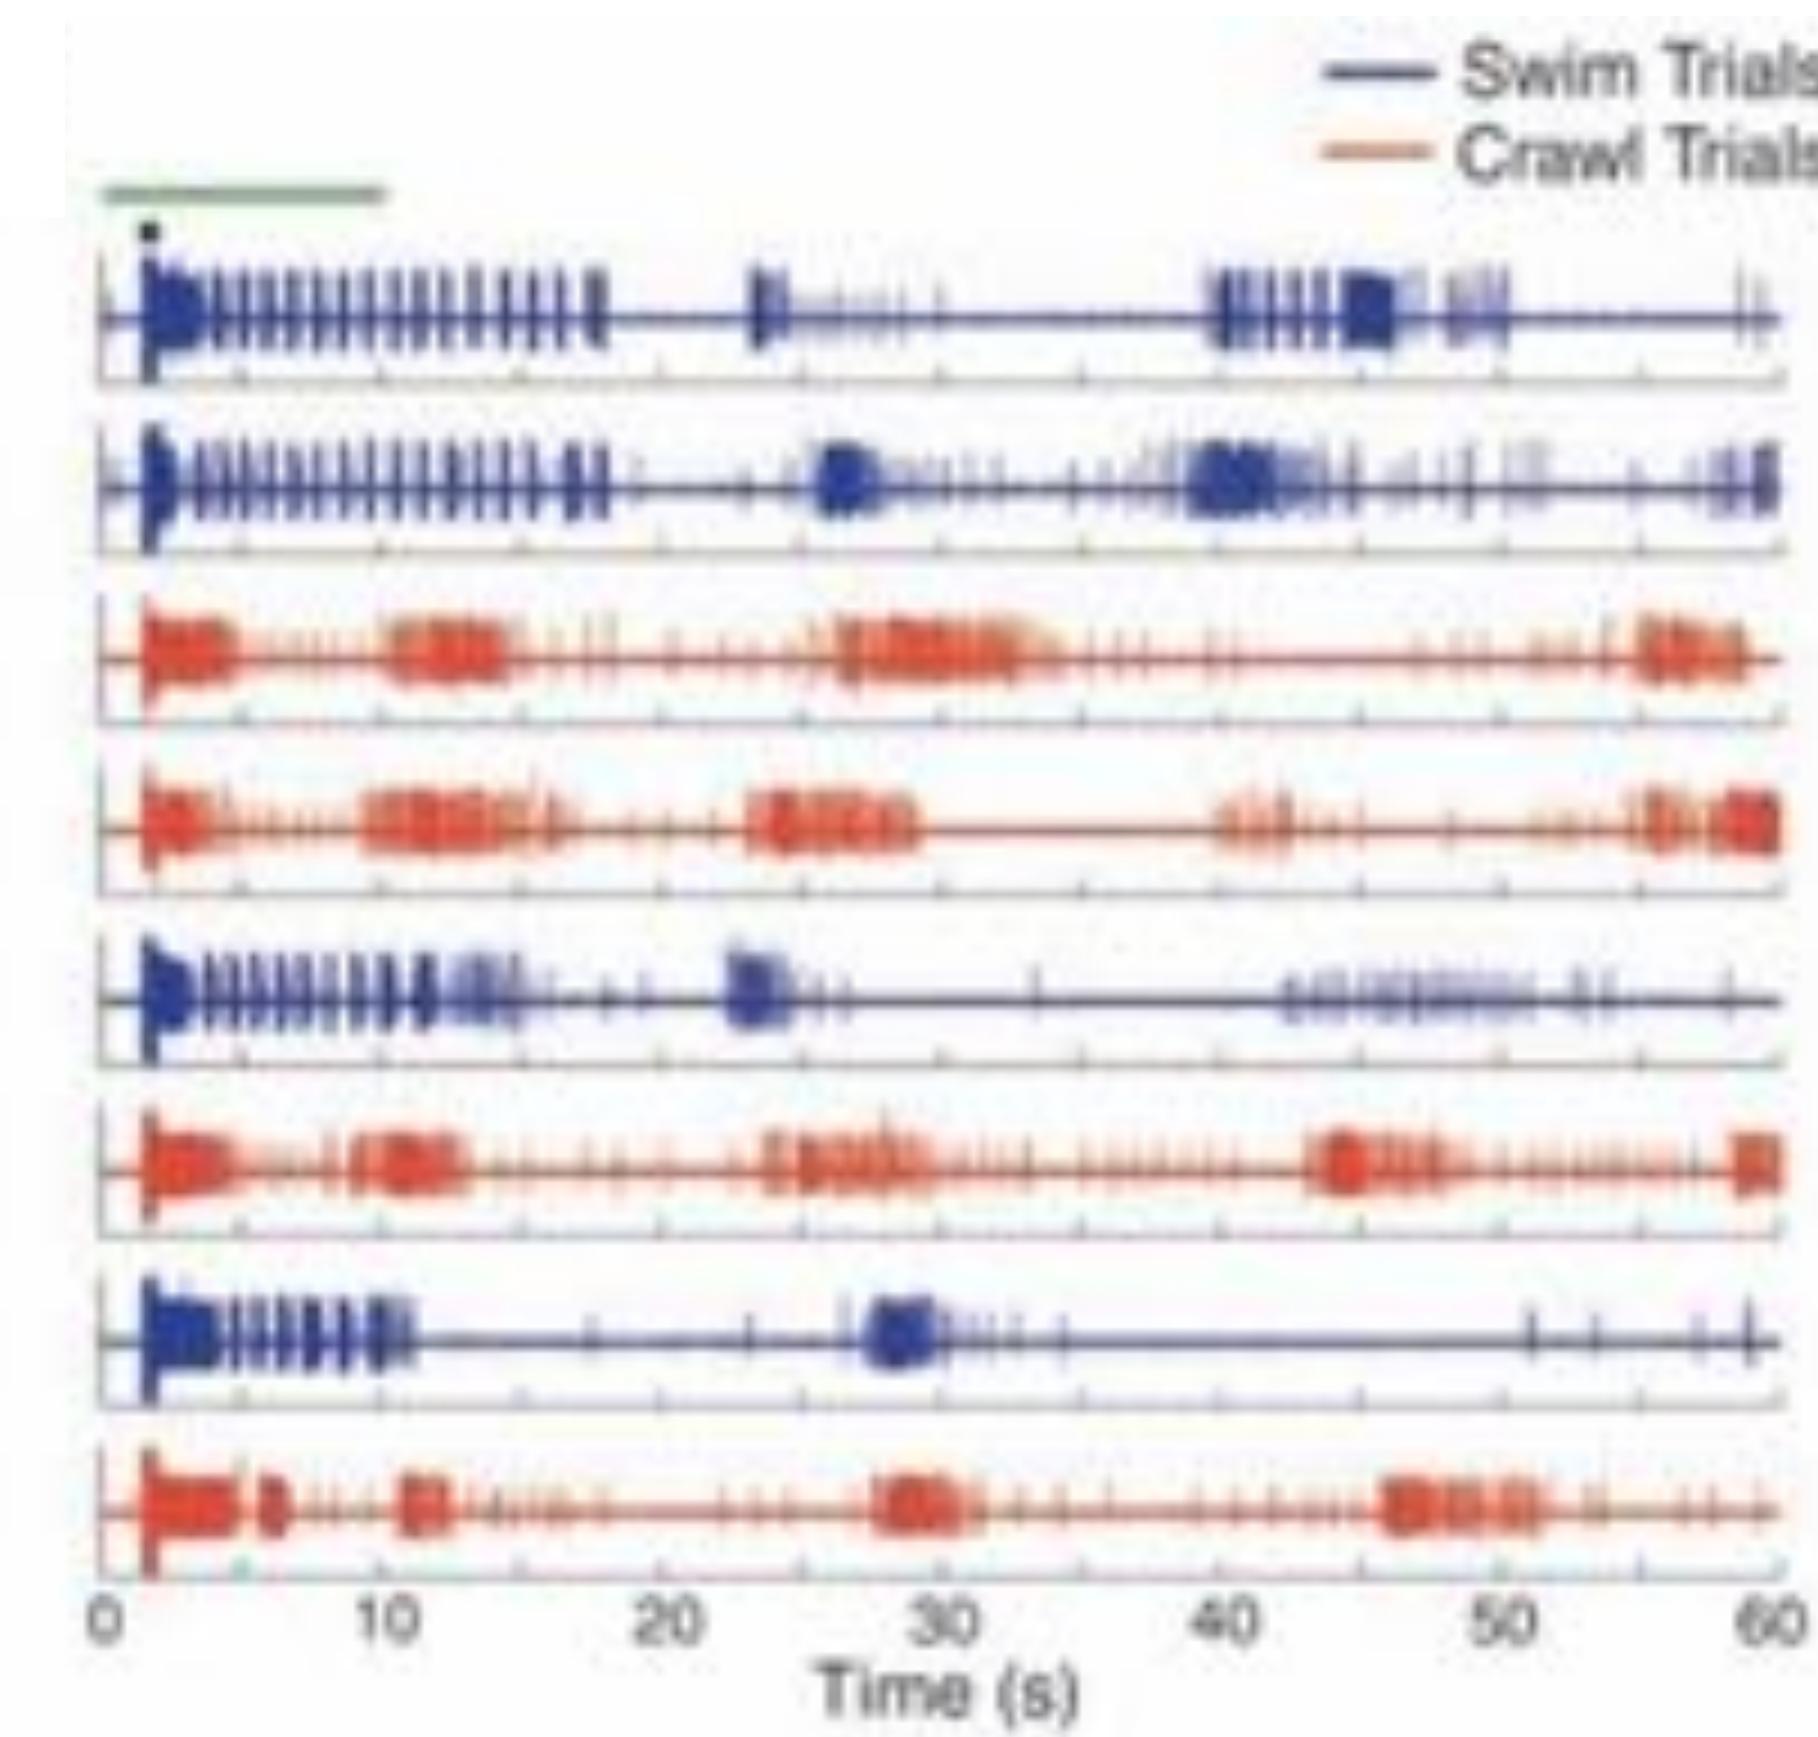
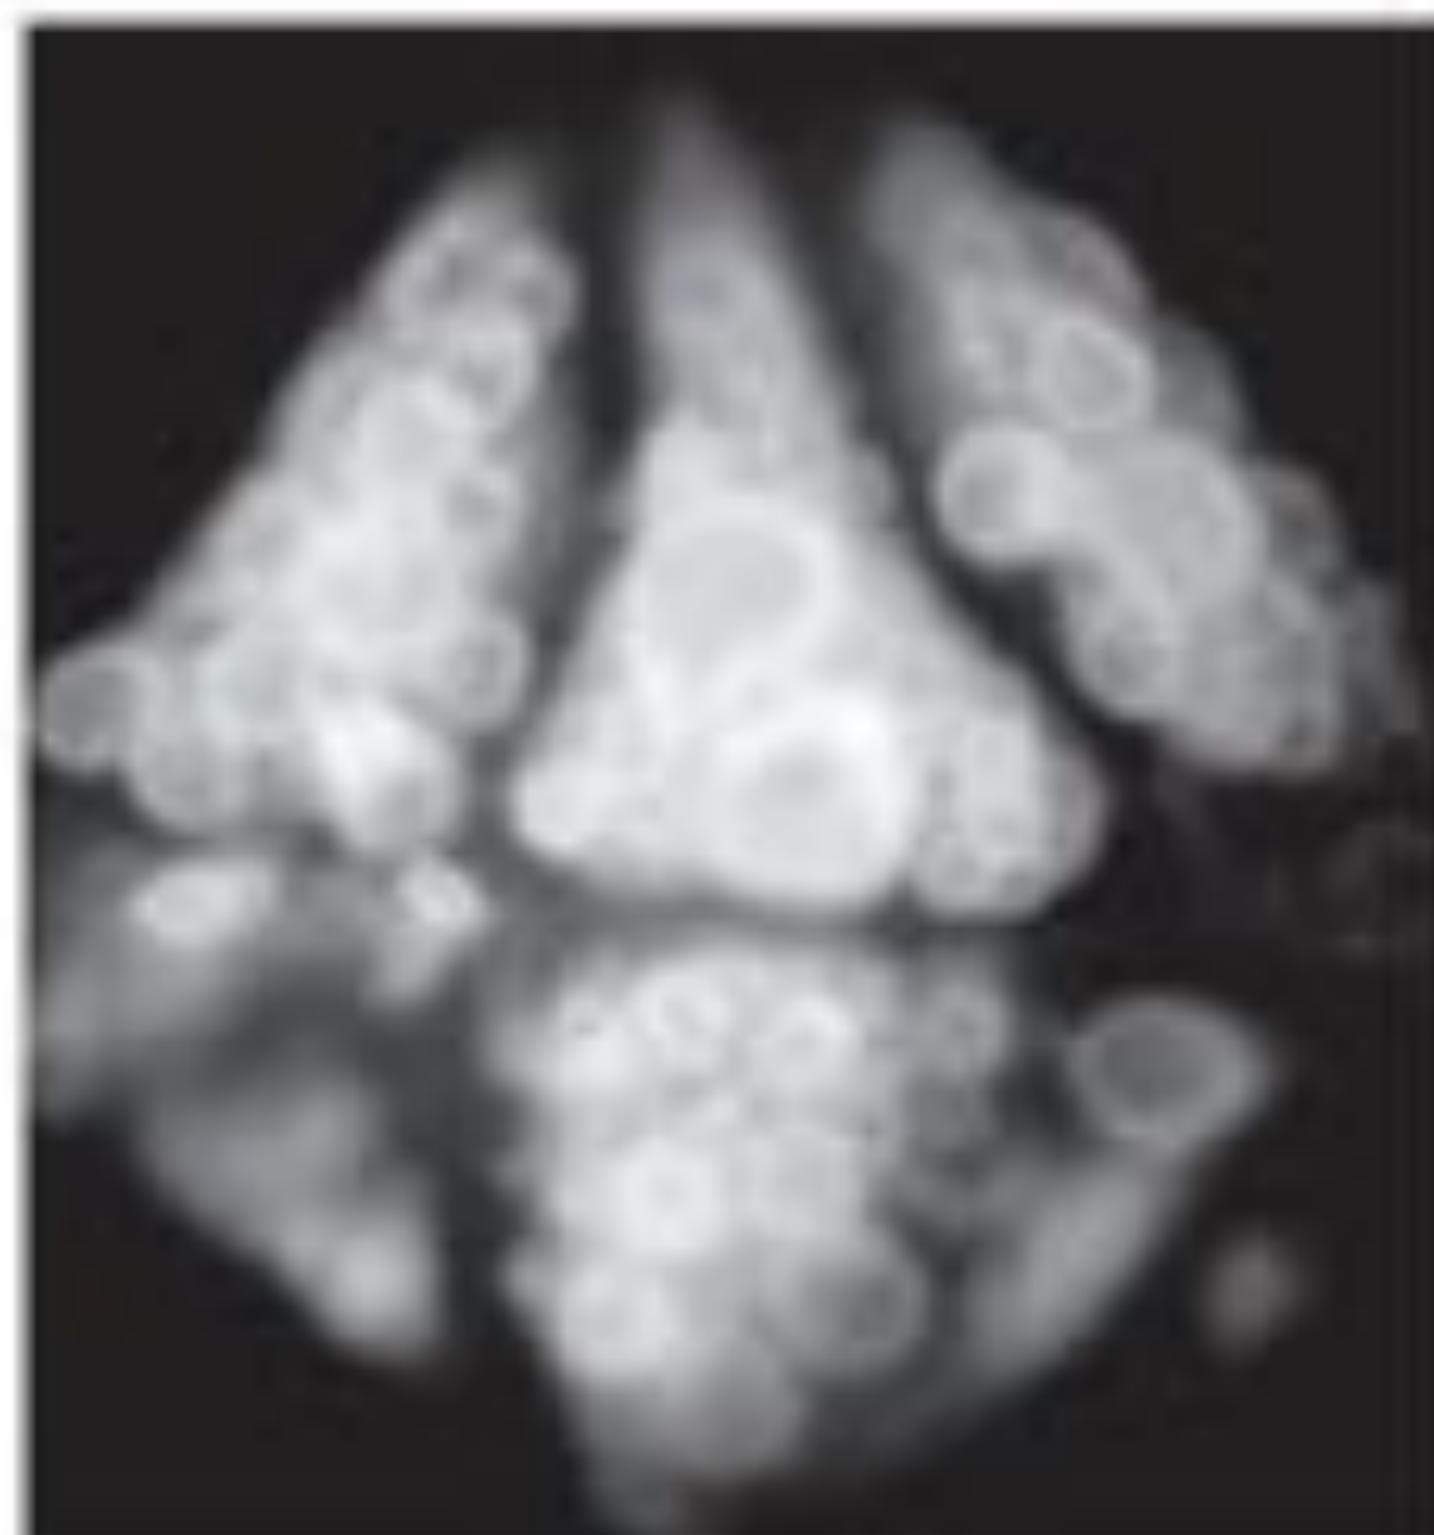


Weights

$$= \begin{bmatrix} 0.8 & -0.6 \\ -0.5 & -0.6 \\ 0.3 & 0.5 \end{bmatrix} * \begin{bmatrix} s_1(t) \\ s_2(t) \end{bmatrix}$$



**Dimensionality
reduction**

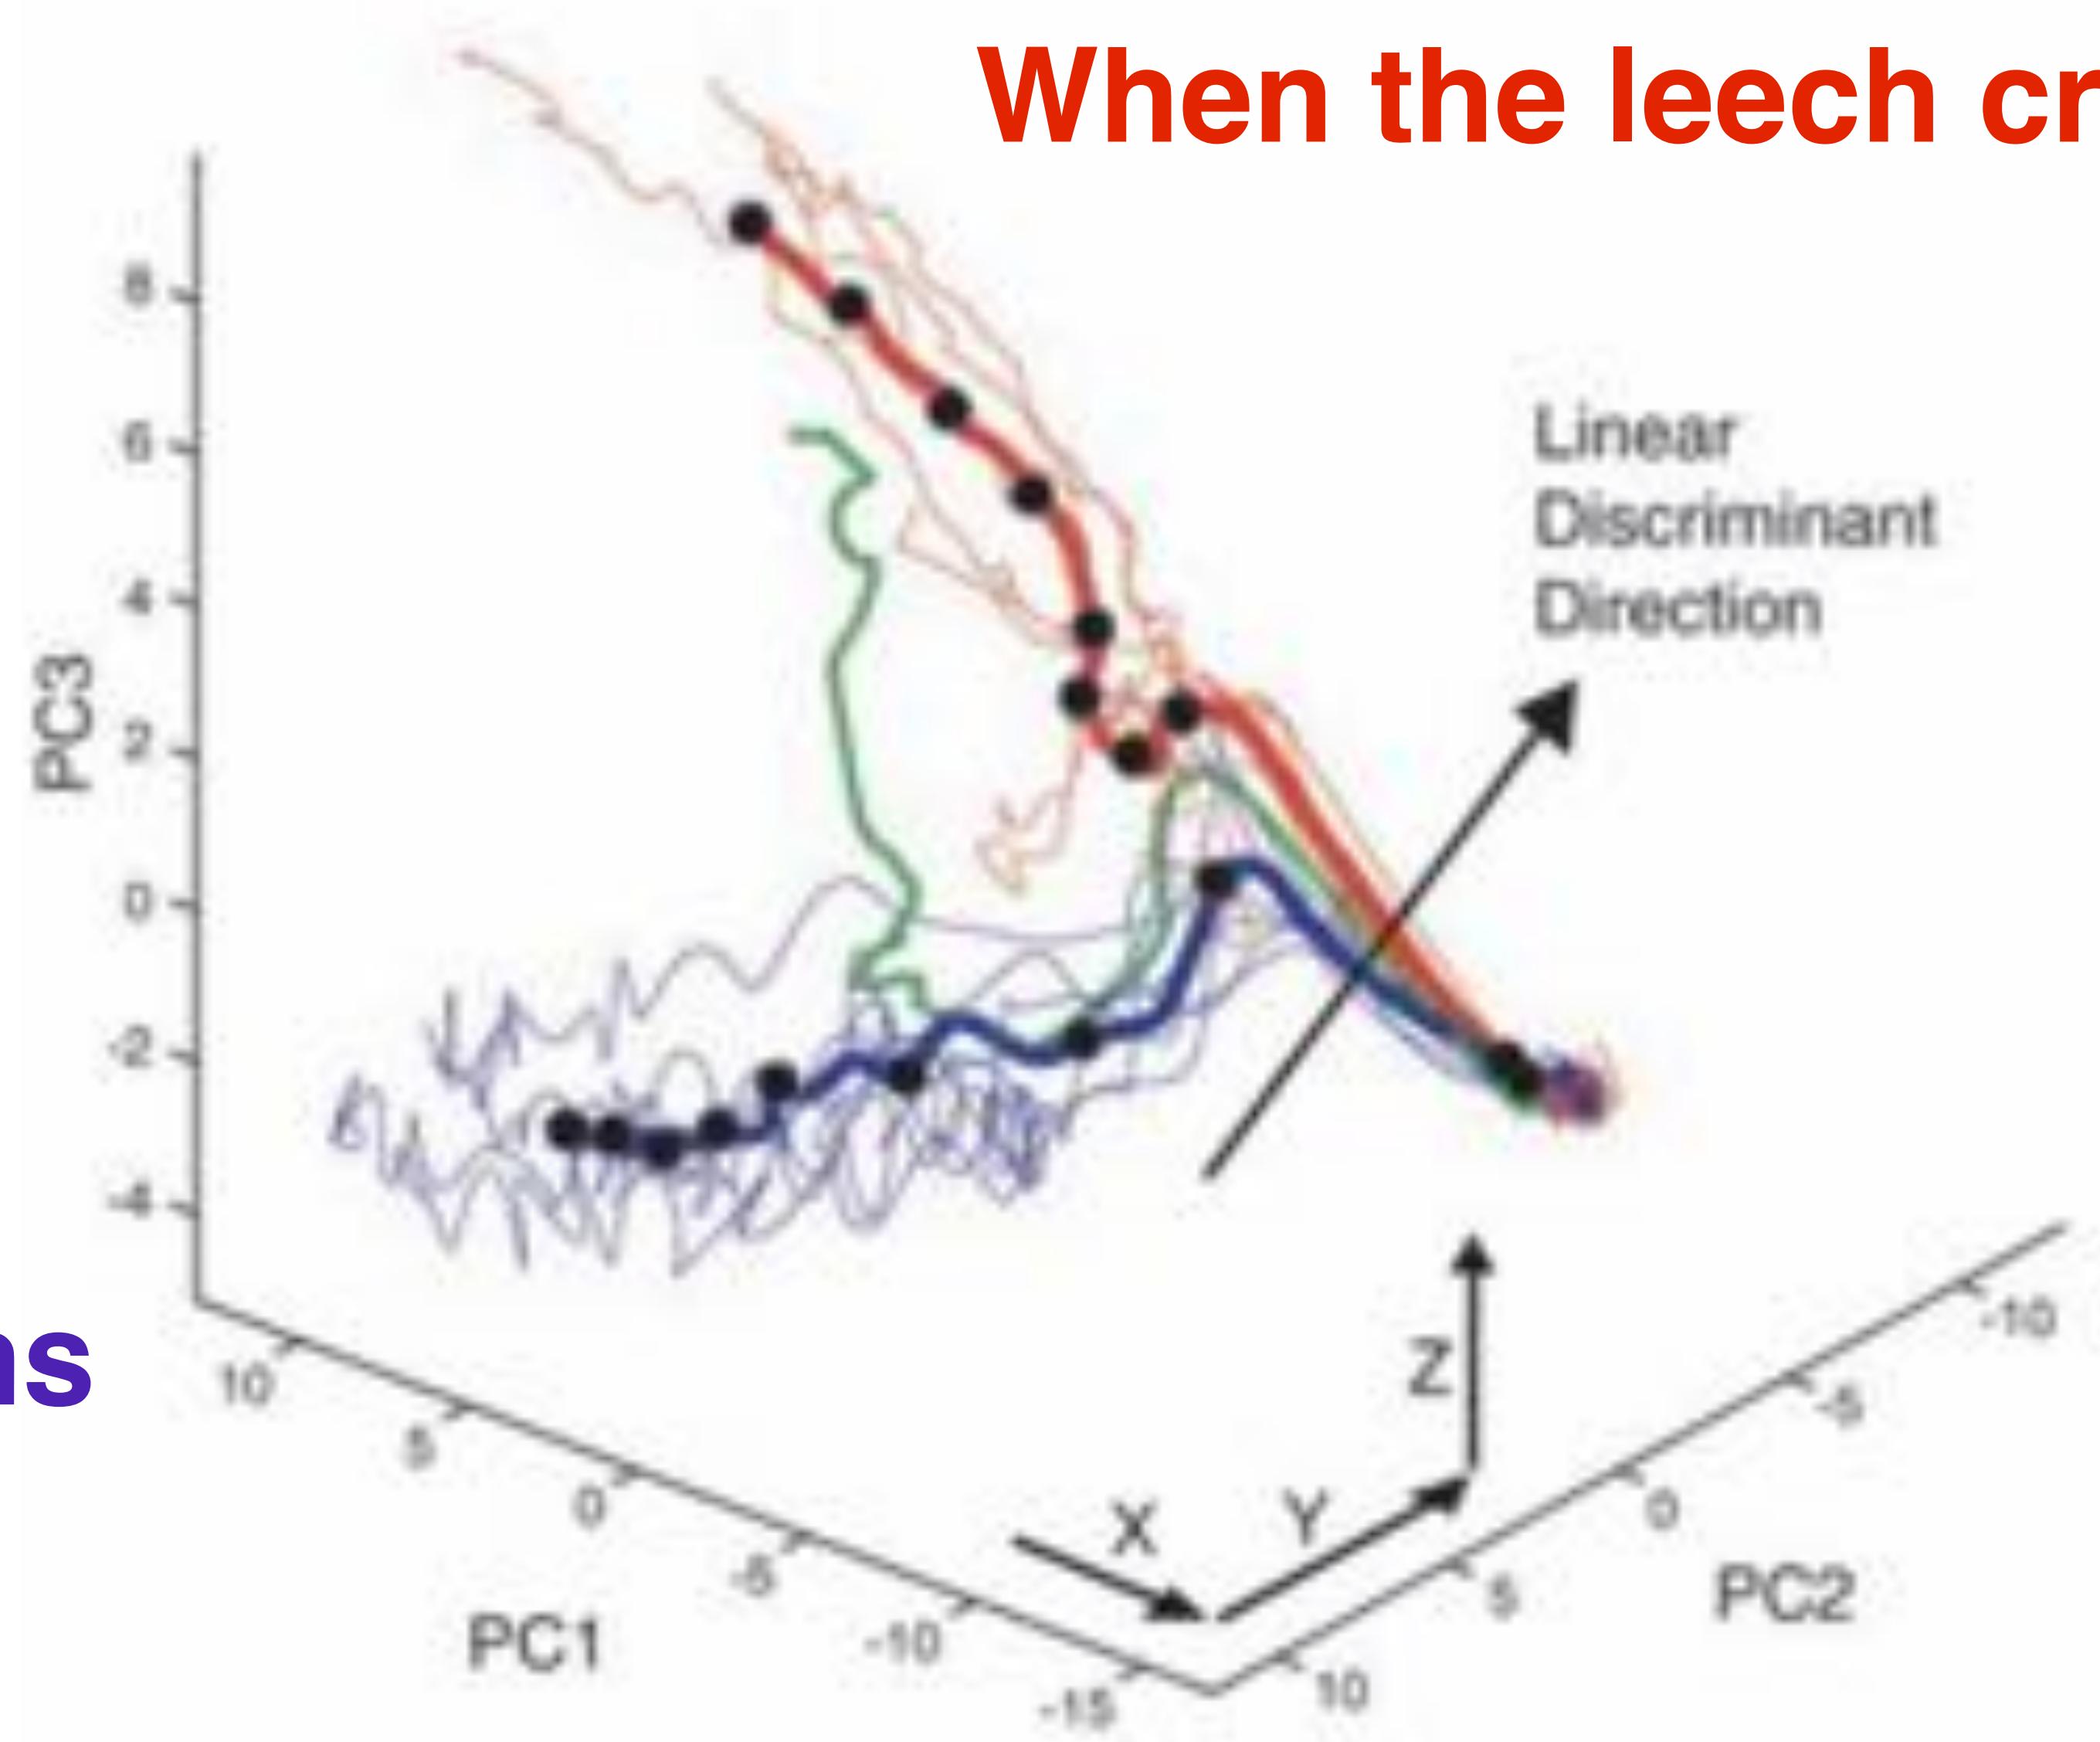


Briggman et al., 2005

**When the
leech changes
its mind!**

**When the
leech swims**

When the leech crawls



Principal Component Analysis (PCA) Overview



Raw data can be Complex, High-dimensional

To understand a phenomenon we measure various related quantities

If we knew what to measure or how to represent our measurements we might find simple relationships

But in practice we often *measure redundant signals*, e.g., US and European shoe sizes

We also *represent data via the method by which it was gathered*, e.g., pixel representation of brain imaging data

Dimensionality Reduction

Issues

- *Measure redundant signals*
- *Represent data via the method by which it was gathered*

Goal: Find a ‘better’ representation for data

- To visualize and discover hidden patterns
- Preprocessing for supervised task, e.g., feature hashing

How do we define ‘better’?

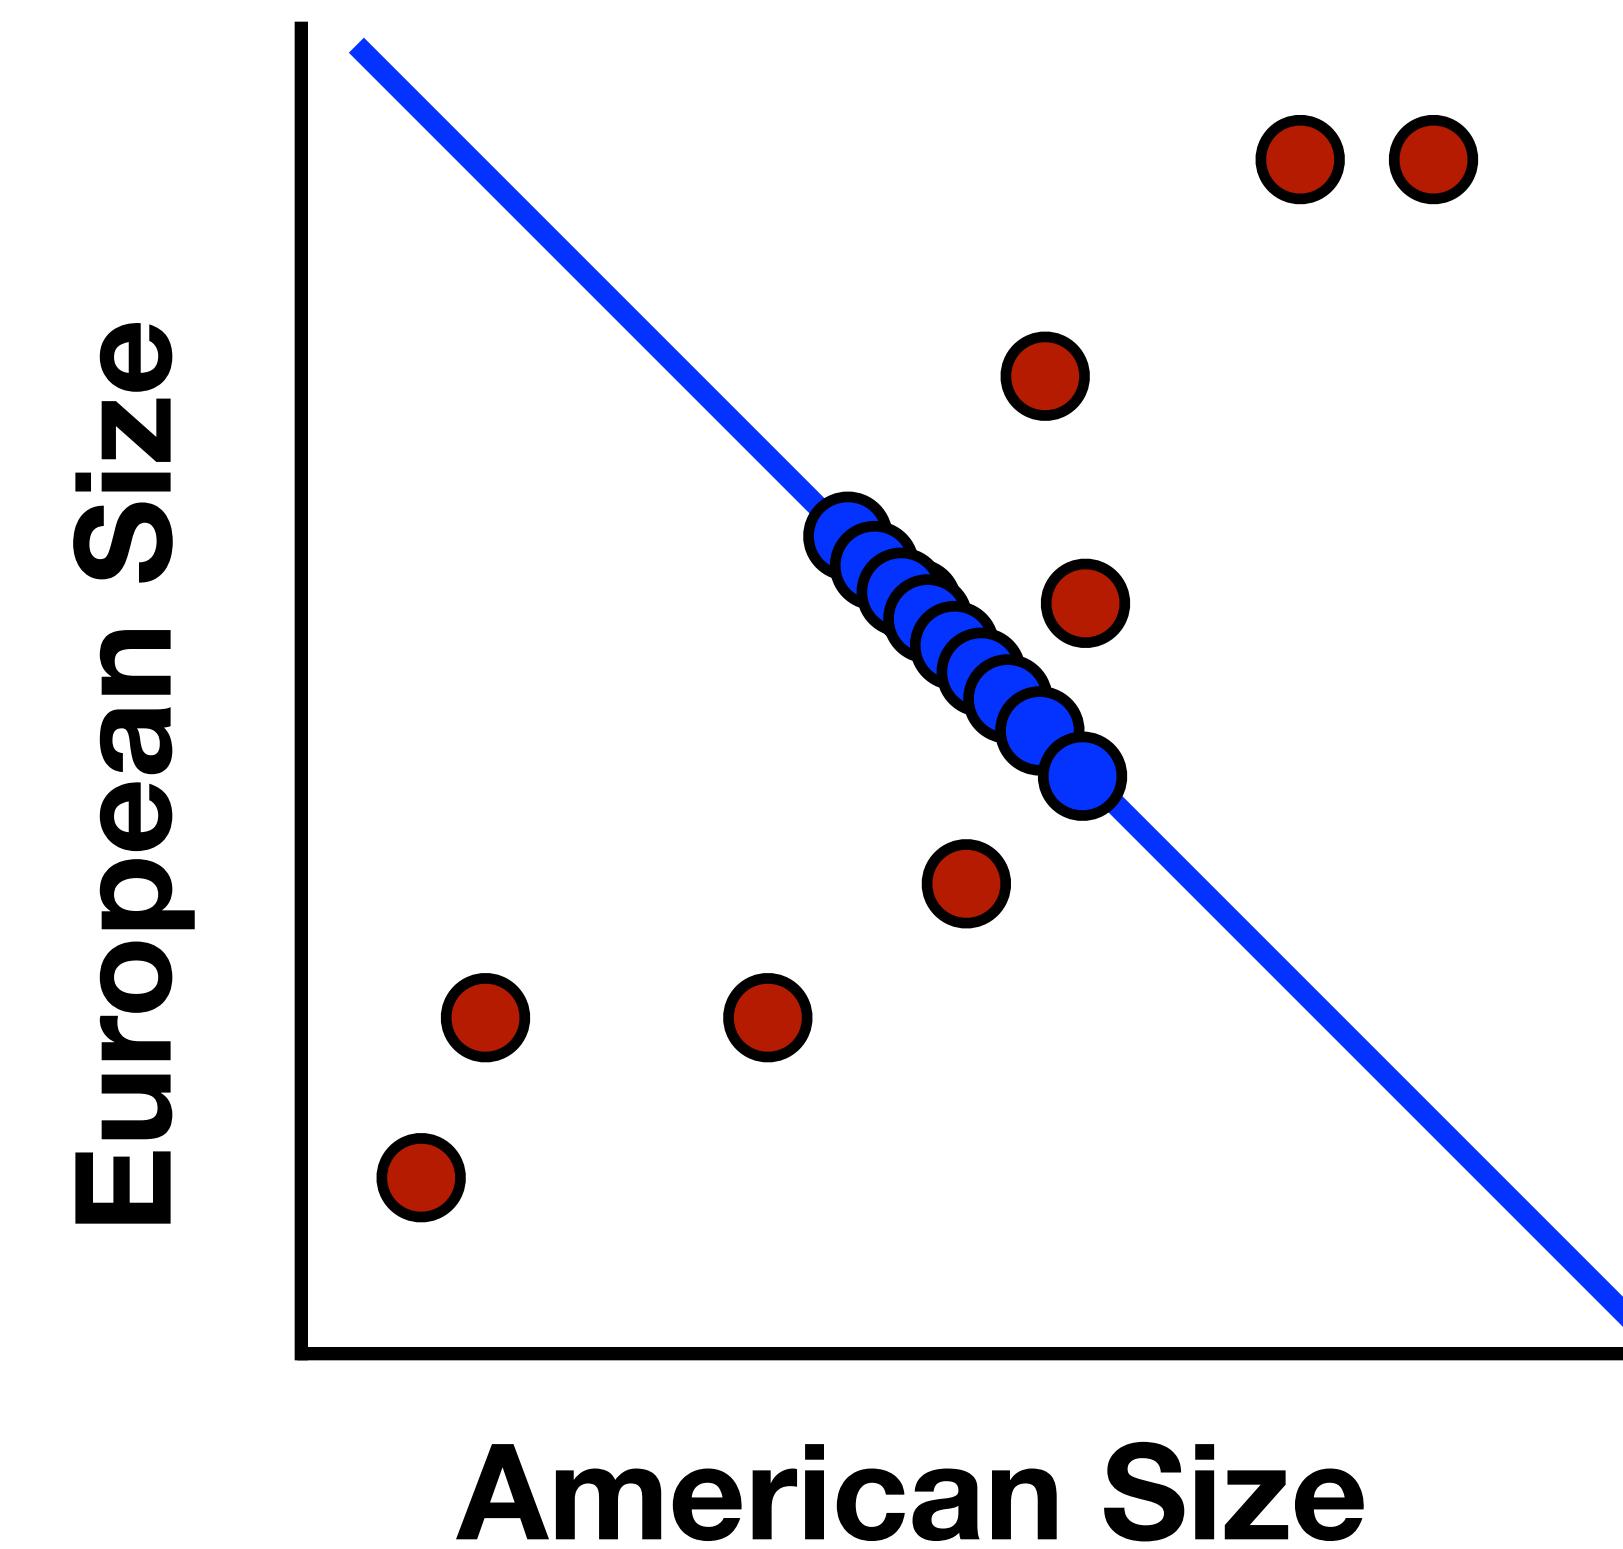
E.g., Shoe Size

We take noisy measurements on European and American scale

- Modulo noise, we expect perfect correlation

How can we do ‘better’, i.e., find a simpler, compact representation?

- Pick a direction and project onto this direction



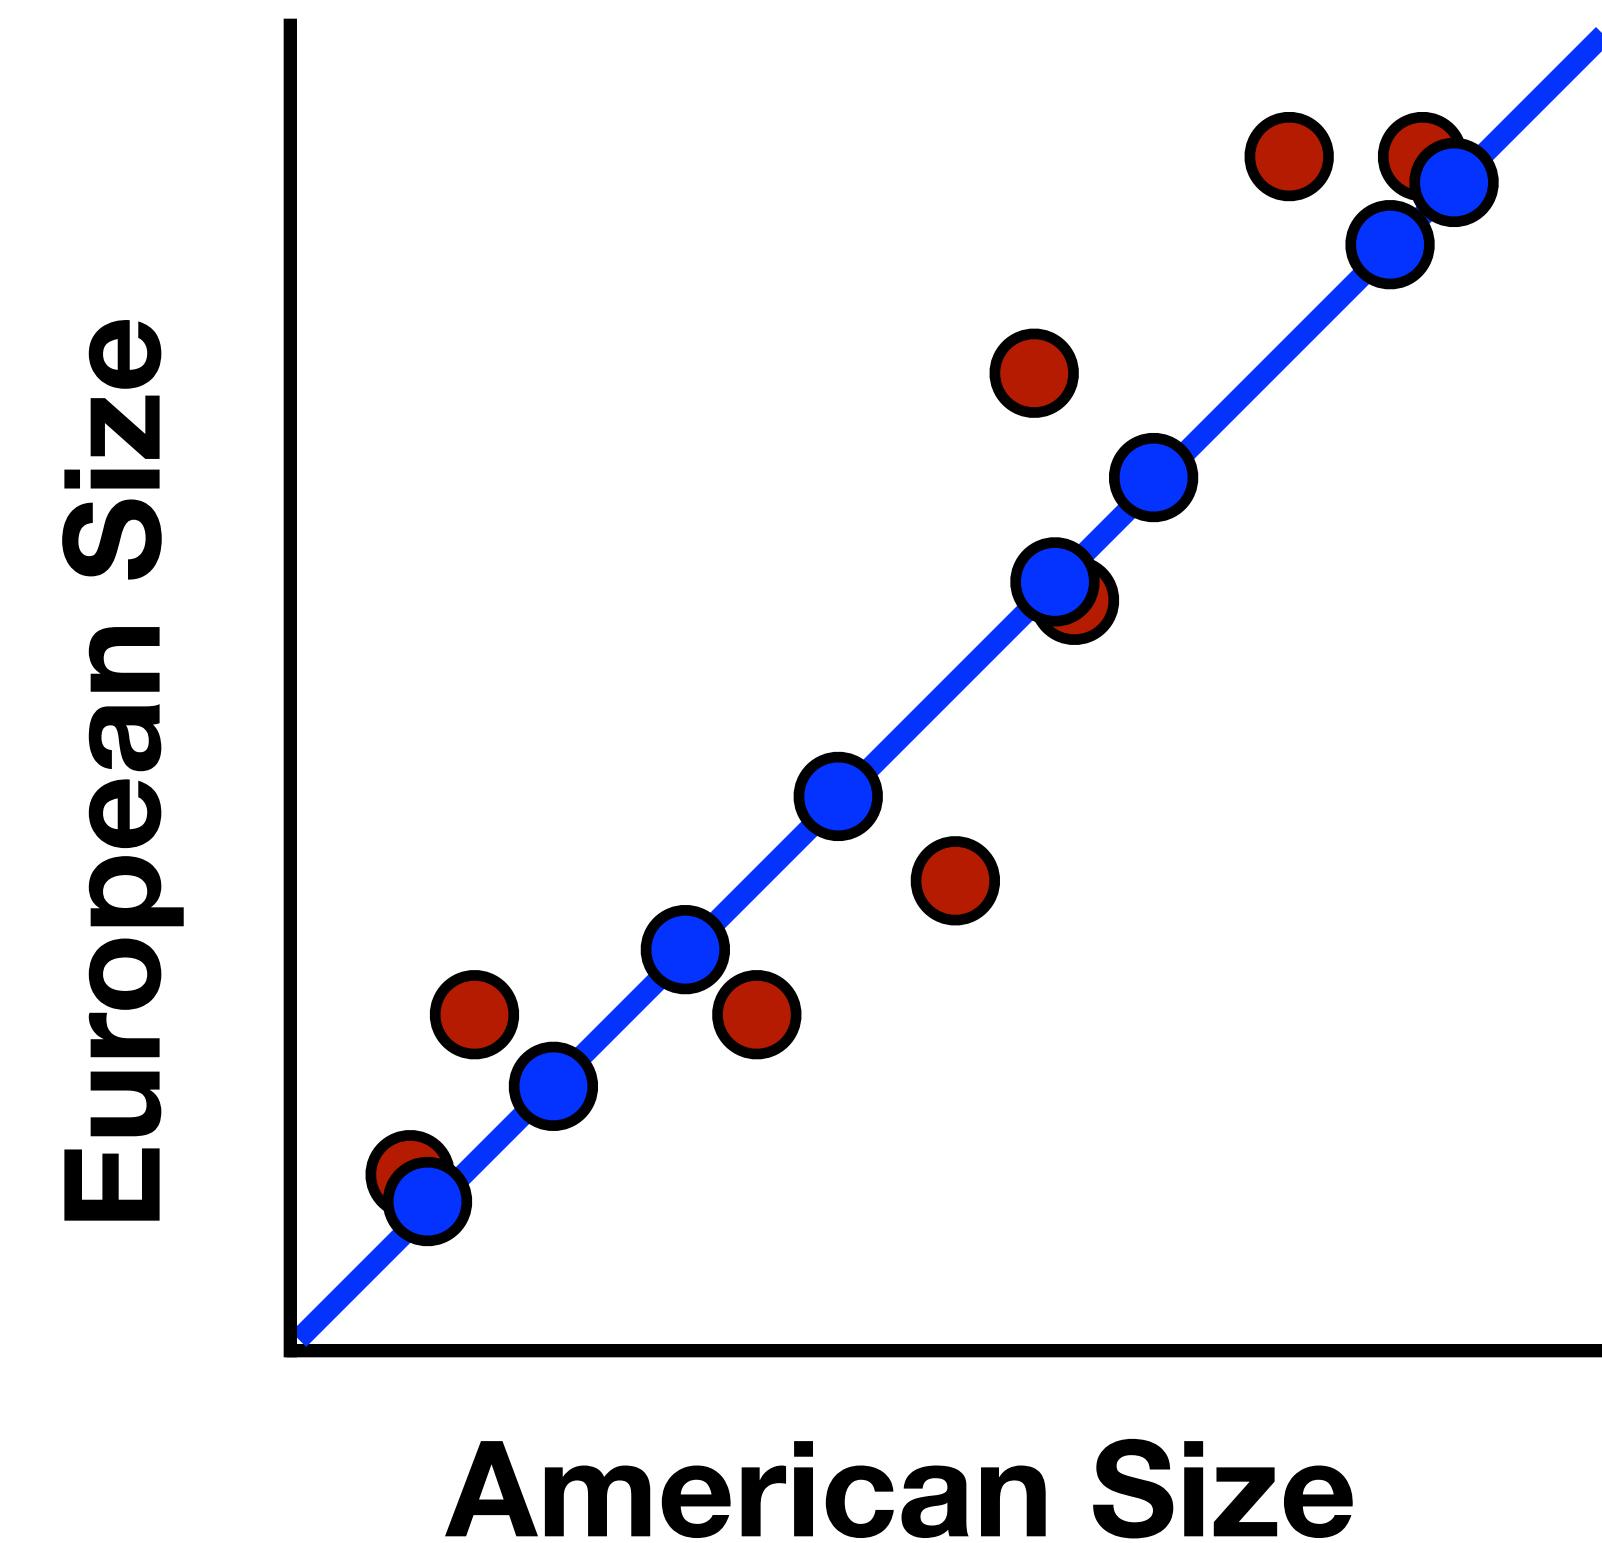
E.g., Shoe Size

We take noisy measurements on European and American scale

- Modulo noise, we expect perfect correlation

How can we do ‘better’, i.e., find a simpler, compact representation?

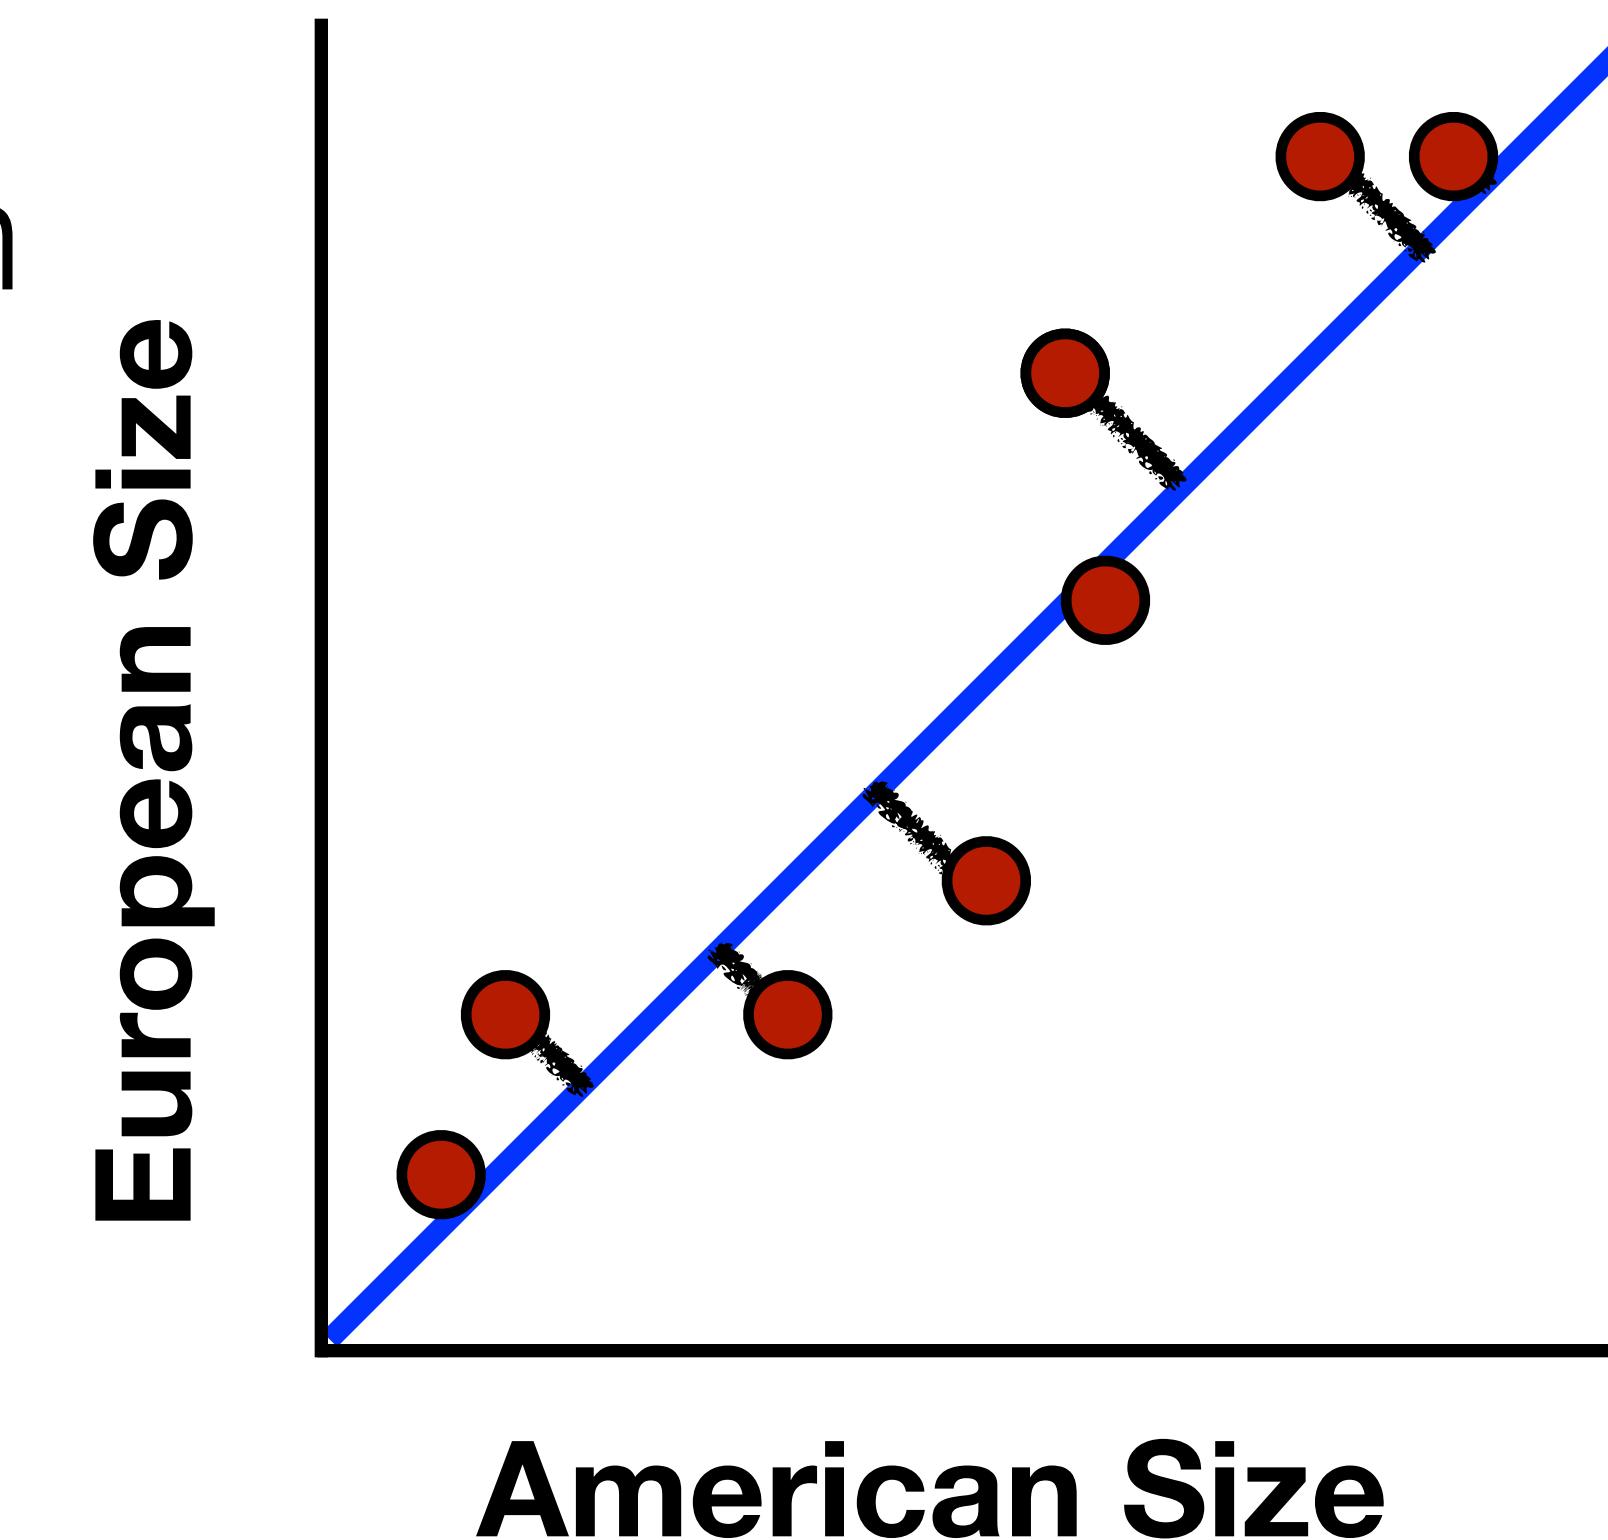
- Pick a direction and project onto this direction



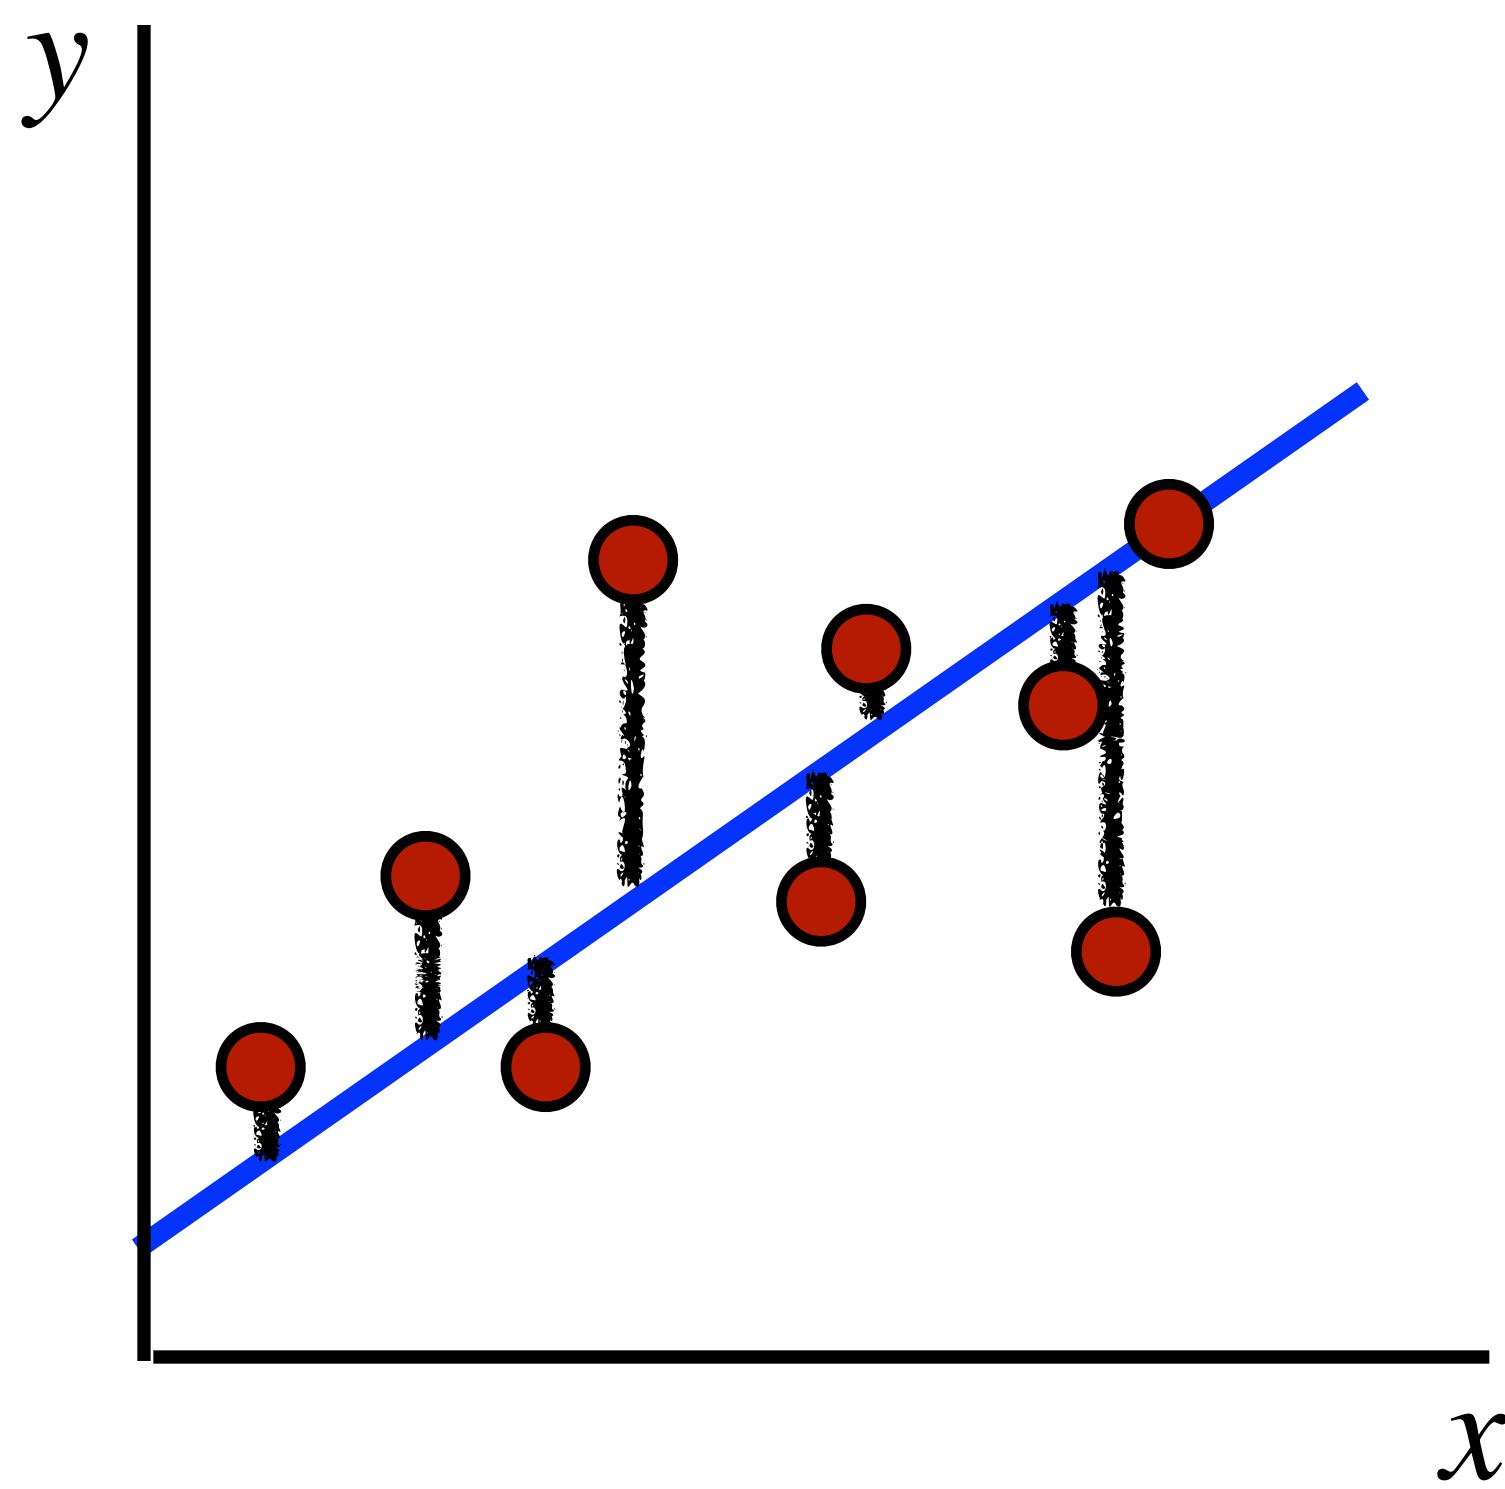
Goal: Minimize Reconstruction Error

Minimize Euclidean distances between original points and their projections

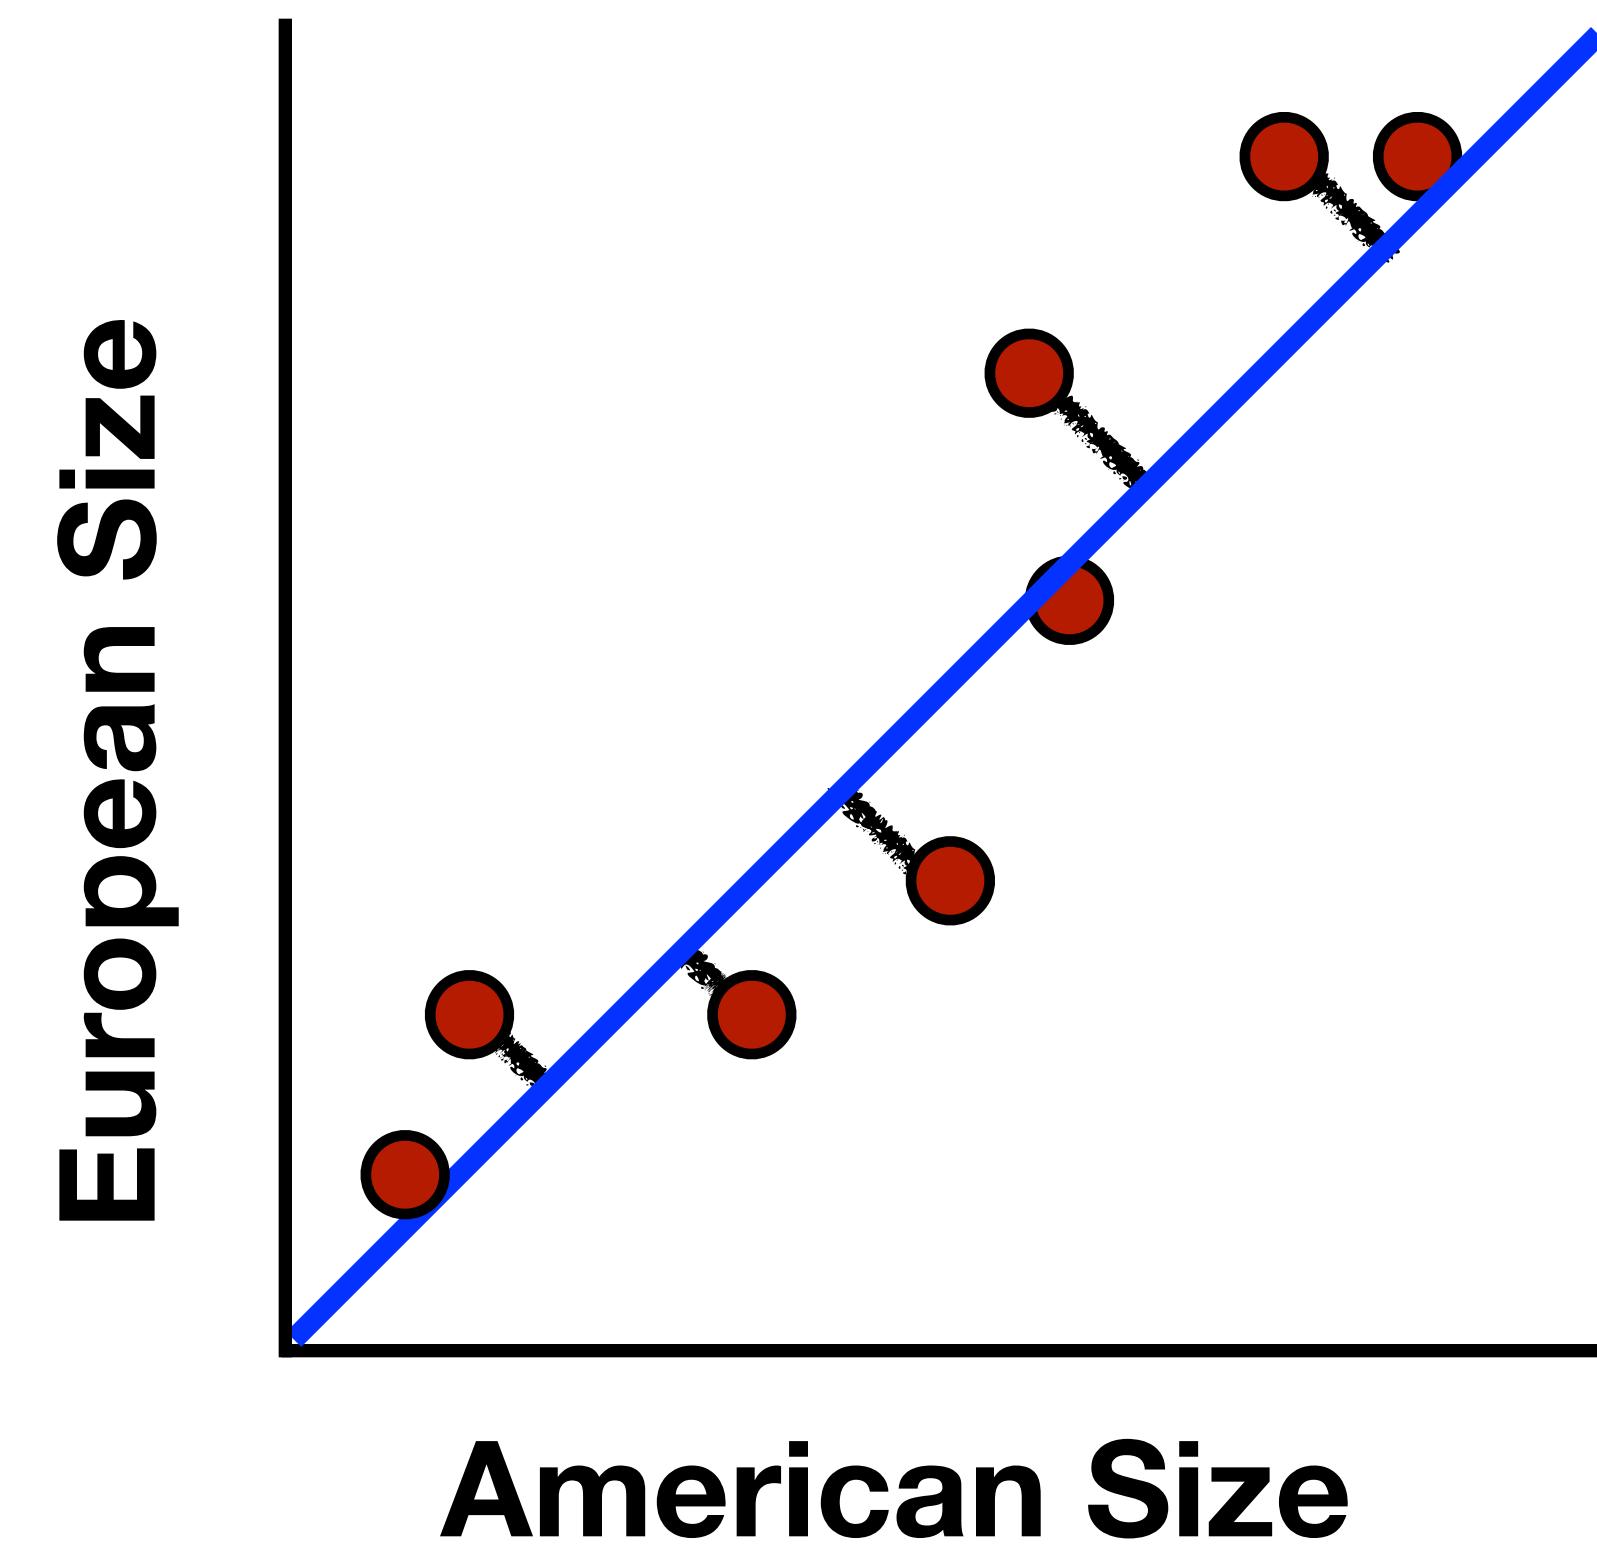
PCA solution solves this problem!



Linear Regression – predict y from x .
Evaluate accuracy of predictions
(represented by blue line) by **vertical**
distances between points and the line



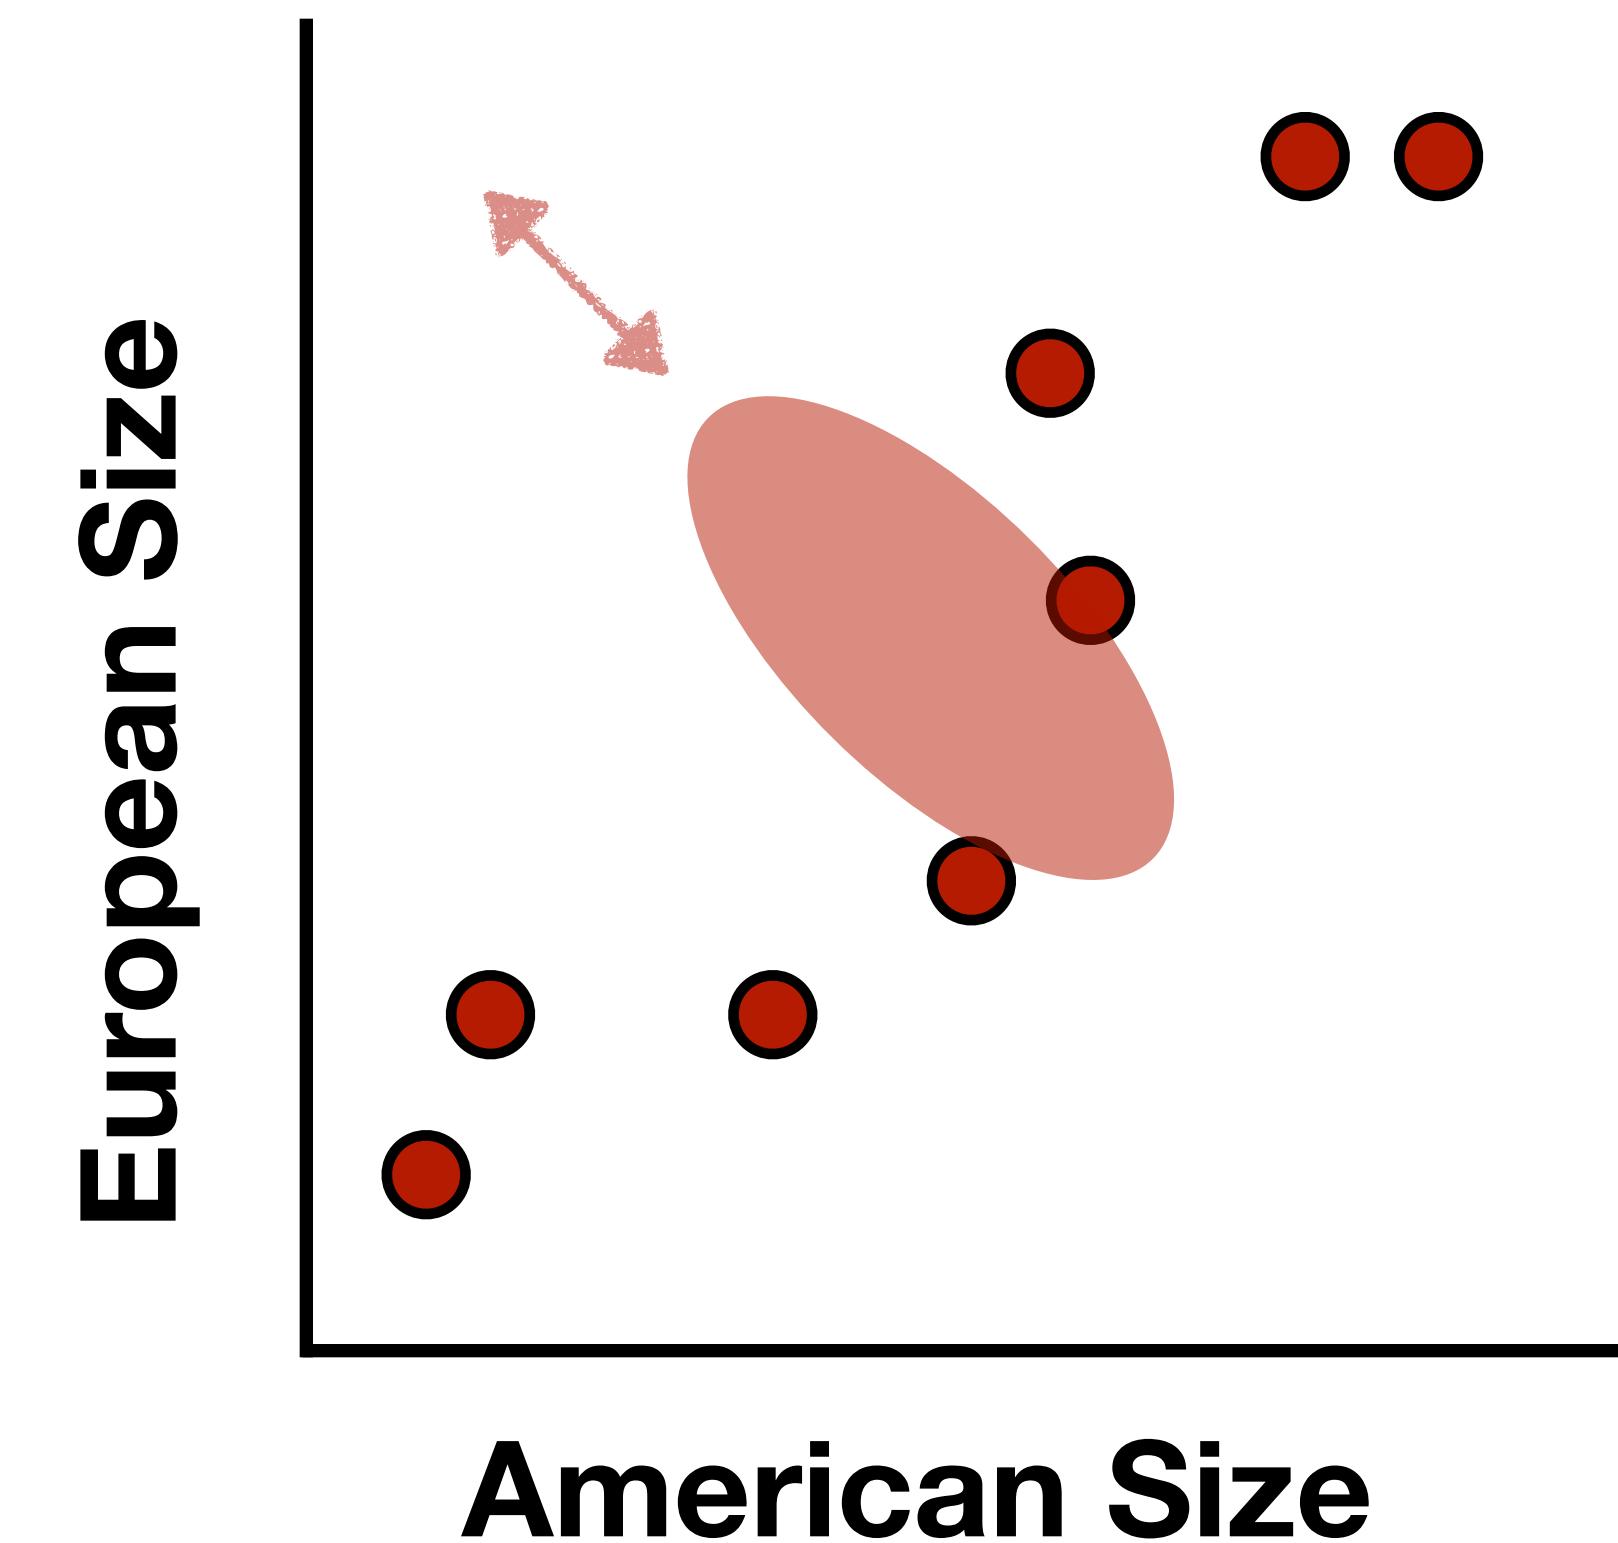
PCA – reconstruct 2D data via 2D
data with single degree of freedom.
Evaluate reconstructions (represented
by blue line) by **Euclidean** distances



Another Goal: Maximize Variance

To identify patterns we want to study variation across observations

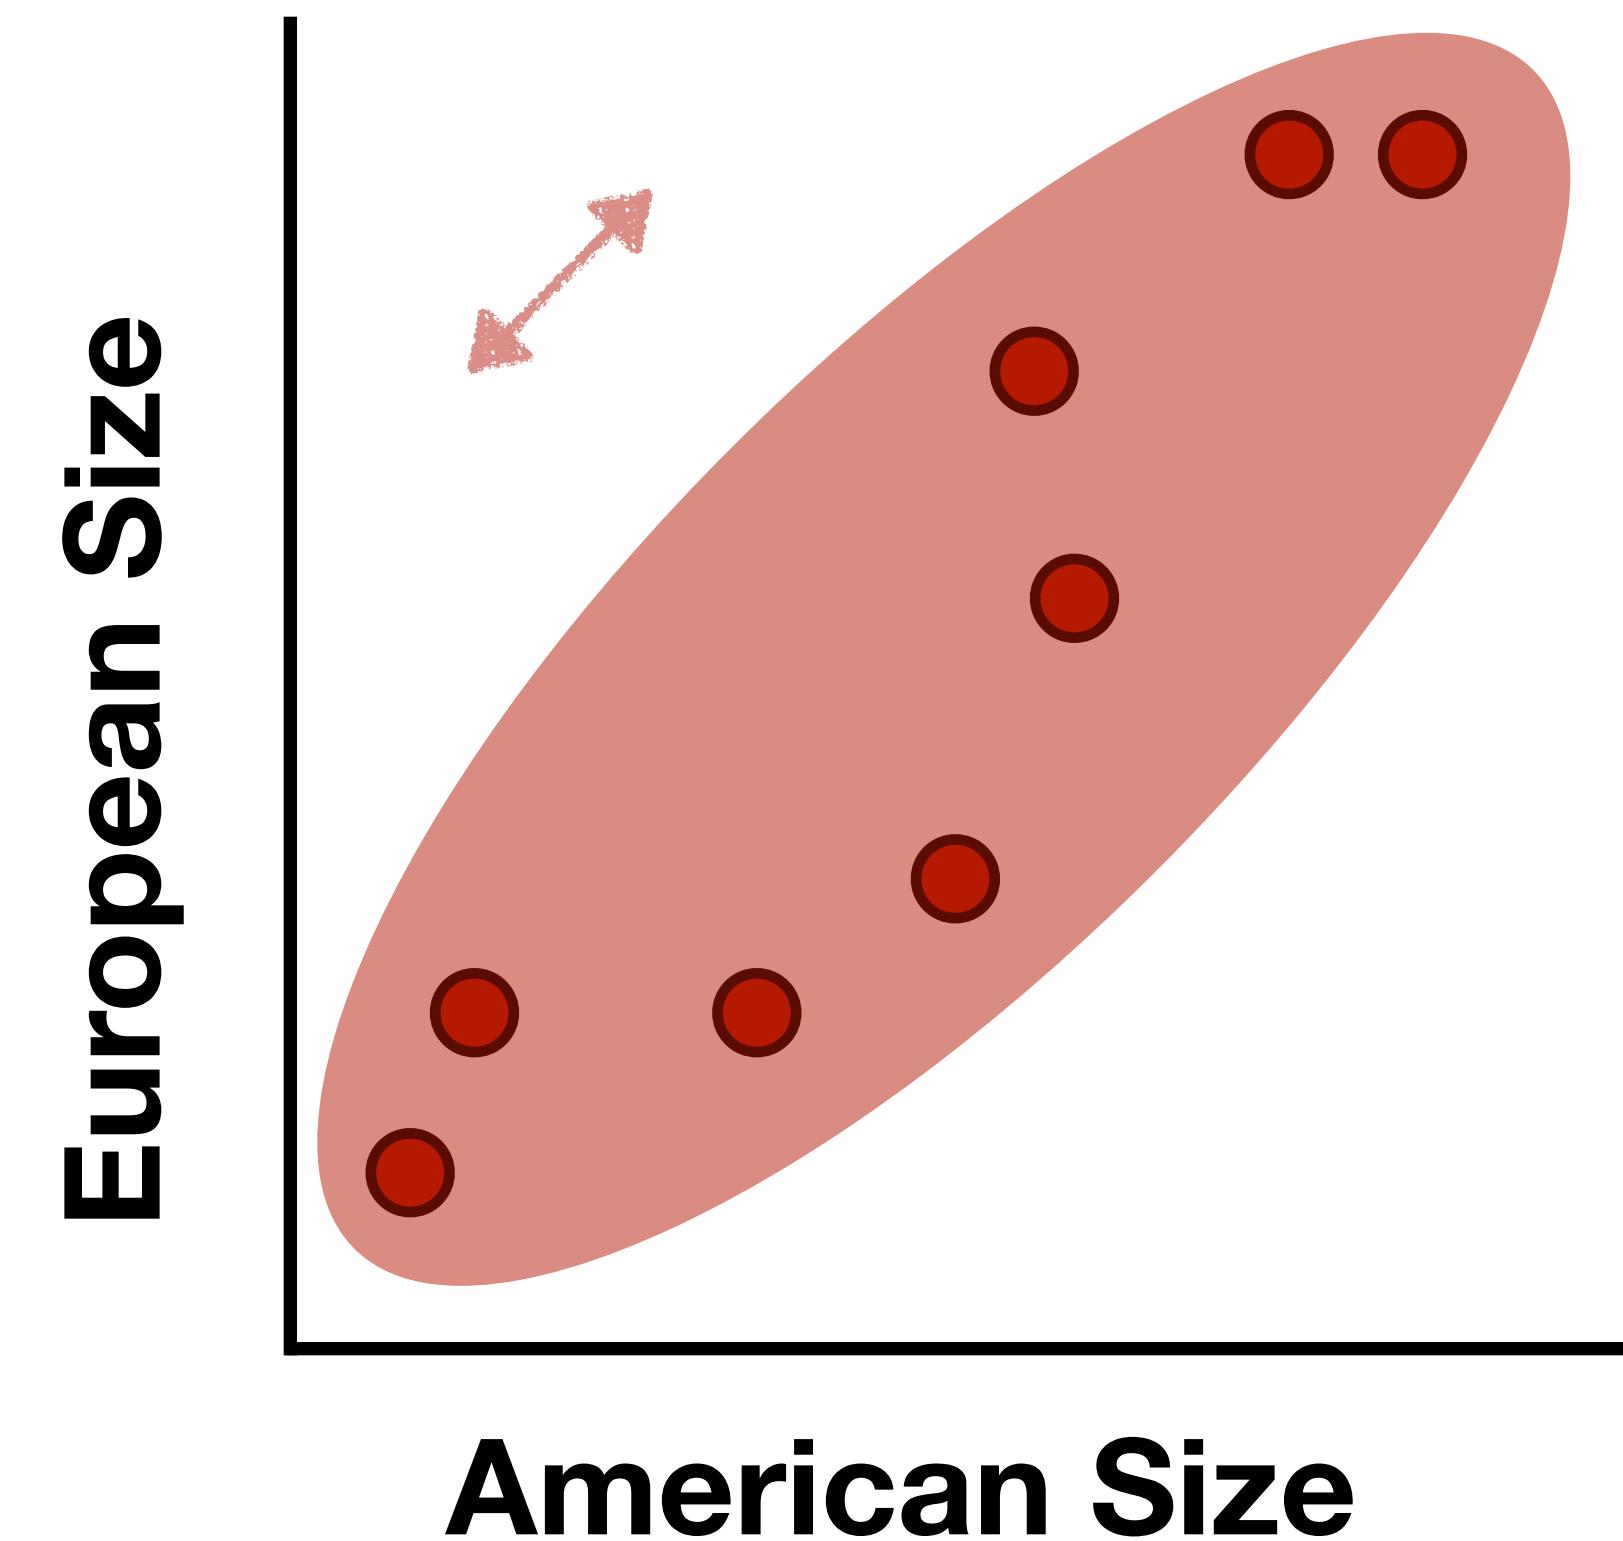
Can we do ‘better’, i.e., find a compact representation that captures variation?



Another Goal: Maximize Variance

To identify patterns we want to study variation across observations

Can we do ‘better’, i.e., find a compact representation that captures variation?

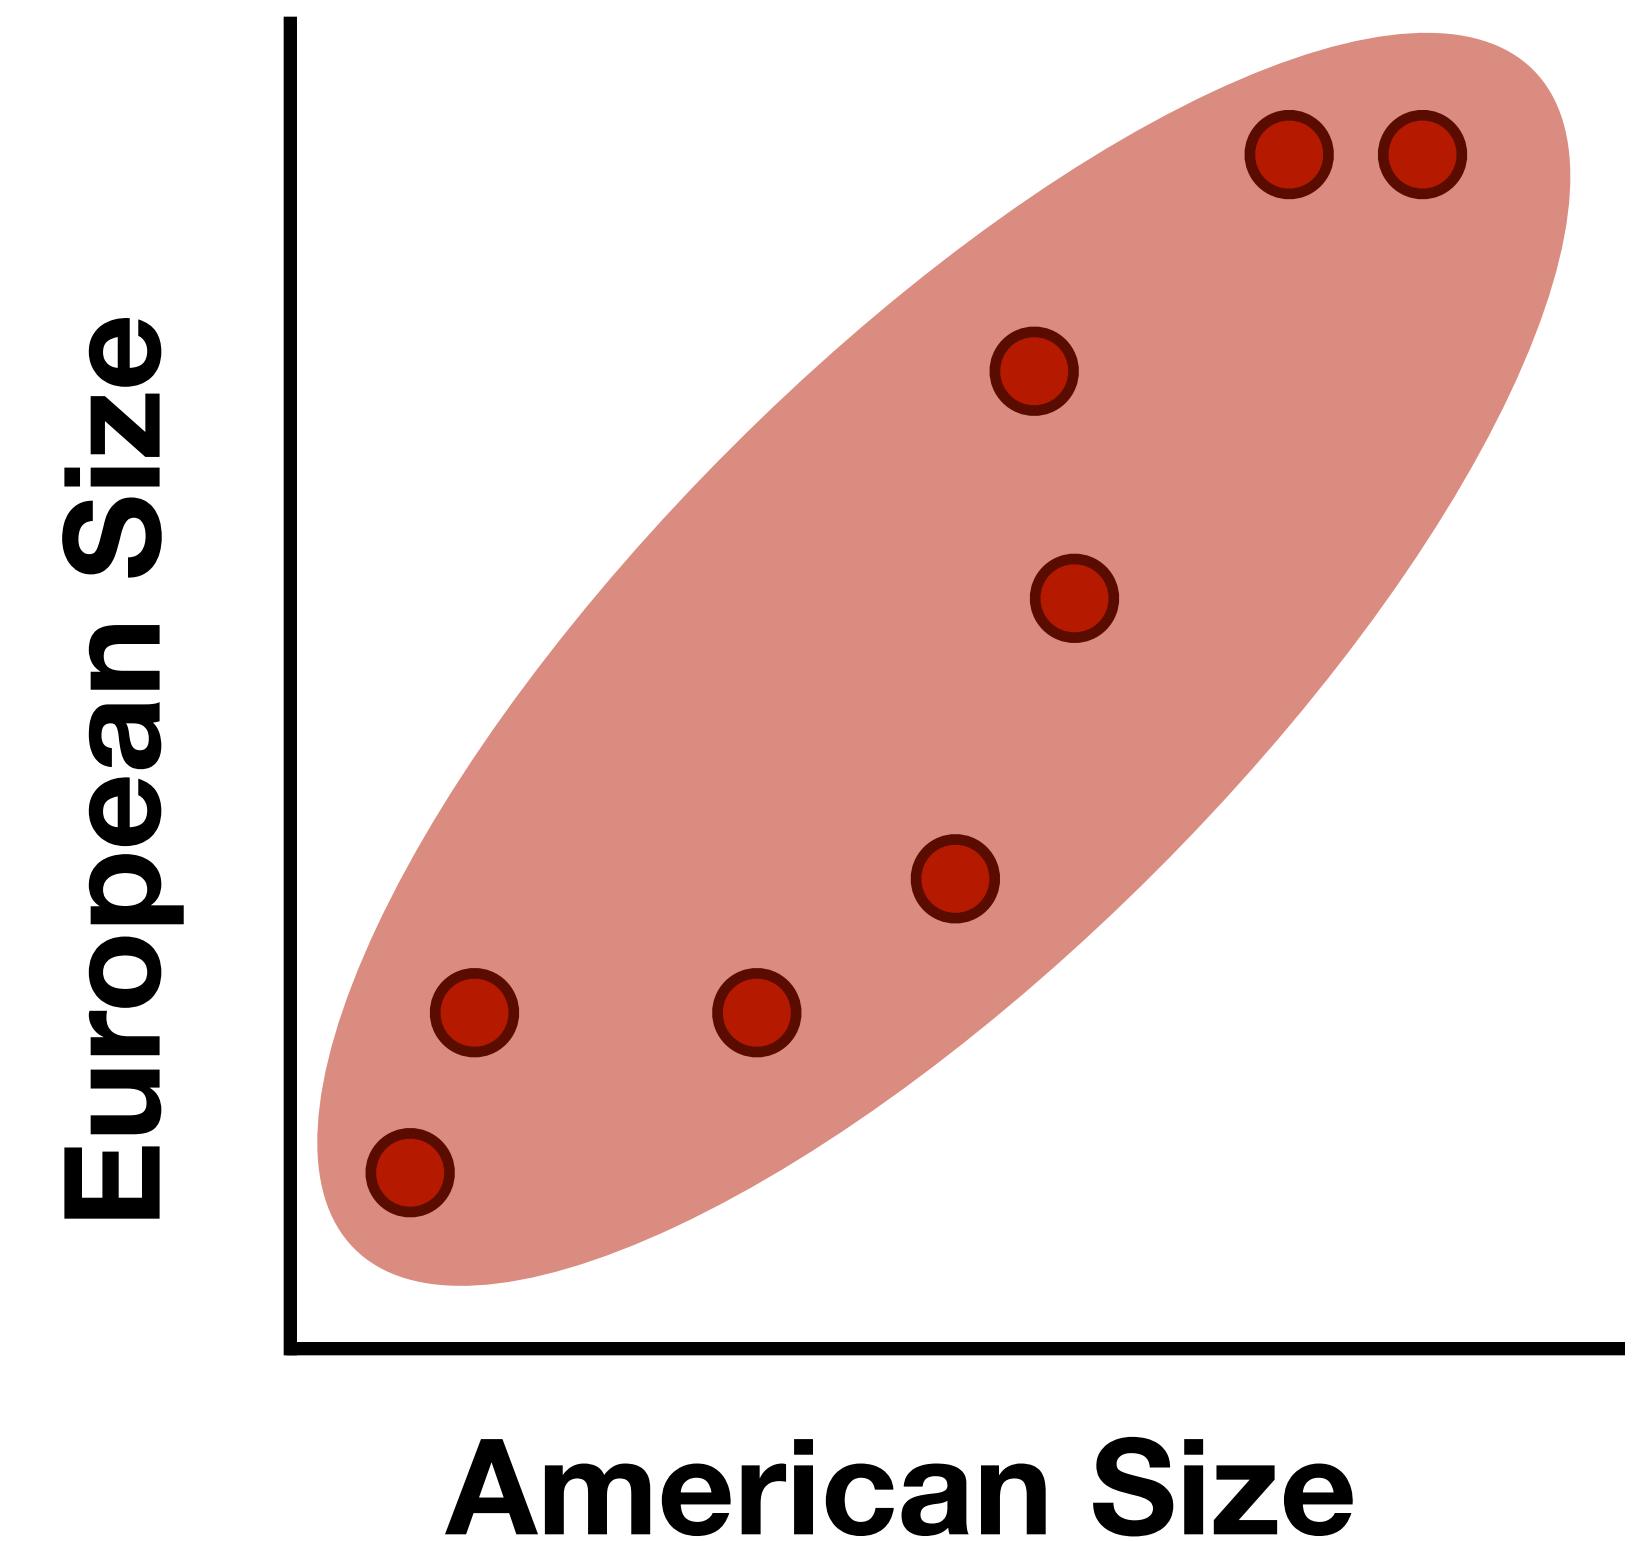


Another Goal: Maximize Variance

To identify patterns we want to study variation across observations

Can we do ‘better’, i.e., find a compact representation that captures variation?

PCA solution finds directions of maximal variance!



PCA Assumptions and Solution



PCA Formulation

PCA: find lower-dimensional representation of raw data

- \mathbf{X} is $n \times d$ (raw data)
- $\mathbf{Z} = \mathbf{XP}$ is $n \times k$ (reduced representation, PCA ‘scores’)
- \mathbf{P} is $d \times k$ (columns are k principal components)
- Variance constraints

Linearity assumption ($\mathbf{Z} = \mathbf{XP}$) simplifies
problem

$$\mathbf{Z} = \mathbf{X} \mathbf{P}$$

Given n training points with d features:

- $\mathbf{X} \in \mathbb{R}^{n \times d}$: matrix storing points
- $x_j^{(i)}$: j th feature for i th point
- μ_j : mean of j th feature

Variance of 1st feature $\sigma_1^2 = \frac{1}{n} \sum_{i=1}^n (x_1^{(i)} - \mu_1)^2$

Variance of 1st feature
(assuming zero mean) $\sigma_1^2 = \frac{1}{n} \sum_{i=1}^n (x_1^{(i)})^2$

Given n training points with d features:

- $\mathbf{X} \in \mathbb{R}^{n \times d}$: matrix storing points
- $x_j^{(i)}$: j th feature for i th point
- μ_j : mean of j th feature

Covariance of 1st and 2nd
features (assuming zero mean)

$$\sigma_{12} = \frac{1}{n} \sum_{i=1}^n x_1^{(i)} x_2^{(i)}$$

- Symmetric: $\sigma_{12} = \sigma_{21}$
- Zero \rightarrow uncorrelated
- Large magnitude \rightarrow (anti) correlated / redundant
- $\sigma_{12} = \sigma_1^2 = \sigma_2^2 \rightarrow$ features are the same

Covariance Matrix

Covariance matrix generalizes this idea for many features

$d \times d$ covariance matrix with
zero mean features

$$\mathbf{C}_{\mathbf{X}} = \frac{1}{n} \mathbf{X}^T \mathbf{X}$$

- i th diagonal entry equals variance of i th feature
- ij th entry is covariance between i th and j th features
- Symmetric (makes sense given definition of covariance)

$$\text{Variance: } \sigma_1^2 = \frac{1}{n} \sum_{i=1}^n (x_1^{(i)})^2$$

$$\begin{bmatrix} 2 & -1 & -1 \\ 3 & 2 & -5 \end{bmatrix} \begin{bmatrix} 2 \\ -1 \\ -1 \end{bmatrix} = \begin{bmatrix} 6 \\ \quad \end{bmatrix}$$

\mathbf{x}^\top \mathbf{x} $\mathbf{x}^\top \mathbf{x}$

$$\text{Variance: } \sigma_1^2 = \frac{1}{n} \sum_{i=1}^n (x_1^{(i)})^2$$

$$\text{Covariance: } \sigma_{12} = \frac{1}{n} \sum_{i=1}^n x_1^{(i)} x_2^{(i)}$$

$$\begin{bmatrix} 2 & -1 & -1 \\ 3 & 2 & -5 \end{bmatrix} \begin{bmatrix} 2 \\ -1 \\ -1 \end{bmatrix} = \begin{bmatrix} 6 & 9 \\ 9 & 38 \end{bmatrix}$$

\mathbf{x}^\top \mathbf{x} $\mathbf{x}^\top \mathbf{x}$

Dividing by n yields
covariance matrix

PCA Formulation

PCA: find lower-dimensional representation of raw data

- \mathbf{X} is $n \times d$ (raw data)
- $\mathbf{Z} = \mathbf{XP}$ is $n \times k$ (reduced representation, PCA ‘scores’)
- \mathbf{P} is $d \times k$ (columns are k principal components)
- Variance / Covariance constraints

What constraints make sense in reduced representation?

- No feature correlation, i.e., all off-diagonals in $\mathbf{C}_\mathbf{Z}$ are zero
- Rank-ordered features by variance, i.e., sorted diagonals of $\mathbf{C}_\mathbf{Z}$

PCA Formulation

PCA: find lower-dimensional representation of raw data

- \mathbf{X} is $n \times d$ (raw data)
- $\mathbf{Z} = \mathbf{XP}$ is $n \times k$ (reduced representation, PCA ‘scores’)
- \mathbf{P} is $d \times k$ (columns are k principal components)
- Variance / Covariance constraints

\mathbf{P} equals the top k eigenvectors of \mathbf{C}_x

$$\mathbf{Z} = \mathbf{X} \mathbf{P}$$

PCA Solution

All covariance matrices have an eigendecomposition

- $\mathbf{C}_x = \mathbf{U}\Lambda\mathbf{U}^\top$ (eigendecomposition)
- \mathbf{U} is $d \times d$ (column are eigenvectors, sorted by their eigenvalues)
- Λ is $d \times d$ (diagonals are eigenvalues, off-diagonals are zero)

The d eigenvectors are orthonormal directions of max variance

- Associated eigenvalues equal variance in these directions
- 1st eigenvector is direction of max variance (variance is λ_1)

In lab, we'll use the `eigh` function from `numpy.linalg`

Choosing k

How should we pick the dimension of the new representation?

Visualization: Pick top 2 or 3 dimensions for plotting purposes

Other analyses: Capture ‘most’ of the variance in the data

- Recall that eigenvalues are variances in the directions specified by eigenvectors, and that eigenvalues are sorted

- Fraction of retained variance: $\frac{\sum_{i=1}^k \lambda_i}{\sum_{i=1}^d \lambda_i}$

Can choose k such that we retain some fraction of the variance, e.g., 95%

Other Practical Tips

PCA assumptions (linearity, orthogonality) not always appropriate

- Various extensions to PCA with different underlying assumptions, e.g., manifold learning, Kernel PCA, ICA

Centering is crucial, i.e., we must preprocess data so that all features have zero mean before applying PCA

PCA results dependent on scaling of data

- Data is sometimes rescaled in practice before applying PCA

PCA Algorithm



Orthogonal and Orthonormal Vectors

Orthogonal vectors are **perpendicular** to each other

- Equivalently, their dot product equals zero
- $\mathbf{a}^\top \mathbf{b} = 0$ and $\mathbf{d}^\top \mathbf{b} = 0$, but \mathbf{c} isn't orthogonal to others



$$\mathbf{a} = [1 \quad 0]^\top$$

$$\mathbf{b} = [0 \quad 1]^\top$$

$$\mathbf{c} = [1 \quad 1]^\top$$

$$\mathbf{d} = [2 \quad 0]^\top$$

Orthonormal vectors are orthogonal and have unit norm

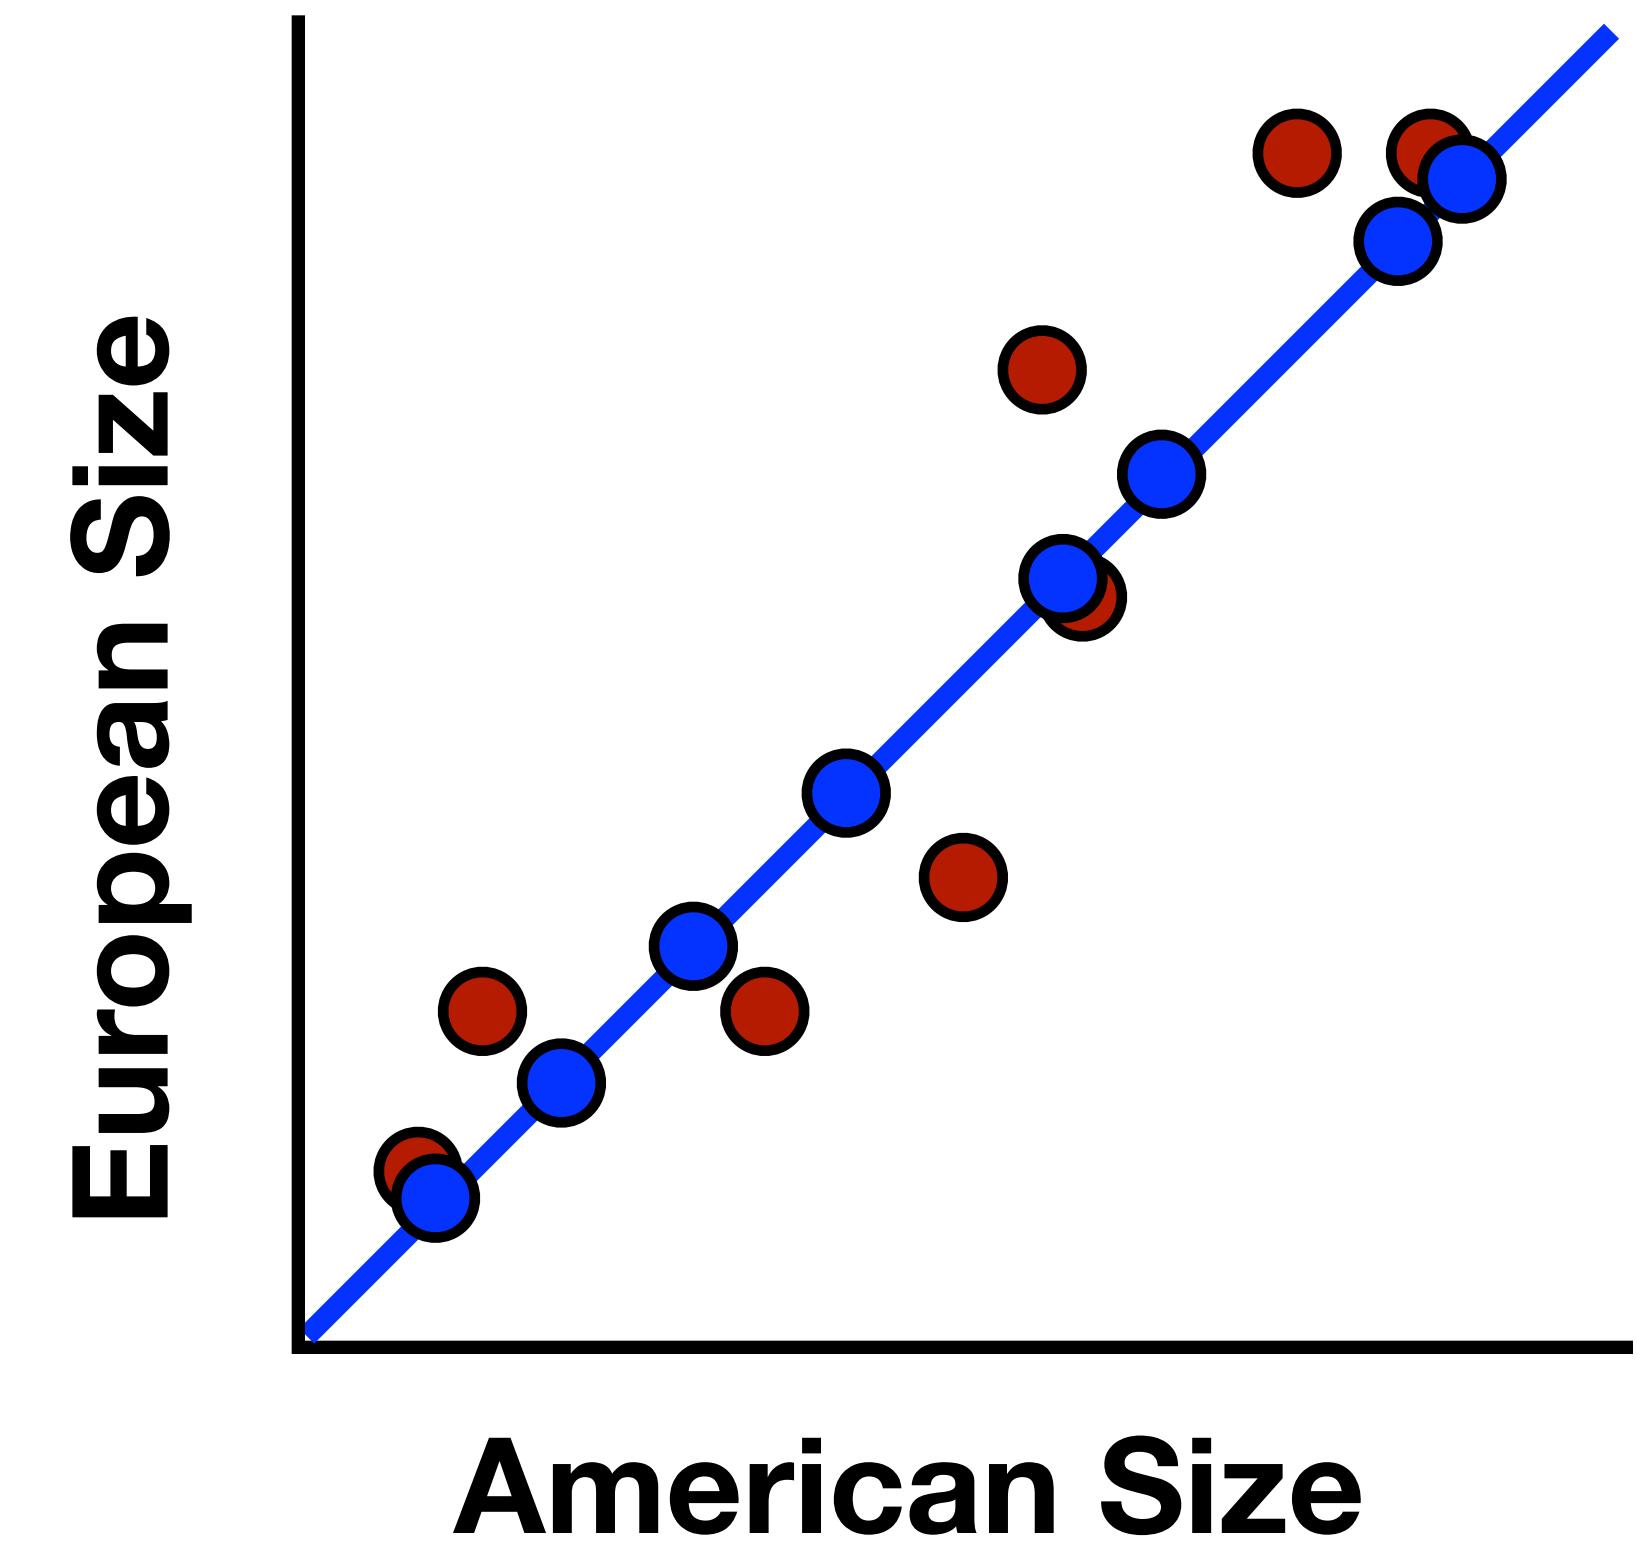
- \mathbf{a} and \mathbf{b} are orthonormal, but \mathbf{b} and \mathbf{d} are not orthonormal

PCA Iterative Algorithm

- $k = 1$: Find direction of max variance, project onto this direction
- Locations along this direction are the new 1D representation

More generally, for i in $\{1, \dots, k\}$:

- Find direction of max variance that is *orthonormal* to previously selected directions, project onto this direction
- Locations along this direction are the i th feature in new representation



PCA Derivation (Optional)



Eigendecomposition

All covariance matrices have an eigendecomposition

- $\mathbf{C}_x = \mathbf{U}\Lambda\mathbf{U}^\top$ (eigendecomposition)
- \mathbf{U} is $d \times d$ (column are eigenvectors, sorted by their eigenvalues)
- Λ is $d \times d$ (diagonals are eigenvalues, off-diagonals are zero)

Eigenvector / Eigenvalue equation: $\mathbf{C}_x \mathbf{u} = \lambda \mathbf{u}$

- By definition $\mathbf{u}^\top \mathbf{u} = 1$ (unit norm)

- Example:
$$\begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} 1 \\ 0 \end{bmatrix} = \begin{bmatrix} 1 \\ 0 \end{bmatrix} \Rightarrow \begin{array}{l} \text{eigenvector: } \mathbf{u} = [1 \quad 0]^\top \\ \text{eigenvalue: } \lambda = 1 \end{array}$$

PCA Formulation

PCA: find lower-dimensional representation of raw data

- \mathbf{X} is $n \times d$ (raw data)
- $\mathbf{Z} = \mathbf{XP}$ is $n \times k$ (reduced representation, PCA ‘scores’)
- \mathbf{P} is $d \times k$ (columns are k principal components)
- Variance / Covariance constraints

$$\begin{bmatrix} \mathbf{Z} \end{bmatrix} = \begin{bmatrix} \mathbf{X} \end{bmatrix} \begin{bmatrix} \mathbf{P} \end{bmatrix}$$

PCA Formulation, $k = 1$

PCA: find **one-dimensional** representation of raw data

- \mathbf{X} is $n \times d$ (raw data)
- $\mathbf{z} = \mathbf{X}\mathbf{p}$ is $n \times 1$ (reduced representation, PCA ‘scores’)
- \mathbf{p} is $d \times 1$ (columns are k principal components)
- Variance constraint

$$\sigma_{\mathbf{z}}^2 = \frac{1}{n} \sum_{i=1}^n (z^{(i)})^2 = \|\mathbf{z}\|_2^2$$

Goal: Maximizes variance, i.e., $\max_{\mathbf{p}} \sigma_{\mathbf{z}}^2$ where $\|\mathbf{p}\|_2 = 1$

Goal: Maximizes variance, i.e., $\max_{\mathbf{p}} \sigma_{\mathbf{z}}^2$ where $\|\mathbf{p}\|_2 = 1$

Relationship between Euclidean distance and dot product

Definition: $\mathbf{z} = \mathbf{X}\mathbf{p}$

$$\begin{aligned}\sigma_{\mathbf{z}}^2 &= \|\mathbf{z}\|_2^2 \\ &= \mathbf{z}^\top \mathbf{z} \\ &= (\mathbf{X}\mathbf{p})^\top (\mathbf{X}\mathbf{p}) \\ &= \mathbf{p}^\top \mathbf{X}^\top \mathbf{X}\mathbf{p} \\ &= \boxed{n}\mathbf{p}^\top \mathbf{C}_x \mathbf{p}\end{aligned}$$

Transpose property: $(\mathbf{X}\mathbf{p})^\top = \mathbf{p}^\top \mathbf{X}^\top$; associativity of multiply

Definition: $\mathbf{C}_x = \frac{1}{n} \mathbf{X}^\top \mathbf{X}$

Restated Goal: $\max_{\mathbf{p}} \mathbf{p}^\top \mathbf{C}_x \mathbf{p}$ where $\|\mathbf{p}\|_2 = 1$

Connection to Eigenvectors

Recall eigenvector / eigenvalue equation: $\mathbf{C}_x \mathbf{u} = \lambda \mathbf{u}$

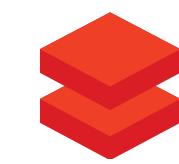
- By definition $\mathbf{u}^\top \mathbf{u} = 1$, and thus $\mathbf{u}^\top \mathbf{C}_x \mathbf{u} = \lambda$
- But this is the expression we're optimizing, and thus maximal variance achieved when \mathbf{p} is top eigenvector of \mathbf{C}_x

Similar arguments can be used for $k > 1$

Restated Goal: $\max_{\mathbf{p}} \mathbf{p}^\top \mathbf{C}_x \mathbf{p}$ where $\|\mathbf{p}\|_2 = 1$

Distributed PCA



UCLA 
 databricks™

Computing PCA Solution

Given: $n \times d$ matrix of uncentered raw data

Goal: Compute $k \ll d$ dimensional representation

Step 1: Center Data

Step 2: Compute Covariance or Scatter Matrix

- $\frac{1}{n} \mathbf{X}^T \mathbf{X}$ versus $\mathbf{X}^T \mathbf{X}$

Step 3: Eigendecomposition

Step 4: Compute PCA Scores

$$\begin{bmatrix} \mathbf{Z} \\ \vdots \end{bmatrix} = \begin{bmatrix} \mathbf{P} \\ \vdots \end{bmatrix}^T \begin{bmatrix} \mathbf{X} \\ \vdots \end{bmatrix}$$

PCA at Scale

Case 1: Big n and Small d

- $O(d^2)$ local storage, $O(d^3)$ local computation, $O(dk)$ communication
- Similar strategy as closed-form linear regression

Case 2: Big n and Big d

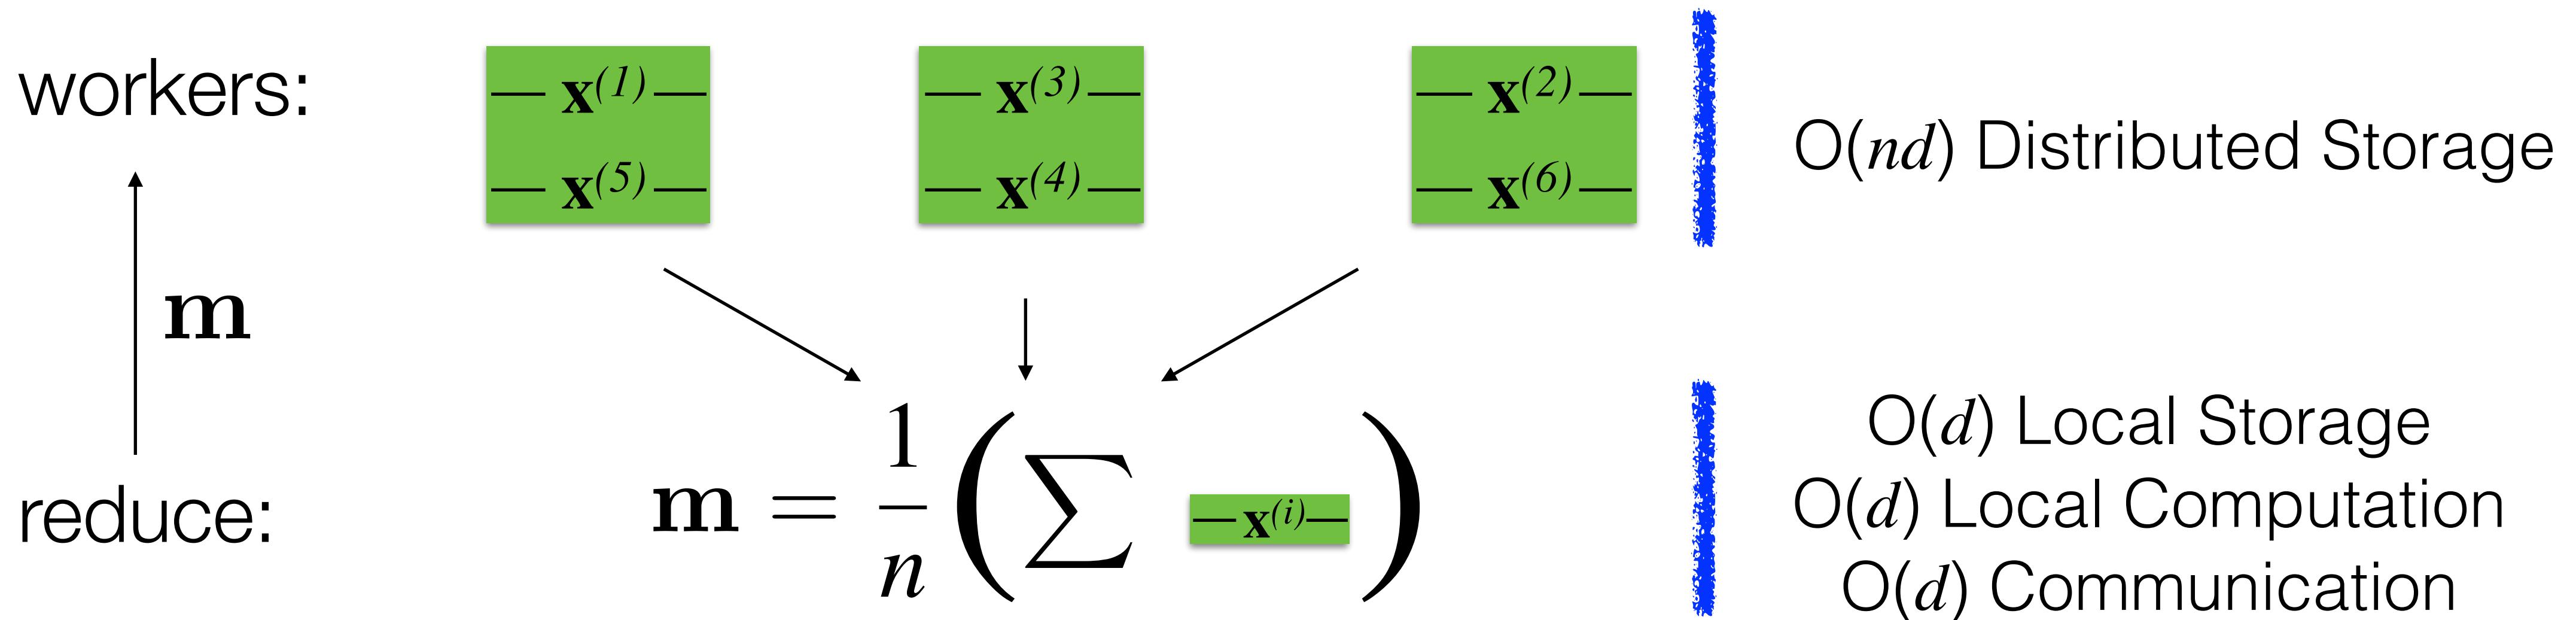
- $O(dk + n)$ local storage, computation;
 $O(dk + n)$ communication
- Iterative algorithm

$$\begin{bmatrix} \mathbf{Z} \end{bmatrix} = \begin{bmatrix} \mathbf{X} \end{bmatrix} \begin{bmatrix} \mathbf{P} \end{bmatrix}$$

Step 1: Center Data

- Compute d feature means, $\mathbf{m} \in \mathbb{R}^d$
- Communicate \mathbf{m} to all workers

Example: $n = 6$; 3 workers



Step 1: Center Data

- Compute d feature means, $\mathbf{m} \in \mathbb{R}^d$
- Communicate \mathbf{m} to all workers
- Subtract \mathbf{m} from each data point

Example: $n = 6$; 3 workers

workers:

$$\begin{array}{c} \text{--- } \mathbf{x}^{(1)} \text{ ---} \\ \text{--- } \mathbf{x}^{(5)} \text{ ---} \end{array}$$

↓

$$\begin{array}{c} \text{--- } \mathbf{x}^{(3)} \text{ ---} \\ \text{--- } \mathbf{x}^{(4)} \text{ ---} \end{array}$$

↓

$$\begin{array}{c} \text{--- } \mathbf{x}^{(2)} \text{ ---} \\ \text{--- } \mathbf{x}^{(6)} \text{ ---} \end{array}$$

↓

$O(nd)$ Distributed Storage

map:

$$\text{--- } \mathbf{x}^{(i)} \text{ ---} - \mathbf{m}$$

$$\text{--- } \mathbf{x}^{(i)} \text{ ---} - \mathbf{m}$$

$$\text{--- } \mathbf{x}^{(i)} \text{ ---} - \mathbf{m}$$

$O(d)$ Local Computation

Step 2: Compute Scatter Matrix ($\mathbf{X}^T \mathbf{X}$)

- Compute matrix product via outer products (just like we did for closed-form linear regression!)

$$\begin{bmatrix} 9 & 3 & 5 \\ 4 & 1 & 2 \end{bmatrix} \begin{bmatrix} 1 & 2 \\ 3 & -5 \\ 2 & 3 \end{bmatrix} = \begin{bmatrix} & \\ & \end{bmatrix}$$

$$\begin{bmatrix} 9 & 18 \\ 4 & 8 \end{bmatrix}$$

Step 2: Compute Scatter Matrix ($\mathbf{X}^\top \mathbf{X}$)

- Compute matrix product via outer products (just like we did for closed-form linear regression!)

$$\begin{bmatrix} 9 & 3 & 5 \\ 4 & 1 & 2 \end{bmatrix} \begin{bmatrix} 1 & 2 \\ 3 & -5 \\ 2 & 3 \end{bmatrix} = \quad \quad \quad$$

$$\begin{bmatrix} 9 & 18 \\ 4 & 8 \end{bmatrix} + \begin{bmatrix} 9 & -15 \\ 3 & -5 \end{bmatrix}$$

Step 2: Compute Scatter Matrix ($\mathbf{X}^\top \mathbf{X}$)

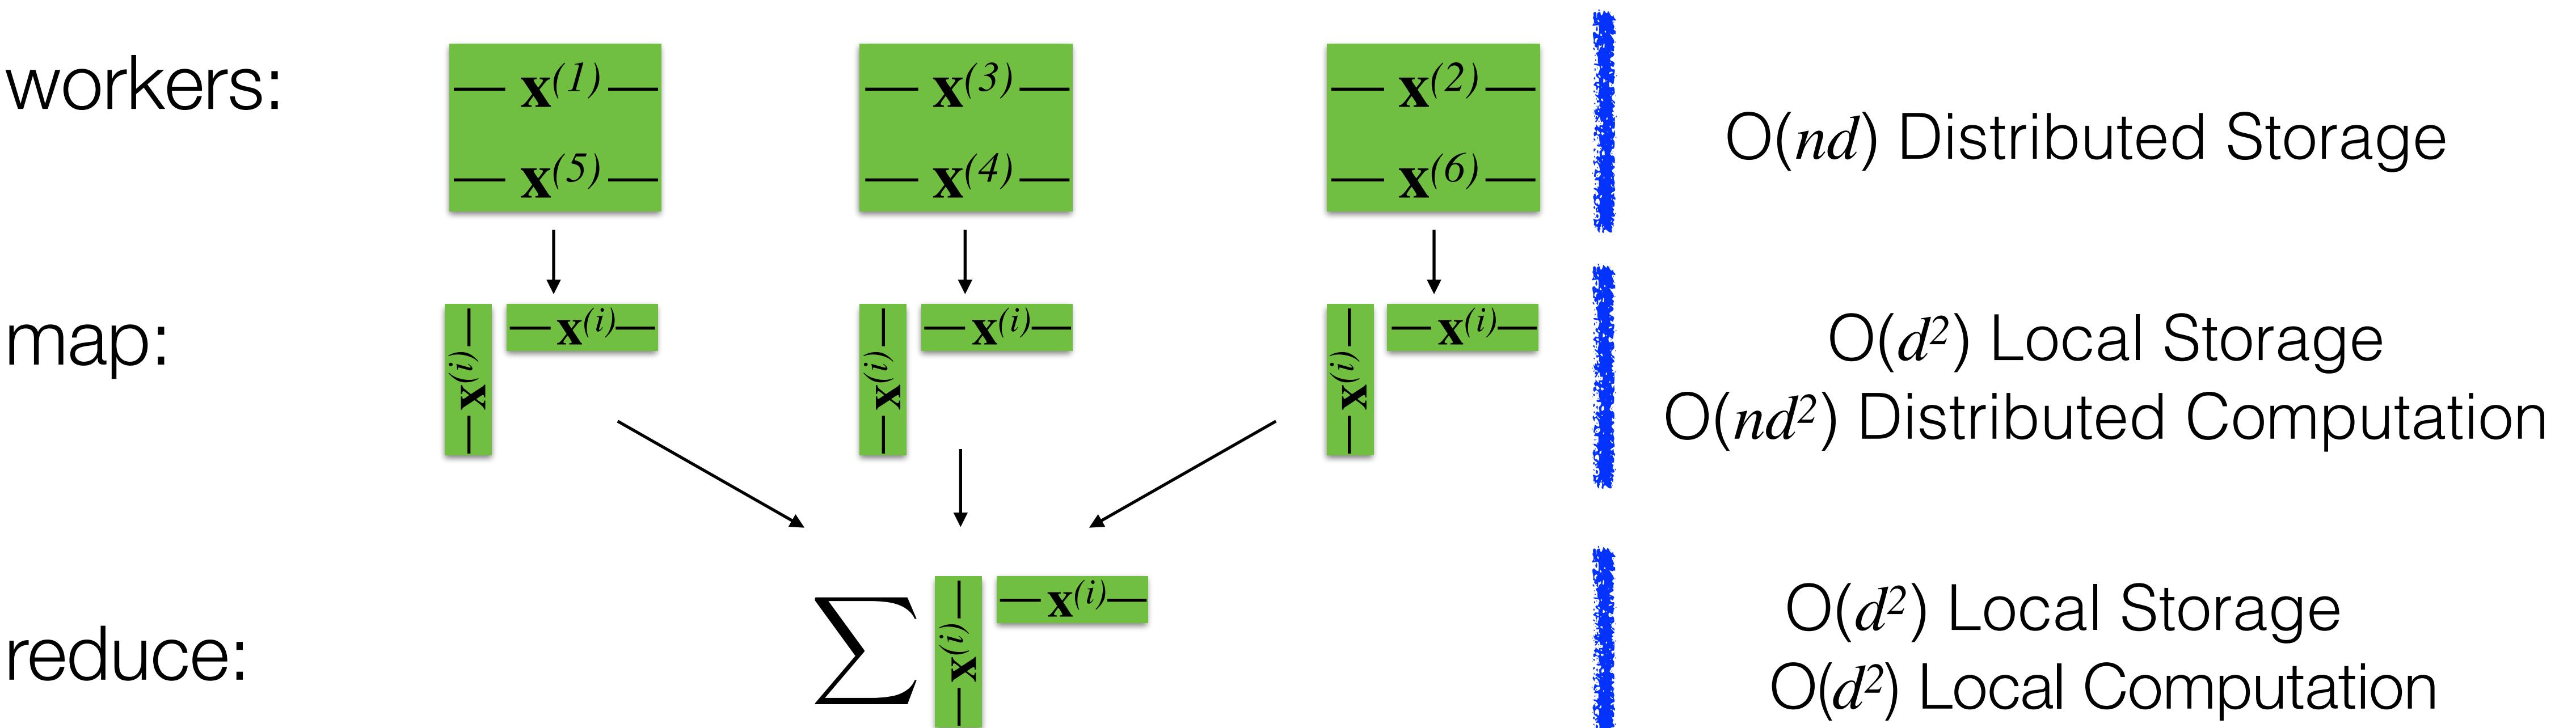
- Compute matrix product via outer products (just like we did for closed-form linear regression!)

$$\begin{bmatrix} 9 & 3 & 5 \\ 4 & 1 & 2 \end{bmatrix} \begin{bmatrix} 1 & 2 \\ 3 & -5 \\ 2 & 3 \end{bmatrix} = \begin{bmatrix} 28 & 18 \\ 11 & 9 \end{bmatrix}$$

$$\begin{bmatrix} 9 & 18 \\ 4 & 8 \end{bmatrix} + \begin{bmatrix} 9 & -15 \\ 3 & -5 \end{bmatrix} + \begin{bmatrix} 10 & 15 \\ 4 & 6 \end{bmatrix}$$

$$\mathbf{X}^\top \mathbf{X} = \begin{matrix} & n \\ d & \left| \begin{array}{c} \mathbf{x}^{(1)} \\ \vdots \\ \mathbf{x}^{(2)} \\ \cdots \\ \mathbf{x}^{(n)} \end{array} \right| \end{matrix} = \sum_{i=1}^n \begin{matrix} & n \\ d & \left| \begin{array}{c} \mathbf{x}^{(1)} \\ \vdots \\ \mathbf{x}^{(2)} \\ \vdots \\ \mathbf{x}^{(n)} \end{array} \right| \end{matrix}$$

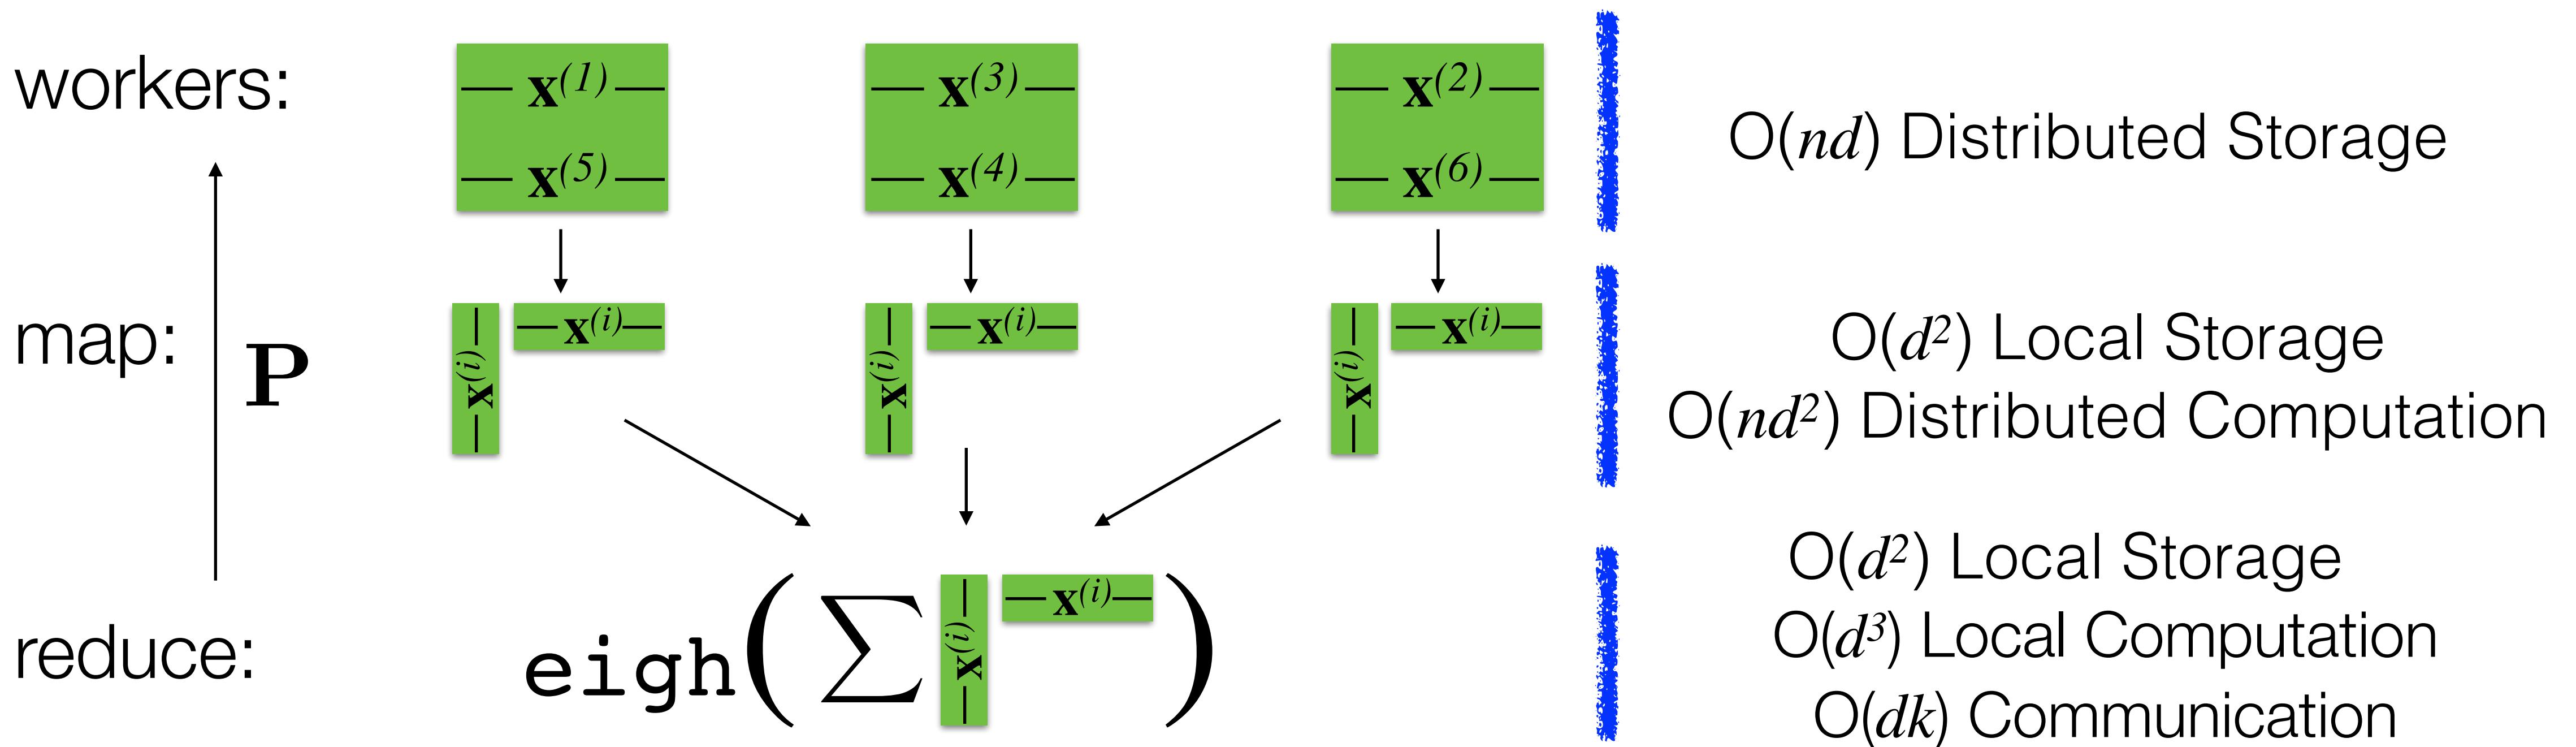
Example: $n = 6$; 3 workers



Step 3: Eigendecomposition

- Perform locally since d is small
- Communicate k principal components ($\mathbf{P} \in \mathbb{R}^{d \times k}$) to workers

Example: $n = 6$; 3 workers



Step 4: Compute PCA Scores

- Multiply each point by principal components, \mathbf{P}

Example: $n = 6$; 3 workers

workers:

$$\begin{array}{c} \text{--- } \mathbf{x}^{(1)} \text{ ---} \\ \text{--- } \mathbf{x}^{(5)} \text{ ---} \end{array}$$

$$\begin{array}{c} \text{--- } \mathbf{x}^{(3)} \text{ ---} \\ \text{--- } \mathbf{x}^{(4)} \text{ ---} \end{array}$$

$$\begin{array}{c} \text{--- } \mathbf{x}^{(2)} \text{ ---} \\ \text{--- } \mathbf{x}^{(6)} \text{ ---} \end{array}$$

map:

$$\begin{array}{c} \text{--- } \mathbf{p}^{(1)} \text{ ---} \\ \text{--- } \mathbf{p}^{(2)} \text{ ---} \end{array} \quad \begin{array}{c} | \\ \mathbf{x}^{(i)} \end{array}$$

$$\begin{array}{c} \text{--- } \mathbf{p}^{(1)} \text{ ---} \\ \text{--- } \mathbf{p}^{(2)} \text{ ---} \end{array} \quad \begin{array}{c} | \\ \mathbf{x}^{(i)} \end{array}$$

$$\begin{array}{c} \text{--- } \mathbf{p}^{(1)} \text{ ---} \\ \text{--- } \mathbf{p}^{(2)} \text{ ---} \end{array} \quad \begin{array}{c} | \\ \mathbf{x}^{(i)} \end{array}$$

$O(nd)$ Distributed Storage

$O(dk)$ Local Computation

Distributed PCA, Part II

(Optional)



PCA at Scale

Case 1: Big n and Small d

- $O(d^2)$ local storage, $O(d^3)$ local computation, $O(dk)$ communication
- Similar strategy as closed-form linear regression

Case 2: Big n and Big d

- $O(dk + n)$ local storage, computation; $O(dk + n)$ communication
- Iterative algorithm

$$\begin{bmatrix} \mathbf{Z} \end{bmatrix} = \begin{bmatrix} \mathbf{X} \end{bmatrix} \begin{bmatrix} \mathbf{P} \end{bmatrix}$$

An Iterative Approach

We can use algorithms that rely on a **sequence of matrix-vector products** to compute top k eigenvectors (\mathbf{P})

- E.g., Krylov subspace or random projection methods

Krylov subspace methods (used in MLlib) iteratively compute $\mathbf{X}^\top \mathbf{X} \mathbf{v}$ for some $\mathbf{v} \in \mathbb{R}^d$ provided by the method

- Requires $O(k)$ passes over the data and $O(dk)$ local storage
- We don't need to compute the covariance matrix!

Repeat for $O(k)$ iterations:

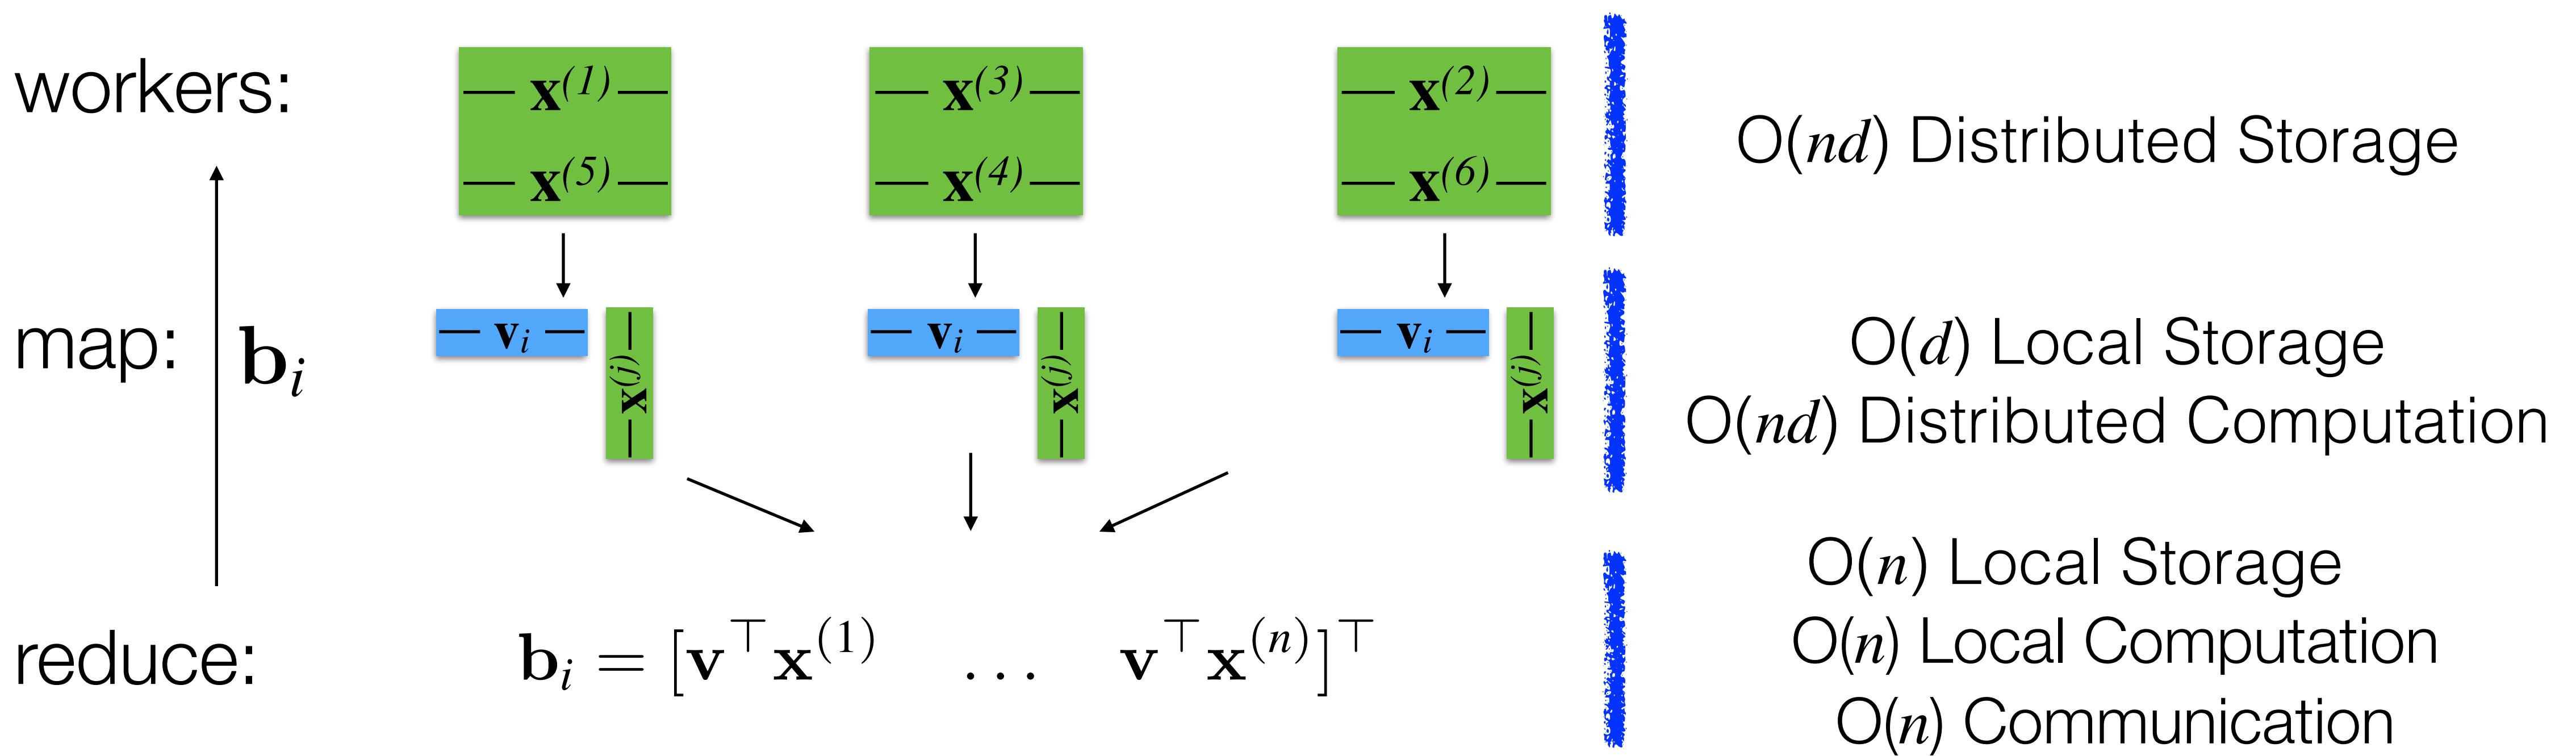
- 1. Communicate $\mathbf{v}_i \in \mathbb{R}^d$ to all workers
- 2. Compute $\mathbf{q}_i = \mathbf{X}^\top \mathbf{X} \mathbf{v}_i$ in a distributed fashion
 - Step 1: $\mathbf{b}_i = \mathbf{X} \mathbf{v}_i$
 - Step 2: $\mathbf{q}_i = \mathbf{X}^\top \mathbf{b}_i$
- 3. Driver uses \mathbf{q}_i to update estimate of \mathbf{P}

Compute $\mathbf{q}_i = \mathbf{X}^\top \mathbf{X} \mathbf{v}_i$ in a distributed fashion

- $\mathbf{b}_i = \mathbf{X} \mathbf{v}_i$: each component is dot product, then concatenate

```
> b = np.array(trainData.map(dotProduct).collect())
```

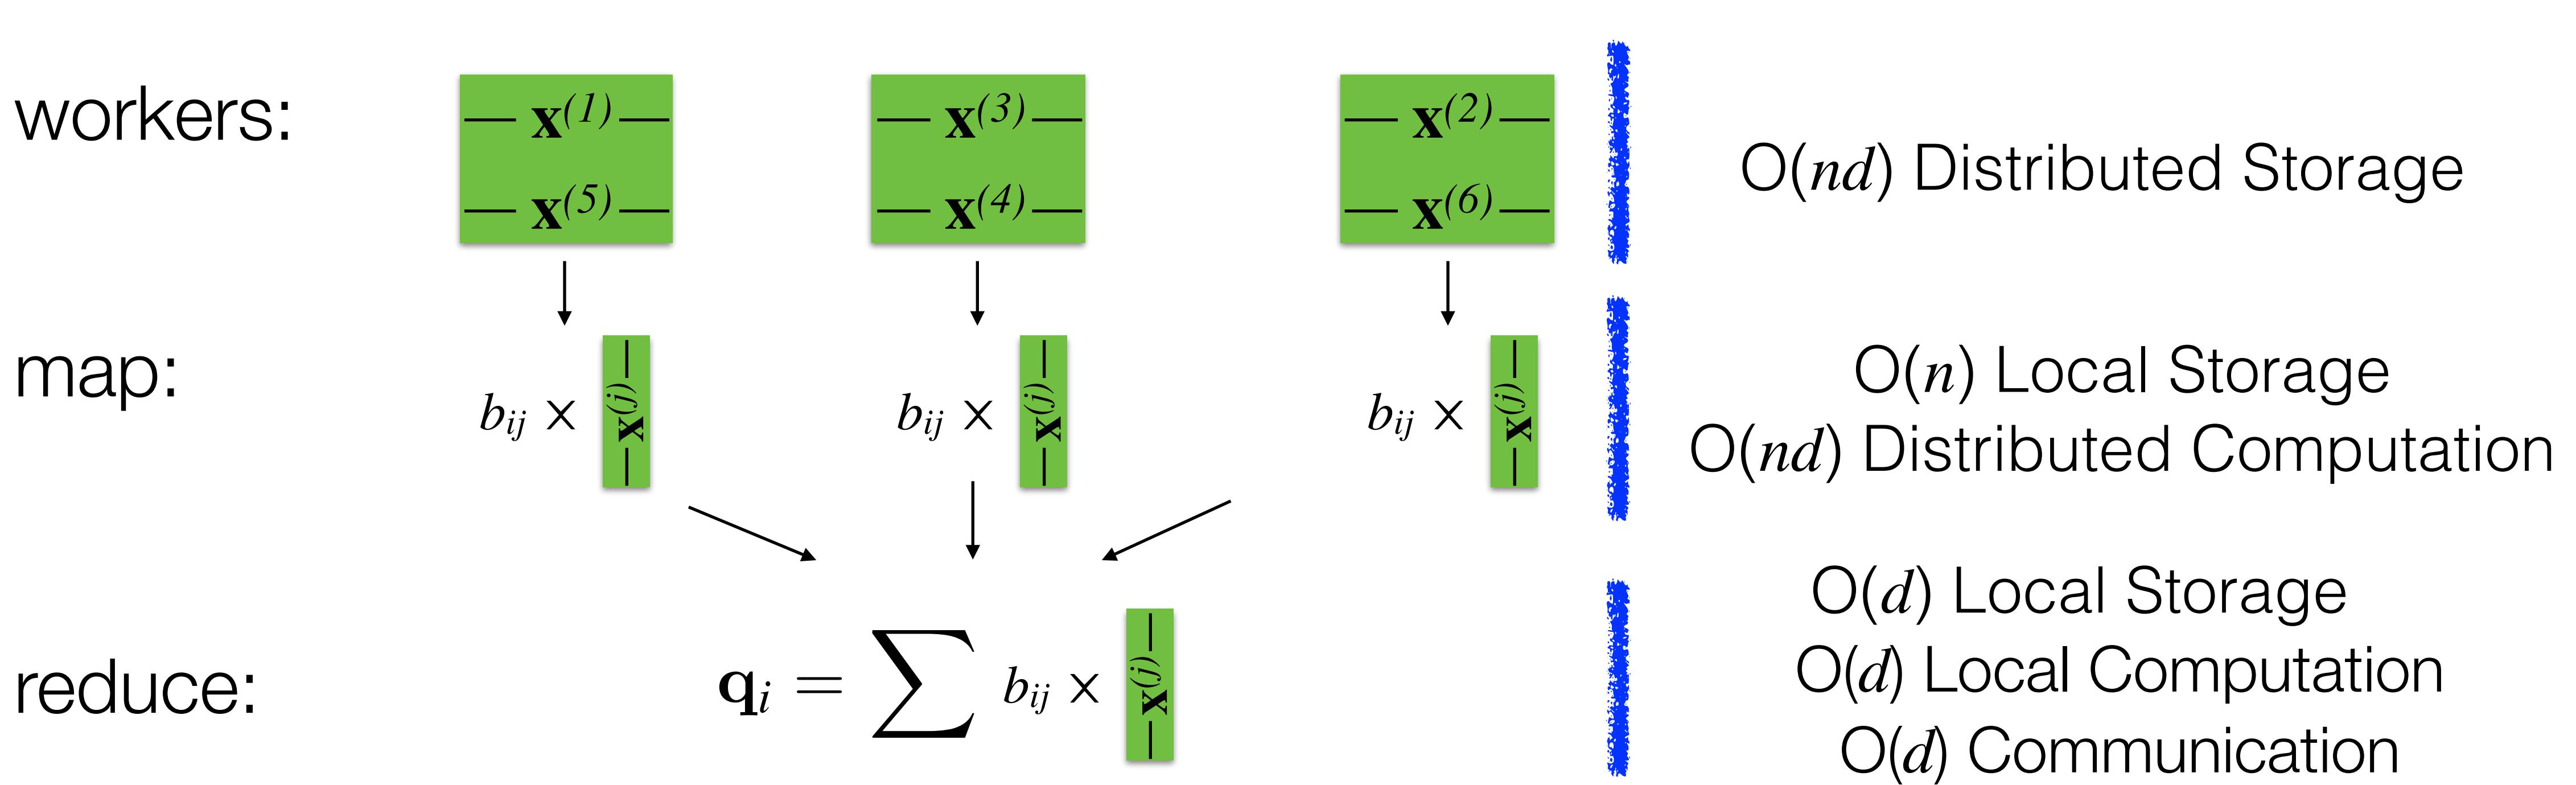
Example: $n = 6$; 3 workers



Compute $\mathbf{q}_i = \mathbf{X}^\top \mathbf{X} \mathbf{v}_i$ in a distributed fashion

- $\mathbf{b}_i = \mathbf{X} \mathbf{v}_i$: each component is dot product, then concatenate
- $\mathbf{q}_i = \mathbf{X}^\top \mathbf{b}_i$: sum of rescaled data points $\mathbf{q}_i = \sum_{j=1}^n b_{ij} \mathbf{x}^{(j)}$

Example: $n = 6$; 3 workers



Compute $\mathbf{q}_i = \mathbf{X}^\top \mathbf{X} \mathbf{v}_i$ in a distributed fashion

- $\mathbf{b}_i = \mathbf{X} \mathbf{v}_i$: each component is dot product, then concatenate
- $\mathbf{q}_i = \mathbf{X}^\top \mathbf{b}_i$: sum of rescaled data points $\mathbf{q}_i = \sum_{j=1}^n b_{ij} \mathbf{x}^{(j)}$

```
> q = trainData.map(rescaleByBi)
    .reduce(sumVectors)
```

workers:

$$\begin{array}{|c|}\hline -\mathbf{x}^{(1)}- \\ \hline -\mathbf{x}^{(5)}- \\ \hline\end{array}$$

$$\begin{array}{|c|}\hline -\mathbf{x}^{(3)}- \\ \hline -\mathbf{x}^{(4)}- \\ \hline\end{array}$$

$$\begin{array}{|c|}\hline -\mathbf{x}^{(2)}- \\ \hline -\mathbf{x}^{(6)}- \\ \hline\end{array}$$

map:

$$b_{ij} \times \begin{array}{|c|}\hline -\mathbf{x}^{(j)}- \\ \hline\end{array}$$

$$b_{ij} \times \begin{array}{|c|}\hline -\mathbf{x}^{(j)}- \\ \hline\end{array}$$

$$b_{ij} \times \begin{array}{|c|}\hline -\mathbf{x}^{(j)}- \\ \hline\end{array}$$

reduce:

$$\mathbf{q}_i = \sum b_{ij} \times \begin{array}{|c|}\hline -\mathbf{x}^{(j)}- \\ \hline\end{array}$$

$O(nd)$ Distributed Storage

$O(n)$ Local Storage

$O(nd)$ Distributed Computation

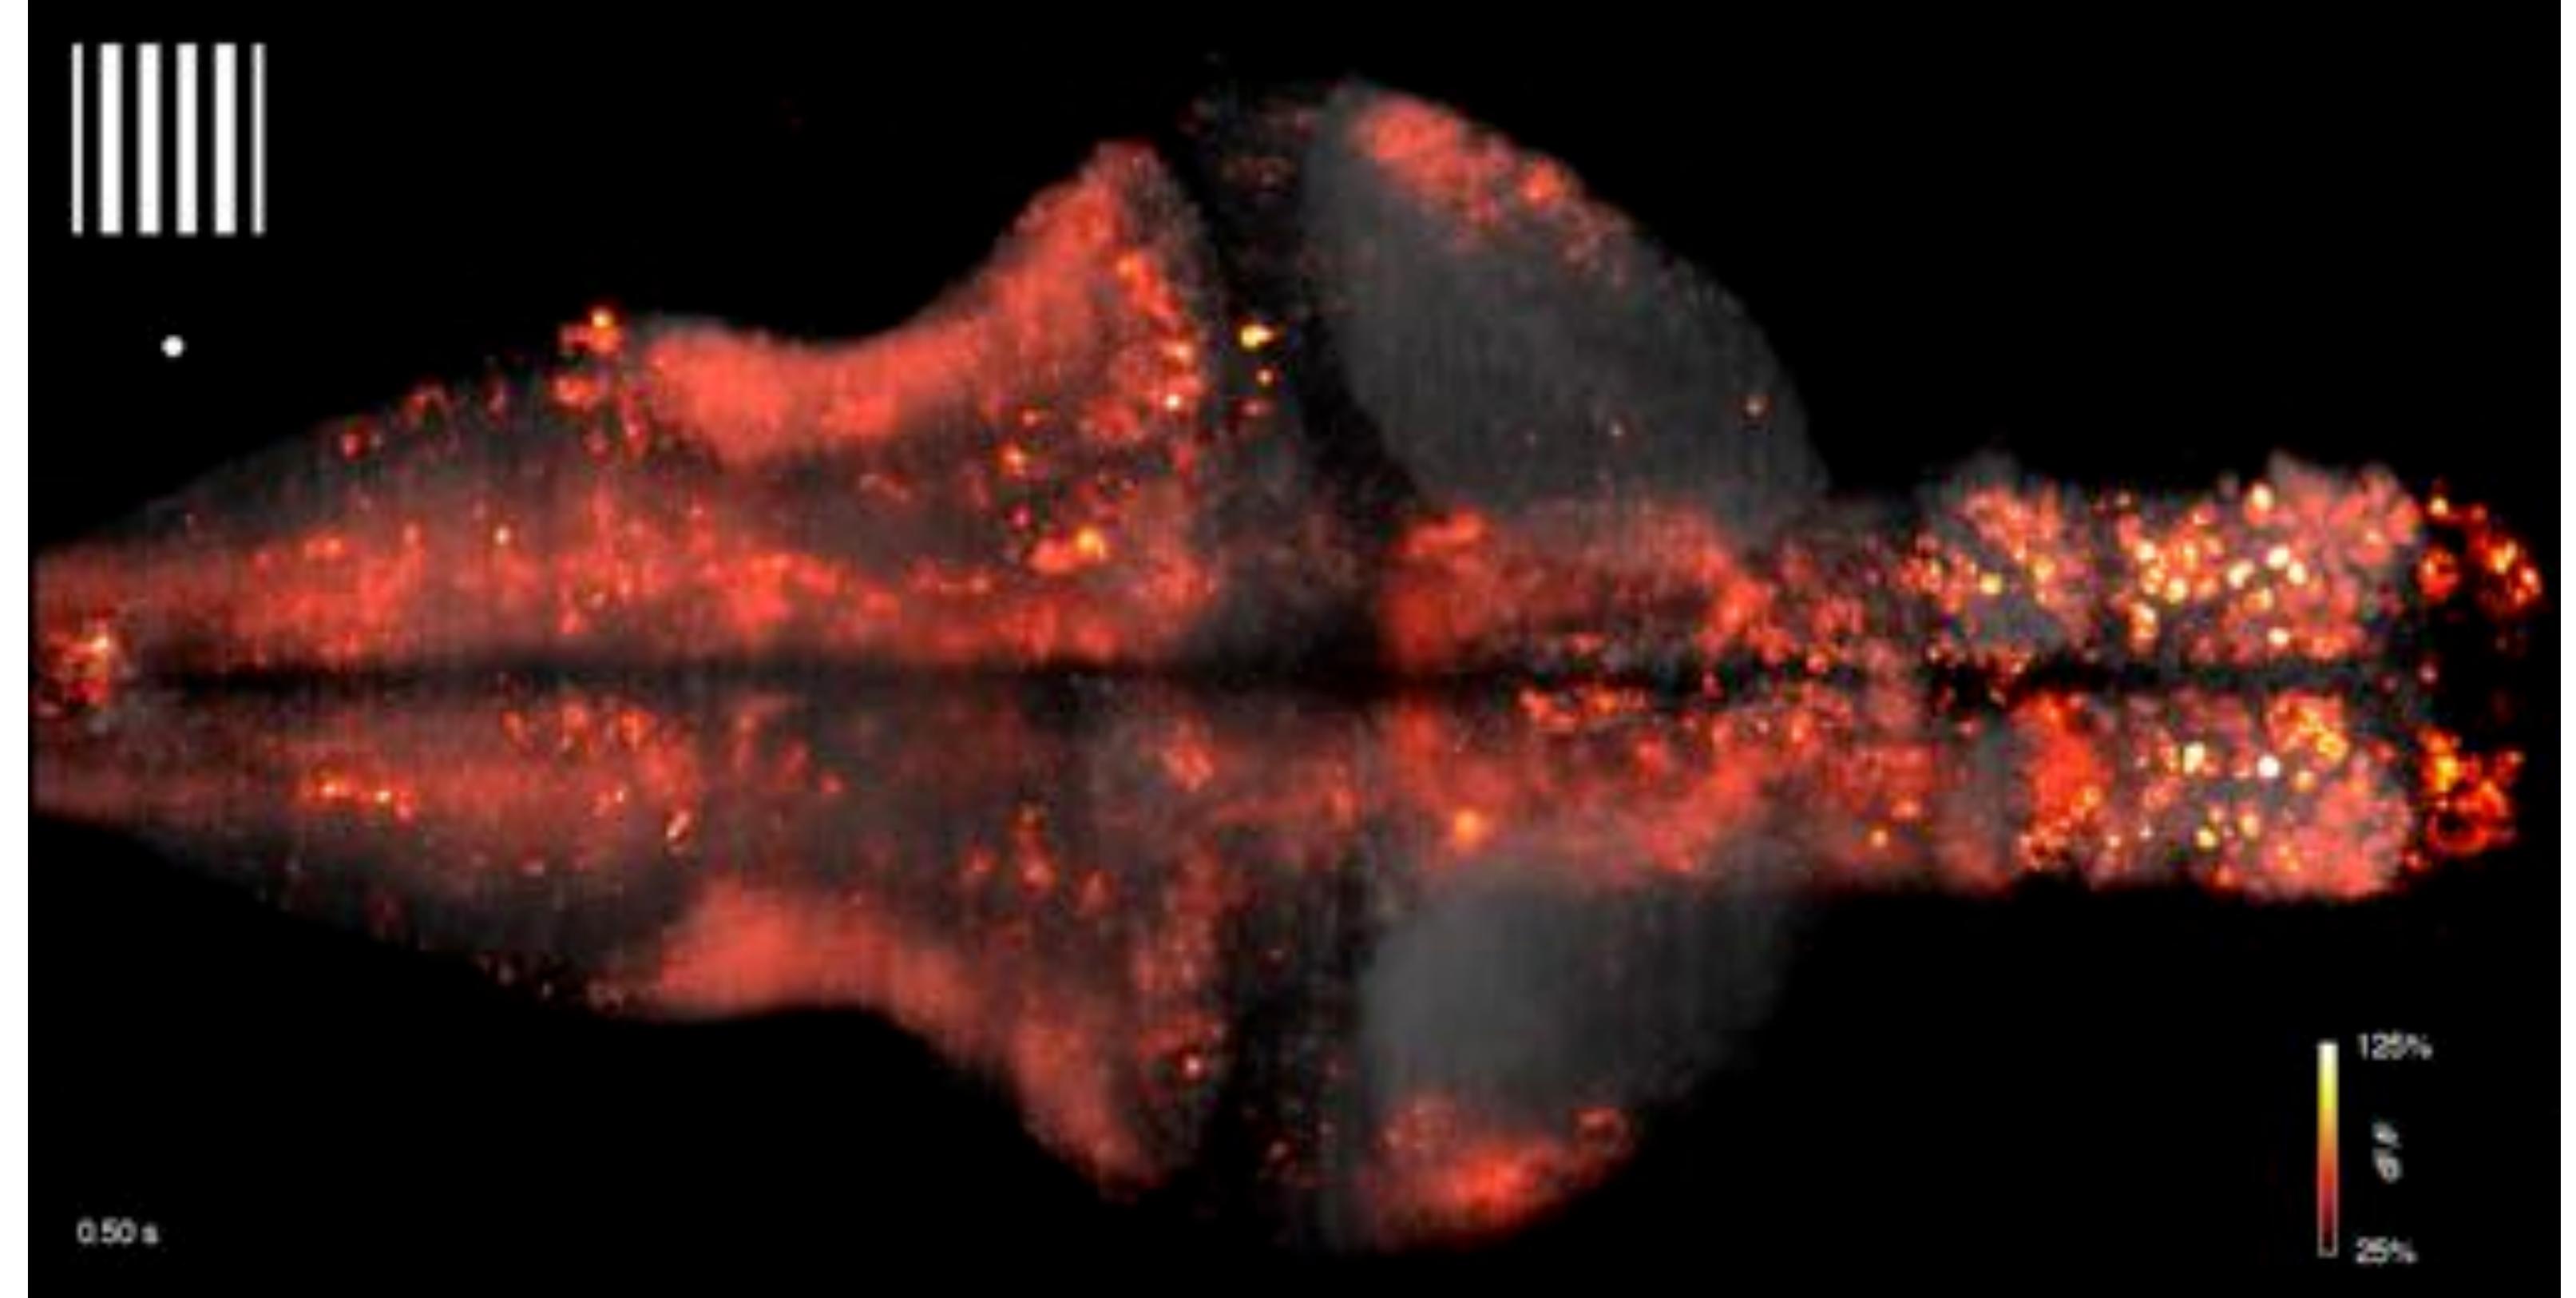
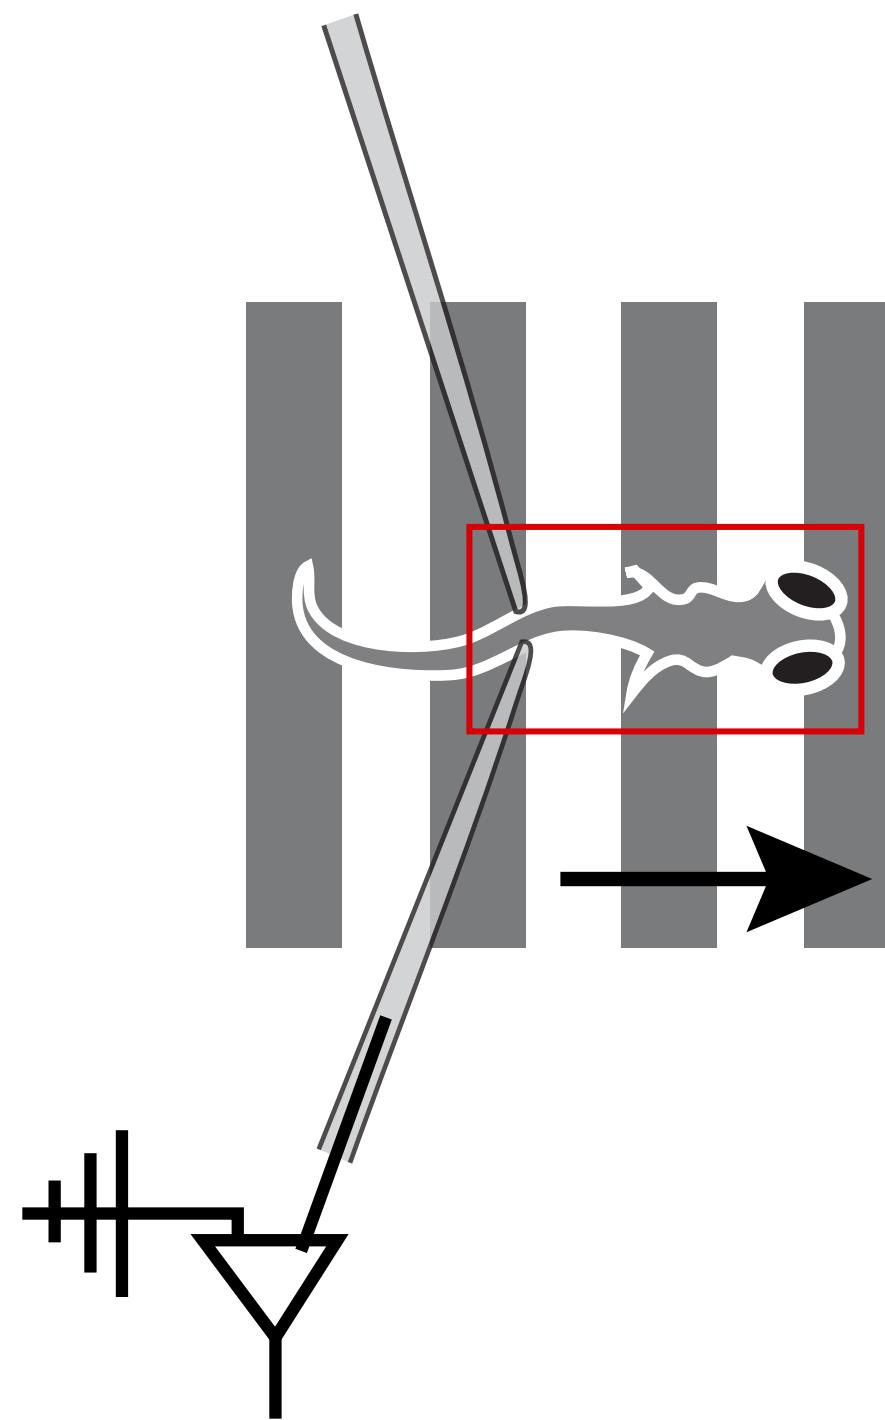
$O(d)$ Local Storage

$O(d)$ Local Computation

$O(d)$ Communication

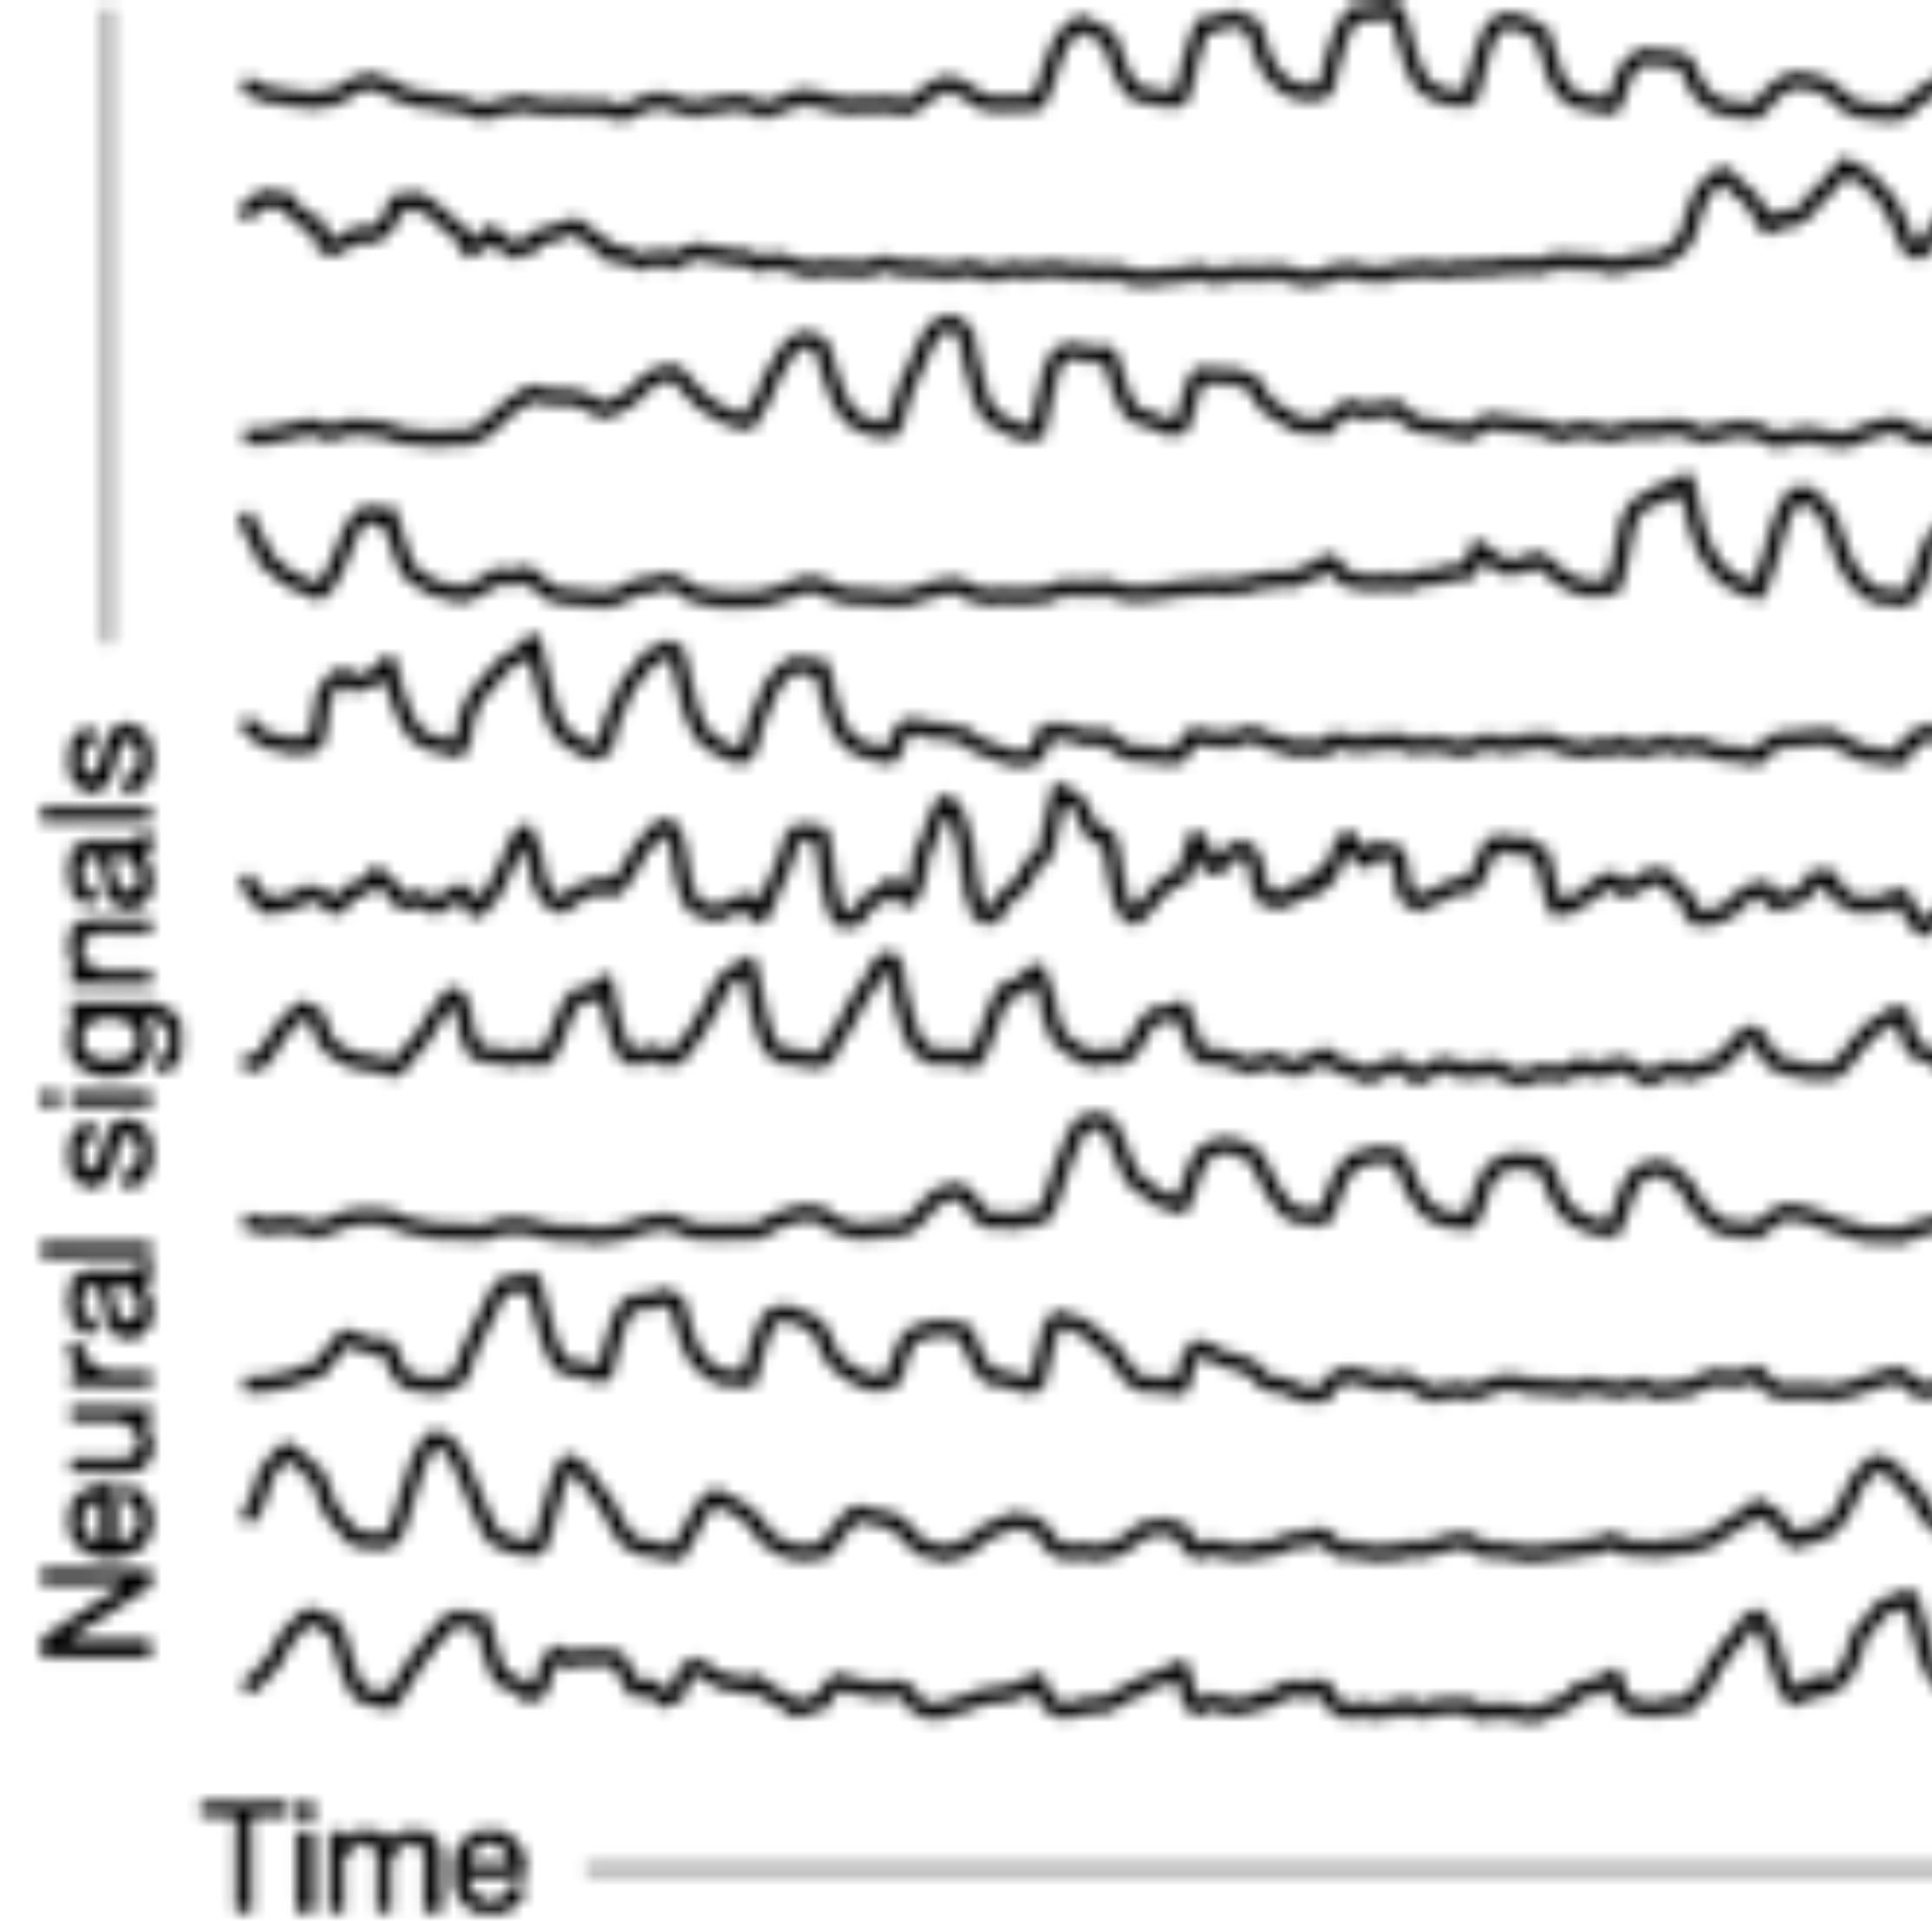
Lab Preview

Vladimirov et al.,
2014



Which areas are active at which times?

Which neuronal populations are activated by different directions of the stimulus?



Given

Collection of neural
time series

Goal

Find representations of
data that reveal how
responses are organized
across space and time

